

Certified support cuts for shared nonlinear expressions and quadratic blocks

Abstract

Global solvers for nonconvex mixed-integer nonlinear programs relax nonlinear expressions one at a time and lose their dependence through shared variables. Support cuts of the hull of their joint graph keep it. We study when such cuts add strength, how to compute them exactly for quadratic blocks, and how to keep them valid. Gluing exact pair hulls on shared moments is known to lose strength; we quantify the loss. For a family of quadratic directions on a three-variable path, the glued relaxation misses either nothing or $\delta^2/2$, depending on whether two point sets at distance δ interleave; an interface that shares κ moments gives no positive bound exactly when the local zero sets alternate at least $\kappa + 1$ times; and from four variables on, joint blocks can be strictly stronger than dense semidefinite relaxations. For quadratic stars with k leaves, m center-leaf rows and m_0 center rows we give an exact rational support algorithm with $O(m_0 + (m + k) \log(m + k))$ operations. Eliminating the model rows gives cuts in the original variables, valid as stored if their bound is certified for the exact binary64 direction, rounding is corrected and the stored row is checked. Every recorded cut passed replay, whereas sample-based constants would have cut off feasible solutions. On MINLPLib models the cuts did not help SCIP. On constructed path and star families the cuts closed nearly all of SCIP’s root gap, and a configuration fixed in advance solved all 20 fresh path instances, against at most 10 for SCIP and 5 for Gurobi.

Keywords: simultaneous convexification; support cuts; nonconvex quadratic programming; certified cuts; safe rounding; mixed-integer nonlinear programming

Mathematics Subject Classification (2020): 90C26, 90C20, 90C11, 65G30

1 Introduction

Global solvers for nonconvex mixed-integer nonlinear programs (MINLPs) bound the optimal value by relaxing every nonlinear expression. The standard construction decomposes each expression into elementary operations, introduces an auxiliary variable for each, and relaxes each operation on its own [90, 117, 119]. This construction is general, but it relaxes expressions as if they were unrelated. In many models they are not: the same variables appear in several nonlinear constraints, and linear constraints restrict how these variables can move together. Pooling and blending models, heat and mass exchange networks and gas networks are typical examples. A relaxation that treats such expressions separately can be much weaker than one that treats them jointly.

The strongest convex relaxation of several functions $f_1, \dots, f_p$ of the same variables x on a common domain D is the convex hull of their joint graph $\{(x, f_1(x), \dots, f_p(x)) : x \in D\}$. Its valid linear inequalities are *support cuts*: for a direction (a, b) , the inequality $a^\top x + b^\top w \geq \min_{x \in D} \{a^\top x + b^\top F(x)\}$ is valid, where w_j stands for $f_j(x)$, and all such inequalities together describe the hull. This simultaneous convexification has been studied for vectors of functions [118], [83, Proposition 1], including vectors of univariate functions [67], for bilinear and composite functions [36, 64, 130], and in applications [83]. Turning it into a solver component raises four questions.

- (1) *Which blocks are worth relaxing jointly?* In particular, is it enough to relax every pair of interacting variables exactly, or do larger blocks carry information that pairs miss?
- (2) *Exactness.* The support value is the minimum of a nonconvex function. For which blocks can it be computed exactly, so that the cuts are as strong as possible?
- (3) *Integration and validity.* Support cuts live in the space of (x, w) , and introducing the variables w changes the model that the solver presolves, propagates and branches on. Can the cuts be stated in the original variables? And since directions are found numerically and the solver stores binary64 rows, which part of the computation must be exact for the stored row to be valid?
- (4) *Usefulness.* Modern solvers already detect quadratic, bilinear and convex structure and have their own cut families for it [13]. Do joint support cuts improve them?

We answer these questions for quadratic blocks in depth and for shared nonlinear expressions more generally. The contributions are the following.

- *Limits of composing pair hulls* (Section 3). On a path of three variables, exact pair hulls identified on the first two moments of the shared variable miss a gap of $1/128$ that a single support cut in the existing coordinates closes. For a family of such paths the gap is positive exactly when two two-point sets interleave, and it then equals $\delta^2/2$, where δ is their distance. An interface that identifies κ moments of the shared variable gives no positive bound exactly when the local zero sets alternate at least $\kappa + 1$ times, a reading of the classical description of Radon partitions on the moment curve, and for every κ there is a quadratic instance with a relative gap of 100%. On a three-variable path, a dense semidefinite relaxation with the McCormick inequalities of all products is exact, by a theorem of Burer et al. [27]; there the advantage of the joint block over its pairs is one of representation, which support cuts obtain in the coordinates that the model already has. From four variables on, the dense relaxation is no longer exact for every objective: on a four-variable path and on our full-gap instances, joint support cuts are strictly stronger, on the latter with a relative gap of nearly 100% for the dense relaxation. Through Lagrangian decomposition we also relate the gain from merging pairs into a star to the glued relaxation.
- *Exact support for constrained stars* (Section 4). For quadratic stars whose affine rows couple the center with one leaf at a time, we give an exact rational algorithm with $O(m_0 + (m + k) \log(m + k))$ operations for k leaves, m center–leaf rows and m_0 rows in the center alone; on instances with $m = \Theta(k)$ this is optimal for algebraic computation trees. The box-constrained case is a special case of the forest algorithm of Del Pia and Khajavirad [37]; to our knowledge, no exact algorithm for the case with center–leaf rows has been published. For polytopes of fixed dimension we state face enumeration together with the degenerate cases that a certificate must cover.
- *Certified cuts in the original variables* (Sections 5 and 6). Minimizing nonnegative combinations of nonlinear rows over their common block domain and eliminating the rows gives cuts in the original variables, with no auxiliary variables. Together these cuts describe the projection of the joint-graph relaxation, and we give conditions under which this projection is the convex hull of the set that the rows define, and examples where it is not. Validity of a stored cut rests on four checks: a certified lower bound over the whole block domain for the exact binary64 direction, exact elimination, bound-corrected rounding, and a comparison with the row that the solver stores; the direction itself may be found by any numerical method.

- *Implementation and evaluation* (Sections 7 and 8). We implemented the cuts as a SCIP separator whose model import is checked to reproduce the exact binary64 source expressions and their domains, and whose every cut can be replayed. We evaluate it with prospective protocols on MINLPLib models and on constructed families with path and star structure, against SCIP’s default, against SCIP’s disabled-by-default nonconvex separators, and, on the constructed families, against Gurobi; an ablation measures what the certificate prevents.

The other ingredients, namely the support description, Bernstein and ball-arithmetic bounds, safe rounding, Lagrangian duality and oracle separation, are classical; we use them and state them precisely where a certificate depends on the details (Sections 5 and 6, Appendices E and F).

Section 8 reports five campaigns, each run with solver code, limits and protocol fixed in advance. All 274,489 cut records of the last three campaigns passed a fresh-process replay, which re-executes the exact support routines against the source model and the row stored by SCIP. Constants taken from the separator’s samples would have been wrong for 38% of the cuts on MINLPLib models and for almost all sampled path-family cuts, and their rows would have cut off feasible solutions; local minimization reduced the share of wrong constants to a few percent but still left rows that cut off known optima, and even a numerical global solver once reported optimality for a wrong constant. On MINLPLib models the cuts did not change which models SCIP solved in campaigns 2 to 4 and left SCIP’s own solving time unchanged; the Python separator made the runs slower, and SCIP’s disabled-by-default separators improved the root bound more often. On 20 larger models that SCIP does not solve within 60 seconds, the separator with default limits stopped during block discovery on 12, and with raised limits its cuts improved one root bound and worsened five. On constructed instances with the path structure of Section 3, SCIP’s root bound is weak because of a presolve step; with that step switched off, SCIP reaches the level of glued pair relaxations and no further. A post hoc diagnostic and a later prospective control showed that most of the remaining closure requires enough cuts per block, and that trying the exact support of each whole row first supplies the rest when the coupling row does not bind. On 20 fresh instances, this configuration, fixed before the instances were generated, closed SCIP’s root gap and solved all 20 within 300 seconds, against at most 10 for SCIP in three settings and 5 for Gurobi; with a binding coupling row the cuts solved all 20 as well, but so did SCIP without its presolve step, faster. On a family of constrained stars with up to 16 leaves, star blocks certified by the star oracle made the root bound exact, which neither smaller blocks nor SCIP achieved; Gurobi, which found good solutions faster, solved more of these instances.

We do not claim a new convexification principle, a general speedup, or a certificate for a complete solve; the certificates cover the cuts and the model they were derived from. Section 2 fixes notation and reviews related work, Sections 3–6 develop the theory, Section 7 describes the implementation, Section 8 reports the experiments, and Section 9 discusses consequences and open problems. The appendices contain proofs and examples, the relation of the cuts to the hull of the feasible set, whole-domain bounds for non-quadratic blocks, complete separation, the first two campaigns, and the test sets with detailed results.

2 Setting and related work

A *block* consists of variables $x \in \mathbb{R}^n$, a nonempty compact domain $D \subset \mathbb{R}^n$ and continuous features $F = (f_1, \dots, f_p) : D \rightarrow \mathbb{R}^p$. In a model, the features are expressions of the same variables, and D is the set on which they must be relaxed, typically a box intersected with the linear rows that involve

only these variables. The *joint graph* of F on D and its convex hull are

$$G_D(F) = \{(x, F(x)) : x \in D\}, \quad H_D(F) = \text{conv } G_D(F), \quad (1)$$

where the coordinates w_j of $(x, w) \in H_D(F)$ stand for the expressions $f_j(x)$. For a direction $(a, b) \in \mathbb{R}^n \times \mathbb{R}^p$ the (lower) *support value* is

$$h_D(a, b) = \min_{x \in D} \{a^\top x + b^\top F(x)\}. \quad (2)$$

Proposition 2.1 (Support description). *For every (a, b) and every $\beta \leq h_D(a, b)$, the inequality $a^\top x + b^\top w \geq \beta$ is valid for $H_D(F)$. Moreover $H_D(F)$ is the intersection of the half-spaces $a^\top x + b^\top w \geq h_D(a, b)$ over all (a, b) .*

Proof. The inequality holds on $G_D(F)$ by definition and is preserved by convex combinations. $G_D(F)$ is compact, so $H_D(F)$ is compact and convex, and a point outside it can be strictly separated by a linear functional; written as a lower bound, the separating inequality is one of the stated half-spaces. $\square$

We call $a^\top x + b^\top w \geq \beta$ a *support cut*, and a *support inequality* when $\beta = h_D(a, b)$. Support inequalities are as strong as possible for the block: no valid inequality in these coordinates is missing from the family. They also combine information that separate relaxations discard. If each f_j is relaxed on its own, the result is the intersection of the hulls of the graphs of (x, f_j) , which can be much larger than $H_D(F)$ when the f_j share variables (Example 5.4). For $n = p = 1$, the family recovers the convex and concave envelopes of a single function. If integer variables occur in a block, we drop their integrality and keep their bounds; support cuts remain valid for the mixed-integer graph, but they need not describe its convex hull.

2.1 Related work

Factorable and simultaneous relaxations. Global solvers relax nonconvex models by decomposing expressions into elementary operations and relaxing each with an auxiliary variable [8, 13, 90, 117, 119]. The resulting relaxations ignore the dependence among expressions that share variables, and they ignore linear constraints that link those variables [130]. Simultaneous convexification addresses the first point. Tawarmalani [118] develops inclusion certificates for the convex hull of several functions of the same variables; Ballerstein [5], as stated by Liers et al. [83, Proposition 1], characterizes the hull of a vector of functions by linear combinations, and Liers et al. [83] separate over it in a gas network application. He and Tawarmalani [64, 65, 66] build composite relaxations with vectors of outer functions, Zhu et al. [130] treat simultaneous factorable graphs with coupling constraints, and Davarnia et al. [36] convexify vectors of bilinear functions simultaneously over polytopes; He et al. [67] describe simultaneous hulls of vectors of univariate functions through moment hulls. For multilinear and bilinear terms, the gap between termwise and joint relaxations is well studied [7, 19, 86, 110]. Several bilinear and quadratic terms are relaxed jointly by the multiterm relaxations of Bao et al. [6] and by the edge-concave relaxations of Misener and Floudas [91]. Reduced-space relaxations avoid auxiliary variables altogether [20, 98, 123]. Our support cuts are an instance of simultaneous convexification; what we add concerns which blocks are worth relaxing jointly, which support problems can be solved exactly, how the cuts enter the original model, and how they are certified.

Lagrangian and aggregation cuts. Cuts in the original variables obtained by minimizing a weighted sum of constraint functions over a block are Lagrangian cuts. They appear in block decomposition methods for MINLP [73, 106, 128], in range reduction [55, 119] and in rigorous filtering [42]. The relaxation they describe, convexification in the joint space of variables and constraint values, is the primal form of the Lagrangian dual [46, 48, 53, 82], and the set version of this statement for Lagrangian cuts is given by Chen and Luedtke [29]. Aggregations of quadratic constraints that describe convex hulls under additional hypotheses [18, 40], surrogate relaxations that keep an aggregated constraint nonconvex [94], and the joint-range inequalities of Xu and Pokutta [129], which convexify the joint range of two quadratic functions and lift the result to sparse cuts, are different objects; Section 5.3 relates them to our cuts.

Exact low-dimensional and structured quadratic hulls. Anstreicher and Burer [3] show that the convex hull of $\{(x, xx^T)\}$ is described by the semidefinite constraints and the constraints of the reformulation-linearization technique (RLT) [113, 115] over a simplex of dimension at most three and over a box in dimension two, and that polytopes of dimension at most three admit an extended formulation built from a triangulation; for the box in dimension three these constraints are not enough [2, 3, 26]. Explicit envelopes through polyhedral subdivisions are given by Tawarmalani et al. [121]. Convex envelopes of bivariate functions over general polygons, with supporting hyperplanes, are computed by Locatelli and Schoen [84]. Minimizing a quadratic over a polytope of fixed dimension by enumerating faces is classical [97, Section 2.9], quadratic programming is in NP [124], and box-constrained quadratic programming on forests is strongly polynomial [37]. Explicit hulls for star-shaped sign patterns are given by Dey and Khajavirad [39] and Khajavirad [76], and parametric dynamic programs over trees solve convex quadratic problems with indicators [15]. The gluing of local relaxations along a sparsity pattern is discussed in Section 3.4.

Rigorous relaxations and safe cuts. Safe bounds and cuts for linear and mixed-integer programs correct floating-point results with variable bounds and directed rounding [34, 45, 102]; exact MIP solvers and VIPR certificates check complete solves against the original instance [31, 35, 44], SCIP 10 offers an exact mode for MILP [68], and Halbig et al. [62] compute optimality certificates for convex MINLPs. For nonlinear problems, rigorous linear relaxations and safe linear estimators are established [21, 41, 74, 105], as are verified Bernstein bounds [50, 51, 95] and ball arithmetic [70]. That rounding in the problem formulation itself must be controlled was pointed out by Neumaier [101, Section 20] and Schichl and Neumaier [112]. Liers et al. [83] mention safe rounding of coefficients without details. Cut generators are usually tested for validity against known feasible solutions [89]; replay of exact certificates is a stronger test for the cuts it covers. SCIP relaxes its own nonlinear cuts with variable bounds before storing them, in floating point, to protect them against the solver’s coefficient rounding [12, Section 4.2.10]. To our knowledge, no published MINLP implementation certifies the exported rows of joint support cuts against the source model and the row stored by the solver.

Separation and oracles. Weak separation and weak optimization are polynomially equivalent for convex bodies given by oracles [58, 59]. Gilbert’s nearest-point method gives an oracle procedure that returns either a nearby point of the set or a separating hyperplane [25, 54], and the local-cut and Fenchel-cut schemes separate by column generation over an optimization oracle, in exact arithmetic when needed [22, 33]; in direction space this column generation is Kelley’s cutting-plane method [30, 75]. Oracle-based cuts for polynomial optimization are studied by Bienstock et al. [17]. Appendix F states the corresponding finite procedure for joint graphs and its certificate.

Native SCIP. SCIP’s nonlinear constraint handler detects quadratic, bilinear, convex, concave and other structures, uses projections of linear rows for bilinear terms [93], and separates RLT cuts [14] and semidefinite minor cuts [13]. If a variable with finite bounds and no objective coefficient occurs in a single polynomial constraint that is concave (convex) in it, with an infinite left-hand (right-hand) side, some optimal solution has the variable at one of its bounds; SCIP’s presolve then makes the variable binary when its bounds are $[0, 1]$, or adds a bound disjunction [12, Section 4.2.7], an idea that goes back to Hansen et al. [63]. Several cut families that target the same joint quadratic structure as our cuts are implemented but disabled by default because they have not been found to pay off in general [13, 68]: edge-concave cuts on aggregated quadratic terms [91], intersection cuts for quadratic constraints from maximal quadratic-free sets [32, 96], interminor cuts, which are intersection cuts for vanishing 2×2 minors of the matrix of products [12, Section 4.11], related to the outer-product-free sets of Bienstock et al. [17], and RLT cuts for hidden products [14]. A new separator therefore has to be evaluated against SCIP’s default and against these separators, not against termwise McCormick relaxations; Section 8 does both.

3 When exact pair hulls do not compose

A solver that relaxes a quadratic expression term by term or pair by pair works with the variables and products that occur in the model. Even if every pair were relaxed exactly, gluing the pair hulls on their shared variables gives in general a weaker set than the hull of the joint graph, because finitely many shared moments do not determine the distribution of a shared variable [47, 80, 103]. In the MINLP literature, Tawarmalani [118, Example 3.8] shows that separate envelopes are weaker than the simultaneous hull because they may represent the same point by different convex combinations. This section quantifies the phenomenon for the smallest quadratic case, a path of three variables with exact pair hulls, and determines exactly when gluing on κ shared moments yields no positive bound. The results explain when a merged block is stronger than its pairs, and they delimit that advantage.

3.1 An example on a path

Consider variables $x, y, z \in [0, 1]$ that interact along the path x – y – z . Write m_u for the coordinate that represents u , s_u for u^2 and p_{uv} for uv . The exact pair hulls are

$$H_{xy} = \text{conv}\{(x, y, x^2, y^2, xy) : x, y \in [0, 1]\}, \quad H_{yz} = \text{conv}\{(y, z, y^2, z^2, yz) : y, z \in [0, 1]\},$$

each described by its semidefinite and RLT constraints [3]. Let R be the set of points

$$(m_x, m_y, m_z, s_x, s_y, s_z, p_{xy}, p_{yz})$$

whose (x, y) and (y, z) parts lie in H_{xy} and H_{yz} ; the two parts share m_y and s_y . We call R the *glued pair relaxation*. The joint hull $H = \text{conv}\{(x, y, z, x^2, y^2, z^2, xy, yz) : x, y, z \in [0, 1]\}$ is contained in R . For a quadratic Φ in (x, y, z) without the monomial xz , the *linearization* $\text{lin } \Phi$ is the affine function of these coordinates obtained by replacing each monomial by its coordinate. Its minimum over H is $\min_{[0,1]^3} \Phi$, and the difference between this minimum and the minimum over R is the *gap* of R in the direction of Φ .

Proposition 3.1 (A strict gap). *The point*

$$\bar{v} = (m_x, m_y, m_z, s_x, s_y, s_z, p_{xy}, p_{yz}) = \left(\frac{1}{2}, \frac{1}{2}, \frac{4}{5}, \frac{1}{2}, \frac{5}{16}, \frac{4}{5}, \frac{3}{8}, \frac{1}{2}\right) \quad (3)$$

belongs to R but not to H . The quadratic

$$\Phi(x, y, z) = (y - \frac{1}{4} - \frac{x}{2})^2 + (y - \frac{5z}{8})^2 + x(1-x) + z(1-z) \quad (4)$$

has minimum $1/128$ on $[0, 1]^3$, attained only at $(1, \frac{11}{16}, 1)$, while $\lim \Phi(\bar{v}) = 0$. The support inequality

$$2s_y - \frac{1}{2}m_y - p_{xy} - \frac{5}{4}p_{yz} + \frac{5}{4}m_x - \frac{3}{4}s_x + m_z - \frac{39}{64}s_z \geq -\frac{7}{128} \quad (5)$$

is valid for H and violated by $1/128$ at $\bar{v}$.

Proof. The measure with mass $\frac{1}{2}$ on each of $(x, y) = (0, \frac{1}{4})$ and $(1, \frac{3}{4})$, and the measure with mass $\frac{1}{5}$ on $(y, z) = (0, 0)$ and $\frac{4}{5}$ on $(\frac{5}{8}, 1)$, have the moments

$$(m_x, m_y, s_x, s_y, p_{xy}) = (\frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{5}{16}, \frac{3}{8}), \quad (m_y, m_z, s_y, s_z, p_{yz}) = (\frac{1}{2}, \frac{4}{5}, \frac{5}{16}, \frac{4}{5}, \frac{1}{2}).$$

They agree on (m_y, s_y) , so $\bar{v} \in R$. The two pair parts of Φ , $(y - \frac{1}{4} - \frac{x}{2})^2 + x(1-x)$ and $(y - \frac{5z}{8})^2 + z(1-z)$, vanish on the supports of the respective measures, so $\lim \Phi(\bar{v}) = 0$. The minimum is the case $A = \{\frac{1}{4}, \frac{3}{4}\}$, $C = \{0, \frac{5}{8}\}$ of Theorem 3.2 below, with $\delta = \frac{3}{4} - \frac{5}{8} = \frac{1}{8}$. Expanding $\Phi \geq \frac{1}{128}$ gives (5), since the constant term of Φ is $\frac{1}{16}$. $\square$

The inequality (5) uses only coordinates that occur in the model; it needs no product xz . It is the support inequality of the three-variable block in the direction of Φ , and both oracles of Section 4 compute it exactly.

3.2 An exact description for a path family

The example belongs to a family whose behavior can be described completely. Fix an interval $[\ell, u]$ for y and two-point sets $A = \{a_1, a_2\}$ and $C = \{c_1, c_2\}$ in $[\ell, u]$ with $a_1 \neq a_2$ and $c_1 \neq c_2$. For weights $\omega_A \geq (a_2 - a_1)^2$ and $\omega_C \geq (c_2 - c_1)^2$ let

$$\Phi_A(x, y) = (y - a_1 - (a_2 - a_1)x)^2 + \omega_A x(1-x), \quad \Phi_C(y, z) = (y - c_1 - (c_2 - c_1)z)^2 + \omega_C z(1-z), \quad (6)$$

with $x, z \in [0, 1]$, and $\Phi = \Phi_A + \Phi_C$. Define R as above, on the box $[0, 1] \times [\ell, u] \times [0, 1]$, and let $\delta = \min\{|s - t| : s \in A, t \in C\}$. The sets A and C *interleave* if they are disjoint and their elements alternate in sorted order.

Theorem 3.2 (Interleaving). (i) *The minimum of Φ over the box is $\delta^2/2$.*

(ii) *The minimum of $\lim \Phi$ over R is 0 if $A \cap C \neq \emptyset$ or A and C interleave, and $\delta^2/2$ otherwise.*

Hence the gap of R in the direction of Φ is either 0 or $\delta^2/2$; it is positive exactly when A and C interleave, and it then equals $\delta^2/2$.

Proof. Since $\omega_A \geq (a_2 - a_1)^2$, the function $\Phi_A(\cdot, y)$ is concave or affine, so its minimum over $[0, 1]$ is attained at $x \in \{0, 1\}$, where it equals $(y - a_1)^2$ or $(y - a_2)^2$. Hence $\min_{x \in [0, 1]} \Phi_A(x, y) = \text{dist}(y, A)^2$, and similarly for Φ_C .

(i) For fixed y the minimum over x and z is $\min_{s \in A, t \in C} \{(y - s)^2 + (y - t)^2\}$, and $(y - s)^2 + (y - t)^2 = 2(y - \frac{s+t}{2})^2 + \frac{1}{2}(s - t)^2$ with $\frac{s+t}{2} \in [\ell, u]$.

(ii) A point of H_{xy} is the moment vector of a finitely supported probability measure π on $[0, 1] \times [\ell, u]$, and the linearization of Φ_A at that point is $\int \Phi_A d\pi \geq \int \text{dist}(y, A)^2 d\nu$, where ν is the y -marginal of π ; equality holds when π puts its mass at a minimizing $x \in \{0, 1\}$ for each y . Hence the minimum over R is the minimum of $\int \text{dist}(y, A)^2 d\nu_A + \int \text{dist}(y, C)^2 d\nu_C$ over probability measures ν_A, ν_C on $[\ell, u]$ with equal first and second moments. It is zero if and only if some measure




Figure 1: Moment pairs $(\mathbb{E}y, \mathbb{E}y^2)$ of distributions on a two-point set form the chord of the parabola between the two points. Solid: A ; dashed: C . (a) The sets $A = \{\frac{1}{4}, \frac{3}{4}\}$ and $C = \{0, \frac{5}{8}\}$ of Proposition 3.1 interleave, the chords cross at $(\frac{1}{2}, \frac{5}{16})$, and the glued pair relaxation has value 0. (b) Nested sets $A = \{0, \frac{3}{4}\}$ and $C = \{\frac{1}{4}, \frac{1}{2}\}$: the chords do not meet, and by Theorem 3.2 the glued relaxation attains the true minimum $\delta^2/2$ in the direction of Φ .

on A and some measure on C have equal first two moments, that is, if the chords of the parabola $\{(t, t^2)\}$ spanned by A and by C meet. Two such chords meet if they share an endpoint; otherwise, since no three points of a parabola are collinear and (t, t^2) lies below the chord spanned by A exactly when t lies between a_1 and a_2 , they meet if and only if A and C interleave. When the chords do not meet, A and C are either separated (one lies below the other) or nested (one lies between the two points of the other). In both cases Appendix A.1 gives a polynomial p of degree at most two such that $\text{dist}(y, A)^2 - p(y) \geq \eta_A$ and $\text{dist}(y, C)^2 + p(y) \geq \eta_C$ for all y , with $\eta_A + \eta_C = \delta^2/2$. Integrating against ν_A and ν_C , whose moments of order at most two agree, shows that the minimum over R is at least $\delta^2/2$, and graph points show that it is at most $\delta^2/2$. $\square$

Proposition 3.1 is the case $[\ell, u] = [0, 1]$, $A = \{\frac{1}{4}, \frac{3}{4}\}$, $C = \{0, \frac{5}{8}\}$ and $\omega_A = \omega_C = 1$, with sorted order $0, \frac{1}{4}, \frac{5}{8}, \frac{3}{4}$. Figure 1 shows what goes wrong. In R the two pairs may represent y by different distributions that agree only in their first two moments. Each pair places y on its own preferred two-point set, and whenever the chords cross, the two distributions have the same mean and second moment. The theorem concerns the directions Φ of the family (6); in other directions on the path the gap can take other values, and Proposition 3.5(ii) expresses it for every direction.

3.3 Matching more moments

The same mechanism defeats any interface that identifies a fixed finite number of moments of the shared variable. Let X_L and X_R be compact sets of points (x, y) and (y, z) , with y scalar, and let H_L and H_R be the convex hulls of continuous feature maps on them whose first κ coordinates are $y, y^2, \dots, y^\kappa$; the other local features are arbitrary. Let R_κ be the set of pairs of points of H_L and H_R that agree in these κ coordinates. Let $\Phi_L \geq 0$ on X_L and $\Phi_R \geq 0$ on X_R be affine functions of the local features, with zero sets Z_L and Z_R , and let A and C be the projections of Z_L and Z_R onto y . For disjoint compact $A, C \subset \mathbb{R}$, the *alternation number* $\text{alt}(A, C)$ is the largest r for which there are points $t_0 < \dots < t_r$ of $A \cup C$ with consecutive points in different sets, and $\text{alt}(A, C) = 0$ if $A \cup C$ has at most one point.

Theorem 3.3 (Alternation). *Suppose A and C are disjoint, so that $\Phi_L + \Phi_R > 0$ at every point (x, y, z) with $(x, y) \in X_L$ and $(y, z) \in X_R$. The linearization of $\Phi_L + \Phi_R$ vanishes at some point of R_κ if and only if $\text{alt}(A, C) \geq \kappa + 1$. In particular, gluing on κ moments always gives a positive*

bound when $|A| + |C| \leq \kappa + 1$, and, for finite sets with $|A| + |C| = \kappa + 2$, gives the bound zero exactly when A and C alternate.

Proof. As in the proof of Theorem 3.2, the linearization vanishes at a point of R_κ if and only if there are probability measures on A and on C with equal moments of orders $1, \dots, \kappa$; these measures lift to Z_L and Z_R atom by atom. If $\text{alt}(A, C) = r \leq \kappa$, there is a polynomial of degree r that is positive on A and negative on C (Appendix A.2), and its integrals against two such measures would have to be equal, a contradiction. If $\text{alt}(A, C) \geq \kappa + 1$, take an alternating chain $t_0 < \dots < t_{\kappa+1}$ and the weights $w_i = \prod_{j \neq i} (t_i - t_j)^{-1}$ of the divided difference of order $\kappa + 1$. They annihilate all polynomials of degree at most κ , and their signs alternate, so their positive and negative parts, normalized, are probability measures on the A -points and on the C -points with equal moments of orders $0, \dots, \kappa$. $\square$

For finite A and C the criterion is the classical description of Radon partitions of points on the moment curve: the convex hulls of $\{(t, \dots, t^\kappa) : t \in A\}$ and $\{(t, \dots, t^\kappa) : t \in C\}$ meet exactly when A and C alternate at least $\kappa + 1$ times [24]. By Gordan's alternative it is equivalent to the fact that the least degree of a polynomial that is positive on A and negative on C is $\text{alt}(A, C)$, which rests on the sign-change properties of Chebyshev systems [71]. Theorem 3.2(ii) is the case $\kappa = 2$. What the theorem adds is its reading as a criterion for gluing interfaces.

Theorem 3.3 characterizes when the glued bound is zero, not when it is exact. For $\kappa \geq 3$ the glued bound can lie strictly between 0 and the minimum. For $\Phi_L = (y - 2x)^2 + 4x(1 - x)$ and $\Phi_R = (y - 1 - 2z)^2 + 4z(1 - z)$ with $x, z \in [0, 1]$ and $y \in [0, 3]$ we have $A = \{0, 2\}$, $C = \{1, 3\}$ and $\text{alt}(A, C) = 3$, so gluing on $\kappa = 3$ moments gives a positive bound. The probability measure ν_A with masses $\frac{57}{112}$ and $\frac{55}{112}$ at $\frac{15}{32}$ and $\frac{57}{32}$, and the probability measure ν_C with masses $\frac{22}{117}$, $\frac{1045}{1456}$ and $\frac{95}{1008}$ at 0 , $\frac{39}{32}$ and $\frac{81}{32}$, have equal moments of orders up to three, and $\int \text{dist}(y, A)^2 d\nu_A + \int \text{dist}(y, C)^2 d\nu_C = 6103/16128 \approx 0.378$, whereas $\min(\Phi_L + \Phi_R) = 1/2$. The dichotomy of Theorem 3.2 is therefore specific to two shared moments and two-point sets. For every κ , however, some quadratic configuration has a full gap.

Proposition 3.4 (A full gap for every κ). *Let $\kappa \geq 1$ and choose $r \geq 1$ with $\kappa \leq 2^{r+1} - 2$. On $x, z \in [0, 1]^r$ and $y \in [0, 2^{r+1} - 1]$ let*

$$\Phi_L = \left(y - \sum_{j=1}^r 2^j x_j\right)^2 + \sum_{j=1}^r 4^j x_j(1 - x_j), \quad \Phi_R = \left(y - 1 - \sum_{j=1}^r 2^j z_j\right)^2 + \sum_{j=1}^r 4^j z_j(1 - z_j).$$

Let H_L be the convex hull of the features $y, \dots, y^{\max\{\kappa, 2\}}, x_j, yx_j$ and $x_i x_j$ ($i < j$) over $[0, 1]^r \times [0, 2^{r+1} - 1]$, and H_R the analogous hull in (y, z) . Then $\min(\Phi_L + \Phi_R) = 1/2$, while the linearization of $\Phi_L + \Phi_R$ vanishes at a point of R_κ .

Proof. The coefficient of x_j^2 in Φ_L is $4^j - 4^j = 0$, so Φ_L is affine in the features and multilinear in x . Its minimum over x is therefore attained at a vertex, where it equals $(y - \sum_{j \in J} 2^j)^2$ for a subset J ; hence $\min_x \Phi_L = \text{dist}(y, A)^2$ with $A = \{0, 2, \dots, 2^{r+1} - 2\}$, and similarly $\min_z \Phi_R = \text{dist}(y, C)^2$ with $C = \{1, 3, \dots, 2^{r+1} - 1\}$. By the argument of Theorem 3.2(i), with $\delta = 1$, the minimum is $1/2$. The points $0, 1, \dots, \kappa + 1$ alternate between A and C , so Theorem 3.3 applies; explicitly, the binomial weights $\binom{\kappa+1}{i} 2^{-\kappa}$ on the even and on the odd integers $i \leq \kappa + 1$ have equal moments of orders up to κ . $\square$

Thus no fixed number of shared moments makes the gluing exact, even for quadratic blocks, and the relative gap can be 100%. As $\kappa \rightarrow \infty$ the glued bounds do converge to the true minimum, because measures on a compact interval are determined by their moments and measures with equal marginals glue [80, Lemma 6.3]; the convergence is not uniform over instances.

3.4 Relation to known results

Several earlier examples show that sparse relaxations on a path can be strictly weaker than dense ones. Nie and Demmel [103, Example 3.5] give a numerical quartic example, and Nie et al. [104, Example 6.7] a box-constrained quartic for which the sparse moment hierarchy is never tight. Finite gluing results require rank or flatness conditions on the overlaps [80, Theorem 3.7], [47]. For sparse quadratically constrained problems, Kojima et al. [78] assemble exactness of the semidefinite relaxation from exact clique-wise subproblems; they assume that maximal cliques meet in at most one node, because shared off-diagonal entries impose product conditions that local exactness cannot control [78, Section 3.3], and the same block-clique structure makes doubly nonnegative relaxations exact [77]. In our path, the homogenized cliques $\{1, x, y\}$ and $\{1, y, z\}$ share the off-diagonal entry m_y as well as s_y , and Proposition 3.1 shows that this obstruction occurs even when the local relaxations are exact hulls. For quadratic problems with sign-oriented squares, Dey and Khajavirad [39, Corollary 1] decompose the convex hull across a complete separator that carries no positive square term (no “plus loop”), and Khajavirad [76] builds on this decomposition. Proposition 3.1 shows that the hypothesis cannot be dropped even for a path of three variables with a single plus loop at the center: the reduced objective $2y^2 - \frac{1}{2}y - xy - \frac{5}{4}yz + \frac{1}{16} + \frac{1}{2}x + \frac{25}{64}z$, which equals Φ at binary leaves, has the same witness and the same gap $1/128$. Their stable-set theorem [39, Theorem 2 and Corollary 4] still gives an exact extended formulation for this sign pattern, because the plus-loop node y is a stable set, but it adds the products among the neighbors of y , here the nonedge product xz ; this matches the next paragraph. To our knowledge, a degree-two example with exact pair hulls, an exact rational gap and a separating cut in the existing coordinates, the value $\delta^2/2$ of the gap in Theorem 3.2, and the use of the alternation criterion to decide when a κ -moment interface gives no positive bound have not been given before.

On a three-variable path, the advantage of the joint block over its pairs does not extend to dense relaxations. Sparse and dense Shor relaxations [115] coincide for chordal sparsity patterns, because a partial positive semidefinite matrix with a chordal pattern has a positive semidefinite completion [57, 81, 127]. Hence R is the projection of the dense moment relaxation of $(1, x, y, z)$ with the RLT inequalities [113, 114] of the edges and squares, and what closes the gap must constrain the nonedge product p_{xz} , for example through its McCormick inequalities. On the whole family (6), with $\xi_A = y - a_1 - (a_2 - a_1)x$ and $\xi_C = y - c_1 - (c_2 - c_1)z$,

$$\Phi - \frac{\delta^2}{2} = \frac{(\xi_A + \xi_C)^2}{2} + \sum_{i,j \in \{0,1\}} F_{ij} \ell_{ij}(x, z) + \left(\omega_A - \frac{(a_2 - a_1)^2}{2}\right)x(1-x) + \left(\omega_C - \frac{(c_2 - c_1)^2}{2}\right)z(1-z), \quad (7)$$

where $\ell_{00} = (1-x)(1-z)$, $\ell_{10} = x(1-z)$, $\ell_{01} = (1-x)z$, $\ell_{11} = xz$, and $F_{ij} = \frac{1}{2}((c_{j+1} - a_{i+1})^2 - \delta^2) \geq 0$. The linearized ℓ_{ij} are the McCormick expressions of the product xz , the linearized $(\xi_A + \xi_C)^2$ is a quadratic form in the dense moment matrix of $(1, x, y, z)$, and the last two terms are nonnegative on the pair hulls. Hence the dense moment matrix together with the McCormick inequalities of the nonedge product xz proves $\Phi \geq \delta^2/2$ for every member of the family, and neither ingredient alone suffices (Appendix A.3). This is a special case of a theorem of Burer et al. [27, Theorem 1]: for box-constrained quadratic programs in at most three variables whose off-diagonal coefficients are nonpositive, the Shor relaxation with the RLT upper bounds is exact. On a path $x-y-z$ over a box, replacing x by its complement or z by its complement where necessary makes the coefficients of xy and yz nonpositive, the coefficient of xz stays zero, and complementation maps McCormick inequalities to McCormick inequalities. The dense moment relaxation with the McCormick inequalities of all products, including xz , is therefore exact for every quadratic objective on a three-variable path over a box, and (7) is an explicit certificate for the family. The support inequality (5) attains the

same bound without the coordinate p_{xz} and without a semidefinite constraint. That matters in practice, because solvers relax the products present in the model and rarely add new ones; but on three variables the advantage is one of representation, not of strength over such relaxations.

The restriction to three variables matters. From four variables on, the dense relaxation with the McCormick inequalities of all products is no longer exact for every objective, and joint blocks can be strictly stronger. The objective of Burer et al. [27, Example 4] is $x^\top Qx + c^\top x$ with the tridiagonal matrix Q that has diagonal $(8, 25, 25, 8)$ and off-diagonal entries $-14, -25, -14$, and with $c = (12, 29, 0, 0)$; its interaction graph is the path $x_1-x_2-x_3-x_4$. Its minimum over $[0, 1]^4$ is 0 (Theorem 4.1), whereas the dense relaxation contains a rational point with objective value $-109/1024$ [27, Section 6.4]; numerically its value is about -0.122 , and that of the glued pair hulls of the three edges about -1.454 . The instances of Proposition 3.4 behave even worse: for $r = 2$ and $r = 3$ (five and seven variables) the dense relaxation contains rational points with objective values below 10^{-6} , which we verified exactly, whereas the minimum is $1/2$; for $r = 1$ (three variables) it attains $1/2$, as the theorem above predicts. In both cases the support inequality of the joint block in the direction of the objective uses only the coordinates of the model and attains the minimum. For blocks of four or more variables, joint support cuts can therefore be much stronger than dense first-level relaxations, not only more economical.

3.5 When gluing is exact, and the gain from merging

Gluing is exact when the shared coordinates determine the shared distribution. Under a running-intersection decomposition, local measures with equal full marginals on the separators glue to a global measure [80, 126], and Tawarmalani [118, Corollary 3.10] glues two functions that share a variable when their inclusion certificates have equal marginals of that variable. If the separator takes finitely many known values $t_1, \dots, t_N$, its moments of orders $0, \dots, N - 1$ determine its distribution through a Vandermonde system; for a binary separator the mean suffices. This is why the McCormick inequalities describe the convex hull of a pure bilinear function on a forest over the unit box [107, Proposition 8]: the extreme points are binary, the McCormick inequalities of an edge are the nonnegativity of the four joint probabilities of its endpoints, and adjacent edges share a binary variable whose distribution is fixed by its mean. Square terms break this argument, because they let a fractional mean be realized by continuous distributions, as in (3).

The gain from merging pair directions into one star direction can be computed exactly, and it is related to the glued relaxation by a re-splitting of the center terms. Part (ii) of the next proposition is the Lagrangian-decomposition duality [53, 60] applied to the pair hulls, with the shared moments (m_y, s_y) as copied variables; in the language of sparse moment relaxations, it says that the glued bound is the best split of the objective with an interface polynomial of degree two in y [56], [104, Theorem 3.1]. What we add is that exact pair and star minima evaluate both sides.

Proposition 3.5 (Gain from merging). *Let $P_i \subset \mathbb{R}^2$ be rational polytopes in the variables (y, x_i) , $i = 1, \dots, k$, whose star domain $\Omega = \{(y, x) : (y, x_i) \in P_i \text{ for all } i\}$ is nonempty, and let $q_i(y, x_i)$ be rational quadratics. Put $\beta_i = \min_{P_i} q_i$, $\beta^* = \min_{\Omega} \sum_i q_i$ and $\Delta = \beta^* - \sum_i \beta_i$.*

- (i) $\Delta \geq 0$ is the largest constant c such that $\sum_i q_i \geq \sum_i \beta_i + c$ on Ω , and $\Delta > 0$ if and only if the projections onto y of the sets $\arg \min_{P_i} q_i$ have empty common intersection.
- (ii) Let R_Ω be the set of k -tuples of points of the pair hulls $\text{conv}\{(y, x_i, y^2, yx_i, x_i^2) : (y, x_i) \in P_i\}$ that agree in (m_y, s_y) . Then $\beta^* - \min_{R_\Omega} \sum_i \text{lin } q_i = \inf \Delta'$, where the infimum is over all re-splittings $q'_i = q_i + \alpha_i y + \gamma_i y^2$ with $\sum_i \alpha_i = \sum_i \gamma_i = 0$, and Δ' is computed as Δ for the q'_i .

Proof. (i) Clearly $\sum_i q_i \geq \sum_i \beta_i$ on Ω , and β^* is the best constant in that direction. If some $\bar{y}$ lies in every projection, choose x_i with $(\bar{y}, x_i) \in \arg \min_{P_i} q_i$; the assembled point lies in Ω and attains $\sum_i \beta_i$. Conversely, if $\Delta = 0$, a minimizer over Ω minimizes every q_i , so its y -coordinate lies in every projection.

(ii) A re-splitting does not change $\sum_i q_i$, so β^* is the same for all of them. The set R_Ω is the intersection of the compact convex product of the pair hulls with the linear equations $m_y^{(i)} = m_y^{(1)}$ and $s_y^{(i)} = s_y^{(1)}$ for $i \geq 2$, and it is nonempty because Ω is. Attach multipliers to these equations. Their Lagrangian is bilinear, and by the minimax theorem for a compact convex set [116] the minimum of $\sum_i \ln q_i$ over R_Ω equals the supremum over the multipliers of the minimum of the Lagrangian over the product of the pair hulls. A choice of multipliers is exactly a re-splitting with $\sum_i \alpha_i = \sum_i \gamma_i = 0$, and the minimum of the Lagrangian then separates into $\sum_i \min_{P_i} q'_i = \sum_i \beta'_i$. Hence $\min_{R_\Omega} \sum_i \ln q_i = \sup \sum_i \beta'_i$, which is the claim. $\square$

Both β_i and β^* are computed exactly by the algorithms of Section 4; in Proposition 3.1 both pair minima are zero and $\Delta = 1/128$. Since each β'_i is a minimum of functions that are affine in (α_i, γ_i) , part (ii) turns the gap of the glued relaxation in any direction into the maximization of a concave function of $2(k-1)$ parameters whose oracle is the pair minimum. The quantity Δ compares the merged cut with separately added pair cuts for one decomposition, and the gap to the glued relaxation is its infimum over re-splittings; Appendix A.4 gives two examples, one in which a re-splitting removes the whole gain and one in which pairwise intersections of the projections in (i) do not suffice. The proposition is a diagnostic for a proposed merge; it does not say which directions to merge or whether merging pays off in a solver.

4 Exact support for quadratic blocks

If every feature f_j is a quadratic polynomial, then every support problem (2) asks for the minimum of a single quadratic function over the block domain. This section gives two exact rational algorithms for that problem. The first works for any bounded polytope of small dimension; the second works for quadratic stars of any size whose constraints couple the center with one leaf at a time. Both return an attained minimizer and a rational minimum, so they serve as exact support oracles and give support inequalities that are tight.

4.1 Bounded polytopes in fixed dimension

Let $P = \{x \in \mathbb{R}^d : Ax \leq b\}$ be a bounded polytope given by m rational rows, which include finite variable bounds, and let $q(x) = \frac{1}{2}x^\top Hx + c^\top x + c_0$ with rational data and $H = H^\top$. The matrix H may be indefinite or singular, P may be lower-dimensional, and rows may be redundant or implied equalities; an equality is given by two opposite rows. For a set I of rows, A_I and b_I denote the corresponding submatrix and subvector, and

$$M_I = \begin{pmatrix} H & A_I^\top \\ A_I & 0 \end{pmatrix}, \quad M_I \begin{pmatrix} x \\ \mu \end{pmatrix} = \begin{pmatrix} -c \\ b_I \end{pmatrix}. \quad (8)$$

The *candidates* are the points x obtained from all sets I with $|I| \leq d$ for which M_I is nonsingular and $Ax \leq b$ holds. The multipliers μ are not sign-restricted: the system expresses stationarity on the affine hull of a face, not a KKT condition.

Theorem 4.1 (Exact quadratic support on a polytope). *If $P \neq \emptyset$, some candidate minimizes q over P ; hence the least candidate value is $\min_P q$, and it is attained at a rational point. If there is*

no candidate, then $P = \emptyset$. Both statements remain true if only sets I of linearly independent rows are enumerated.

The proof (Appendix B) chooses a global minimizer in the relative interior of a face of smallest dimension and shows that H is positive definite on the face’s direction space, so that a nonsingular system (8) of independent active rows has this minimizer as its solution. The sign of the multipliers and the inertia of H need not be tested, because extra candidates, such as saddle points of faces, are feasible and cannot produce a value below the minimum. Dependent rows never give a nonsingular M_I , so duplicate, redundant and opposite rows are harmless. Two parts of the argument cannot be dropped. Boundedness is needed: for $P = \{x \in \mathbb{R} : x \geq 0\}$ and $q = -x$ the enumeration returns 0, and for $P = \mathbb{R}$ and $q = x$ it returns no candidate. The choice of a minimizer on a face of minimal dimension is also needed: for $q = (x - y)^2$ on $[0, 1]^2$, the minimizers (t, t) with $0 < t < 1$ give only singular systems, and the enumeration finds the minimizers $(0, 0)$ and $(1, 1)$ instead.

Corollary 4.2 (Support of a quadratic feature graph). *If $f_1, \dots, f_p$ are quadratic polynomials with rational coefficients and (a, b) is rational, the support value $h_P(a, b)$ is the rational number obtained by applying Theorem 4.1 to $q = a^\top x + b^\top F$, and the support inequality is tight at the returned minimizer.*

There are at most $N(m, d) = \sum_{j=0}^d \binom{m}{j}$ sets I . For each, Gaussian elimination on a system of order at most $2d$ and the feasibility test take $O(d^3 + md)$ arithmetic operations, so the enumeration uses $O(N(m, d)(d^3 + md))$ operations. With fraction-free elimination, every intermediate number is a minor of an integer matrix of order at most $2d$, and Hadamard’s inequality bounds the bit length of every candidate and of the minimum by $O(d(\langle q \rangle + \langle A \rangle + \log d))$, where $\langle q \rangle$ is the encoding length of q and $\langle A \rangle$ the largest encoding length of a row (Appendix B). For fixed d the algorithm is therefore strongly polynomial, as is known for face enumeration [38, 124], and the minimum has polynomial encoding length. The method is total enumeration of faces [97, Section 2.9]; what we add is a statement that covers every degenerate case a certificate meets (redundant, opposite and implied-equality rows, empty and lower-dimensional polytopes, singular faces), together with explicit size bounds.

Without a bound on the dimension the problem is hard. The following facts are folklore [26, Section 2.3]; Appendix B gives the short proofs.

Proposition 4.3 (Hardness in unrestricted dimension). *Computing $\min_{x \in [0, 1]^n} q(x)$ for rational quadratics q is strongly NP-hard, even when q is multilinear with integer coefficients of absolute value at most n . Deciding whether a given inequality $q(x) \geq t$ holds on $[0, 1]^n$ is coNP-complete.*

Unless $\text{NP} = \text{coNP}$, the validity of an exact support cut for a quadratic block of unrestricted size therefore has no short certificate that can be checked in polynomial time, and a certificate that is replayed by repeating the enumeration, as ours is, is natural for this class. Useful exact oracles must restrict the dimension, as in Theorem 4.1, or the structure, as in the next subsection. Restricting the interaction graph alone is not enough: box-constrained quadratic minimization is strongly NP-hard already at treewidth two [37, Theorem 3], and quartic minimization already on paths [37, Theorem 2].

4.2 Quadratic stars with center–leaf constraints

The merged blocks of Section 3 are stars: several pairs that share one variable. For stars there is an exact algorithm whose cost is nearly linear in the size of the block. Let y be the center and

$x_1, \dots, x_k$ the leaves, and consider

$$q(y, x) = c_0 + c_1 y + c_2 y^2 + \sum_{i=1}^k (\alpha_i x_i^2 + (\beta_i y + \gamma_i) x_i) \quad (9)$$

with rational coefficients of arbitrary signs. The domain Ω is given by finite rational bounds on all variables and by rational affine rows, each of which involves y and *at most one* leaf; an equality is given as two rows. Every scalarization $a^\top(y, x) + b^\top F(y, x)$ of a vector F of quadratics with this interaction pattern has the form (9). Rows that involve two or more leaves are excluded, because with them the problem is strongly NP-hard already without a center and with two leaves per row: for a graph G on $n \geq 3$ vertices, the minimum of $\sum_i (n x_i (1 - x_i) - x_i)$ over $0 \leq x \leq 1$ and $x_i + x_j \leq 1$ for the edges ij of G is minus the stability number of G , because the objective is concave and the vertices of this polytope are half-integral [100], and at a vertex with $h \geq 1$ half-integral coordinates the value exceeds that of a stable set by $h(n/4 - 1/2) > 0$. A single row over all leaves gives subset sum [72]; compare Moré and Vavasis [92].

For fixed y , the rows of leaf i restrict x_i to an interval $[L_i(y), U_i(y)]$, where L_i is the maximum of the affine lower bounds and U_i the minimum of the affine upper bounds that these rows and the leaf bounds define. Thus L_i is convex and U_i concave, both piecewise affine. The set $J_i = \{y : L_i(y) \leq U_i(y)\}$ is an interval because $L_i - U_i$ is convex. Given y , the leaves are constrained independently, so the projection of Ω onto y is the interval I obtained by intersecting the bounds of y , the rows in y alone, and all J_i . If I is empty, so is Ω .

Theorem 4.4 (Exact support on constrained stars). *Let m be the number of rows that involve a leaf and m_0 the number of rows in y alone, and suppose Ω is nonempty. Then $\min_\Omega q$ is attained at a rational point, and it can be computed, together with a minimizer and a partition of I into at most $2m + 4k + 1$ intervals on each of which every leaf has an affine minimizer, with $O(m_0 + (m + k) \log(m + k))$ arithmetic operations and comparisons. All numbers involved have bit length polynomial in the input length; the breakpoints have bit length bounded independently of k .*

Proof. Define $\varphi_i(y) = \min\{\alpha_i x_i^2 + (\beta_i y + \gamma_i) x_i : L_i(y) \leq x_i \leq U_i(y)\}$ for $y \in I$, so that $\min_\Omega q = \min_{y \in I} V(y)$ with $V(y) = c_0 + c_1 y + c_2 y^2 + \sum_i \varphi_i(y)$. Let $s_i = m_i + 2$ be the number of affine functions that define L_i and U_i , where m_i is the number of rows of leaf i ; together L_i and U_i have at most $s_i - 2$ breakpoints.

If $\alpha_i > 0$, the minimizer is $\xi_i = \max\{L_i, \min\{v_i, U_i\}\}$ with the affine function $v_i(y) = -(\beta_i y + \gamma_i)/(2\alpha_i)$. The function $v_i - L_i$ is concave and $v_i - U_i$ convex, so each has at most two isolated zeros, and ξ_i is affine between consecutive points of a set of at most $s_i + 2$ breakpoints. If $\alpha_i \leq 0$, a minimizer is an endpoint of the interval, and

$$\psi_i(U_i) - \psi_i(L_i) = (U_i - L_i)(\alpha_i(U_i + L_i) + \beta_i y + \gamma_i), \quad \psi_i(t) = \alpha_i t^2 + (\beta_i y + \gamma_i)t. \quad (10)$$

The first factor is nonnegative on I . On each of the at most $s_i - 1$ intervals where L_i and U_i are affine, the second factor is affine and has at most one zero, which decides the preferred endpoint; this gives at most $2s_i - 3$ breakpoints. Taking all breakpoints together splits I into at most $\sum_i 2s_i + 1 = 2m + 4k + 1$ intervals on which every leaf rule is affine and $V(y) = q(y, \xi(y))$ is a quadratic polynomial in y .

Each φ_i is continuous on I , since it is the minimum of a continuous function over an interval whose endpoints depend continuously on y . Hence V agrees with the quadratic of each interval on its closure, and the minimum of V is the least value of these quadratics at the interval endpoints and, where the leading coefficient is positive, at interior stationary points. All breakpoints and all

endpoints of I are zeros of affine functions with rational coefficients, so every candidate is rational, and the minimizing y together with $x_i = \xi_i(y)$ is a rational minimizer.

For the operation count, intersecting the bounds of y with the m_0 rows in y alone costs $O(m_0)$. The upper envelope of s_i lines is computed by sorting slopes and one stack-based scan in $O(s_i \log s_i)$ operations, and J_i , the breakpoints and the rules of leaf i in $O(s_i)$ more. Sorting all breakpoints costs $O((m+k) \log(m+k))$. A sweep over the sorted breakpoints maintains the three coefficients of V ; at each breakpoint only the leaves whose rule changes are updated, so the sweep takes $O(m+k)$ operations. Bit lengths are bounded in Appendix C. $\square$

The bound on the number of intervals is tight up to a constant factor: a single concave leaf whose bounds zigzag forces $2s_i - 3$ rule changes, and k hinge leaves force k more, so the partition can need $\Theta(m+k)$ intervals. On instances with $m = \Theta(k)$, the operation bound is optimal up to a constant factor for algebraic computation trees: by Ben-Or's theorem, applied to a family of instances with $3r$ leaves and r rows whose decision set has $r!$ connected components, deciding whether the minimum is at least a given value requires $\Omega((m+k) \log(m+k))$ operations (Appendix C).

Relation to forest algorithms. Del Pia and Khajavirad [37] solve box-constrained quadratic programs on forests with $O(n^2)$ arithmetic operations. After an affine scaling of each variable to $[0, 1]$, a box-constrained star is such a forest, and their dynamic program rooted at the center specializes to the box-constrained case of Theorem 4.4; in that case our contribution is only the $O(k \log k)$ count, which replaces their $O(k^2)$ merge at the root. Rows that couple the center with a leaf are outside their model. A shear $x_i \mapsto x_i - \rho y$ turns a leaf domain that is a parallelogram with two vertical sides back into a box and keeps the star structure, but triangles, trapezoids with nonparallel sides and polygons with more sides cannot be reduced in this way. Such rows also break the structural property on which their algorithm rests: the conditional value of a leaf is then no longer a minimum of affine functions of y over a fixed set, and it can have convex kinks. For example, for $x \in [0, 1]$ with rows $x \geq y - \frac{1}{2}$ and $x \geq \frac{1}{2} - y$, the minimum of x is $|y - \frac{1}{2}|$. The second-order cone descriptions of star hulls by Dey and Khajavirad [39] assume sign patterns of the square terms and box domains, and the LP approximations of Bienstock and Muñoz [16] allow constraints but are approximate. Our algorithm conditions on the center, as nonserial dynamic programming does [11], and solves a one-parameter quadratic problem for each leaf, a special case of multiparametric quadratic programming [9]; Bhathena et al. [15] use parametric dynamic programs of this kind on trees for convex quadratic problems with indicators. To our knowledge, no exact algorithm for nonconvex stars with center-leaf rows has been published.

Remark 4.5 (Deeper trees). The argument uses that every vertex other than the center is a leaf. On the path x_1 - x_2 - x_3 over $[0, 1]^3$ with $f = (x_2 - \frac{1}{8})x_1 + x_2^2 - x_2 + x_2x_3 + x_3^2 - x_3$, the conditional value of the subtree $\{x_1, x_2\}$ given $x_3 = t$ is $\min\{-\frac{1}{8}, -\frac{1}{4}(1-t)^2\}$, which switches at the irrational point $t = 1 - \sqrt{2}/2$, although $\min f = -\frac{3}{8}$ is rational (Appendix C). Rooted at x_2 , the same path is a star, and Theorem 4.4 gives a rational partition; the obstruction concerns nesting conditional values along a tree of depth two. Del Pia and Khajavirad handle such breakpoints for boxes with an explicit algebraic representation; we do not know whether edge-local affine rows can be added to deeper trees with a polynomial algorithm.

The certificate of a star support value records the partition of I , the affine rule of every leaf on every piece, the quadratic of V on every piece and the candidates; a checker rebuilds the partition from the rows and coefficients and compares. The star oracle in our implementation is simpler than the algorithm of the proof: it splits the center interval at all pairwise intersections of each leaf's bound lines and, inside each resulting interval, at the points where the unconstrained minimizer of a

convex leaf meets a bound or a concave leaf changes its preferred endpoint, and it then recomputes every leaf rule on every piece. This costs $O((m+k)^3)$ operations, stores $\Theta(k)$ rules per piece, and can produce more pieces than Theorem 4.4 requires. It certified the star cuts of our first implementation, whose stars had at most four leaves (Appendix G), and those of the star blocks with up to 16 leaves in Section 8.6; in the other experiments every block had at most four variables, and the polytope oracle certified every quadratic cut. Section 8.6 also compares the two star algorithms on random stars with up to 10^4 leaves.

5 Cuts in the original variables

Proposition 2.1 produces inequalities in lifted coordinates (x, w) . Using them requires variables w_j with $w_j = f_j(x)$ in the model. If the solver's representation does not already contain exactly these variables, adding them changes the model: new equalities, new variables for branching and propagation, and possibly a different presolve. Our first implementation used such a reformulation and needed a control mode to separate its effect from that of the cuts (Appendix G). This section shows how to avoid the auxiliary variables. The nonlinear functions are minimized jointly over the block domain, and the rows of the model then eliminate them. The underlying facts are classical Lagrangian duality; we state them in the form that the implementation and its certificates need.

5.1 Aggregating nonlinear rows

Let $v \in \mathbb{R}^N$ be the vector of all model variables, including an objective variable if there is one, and let $x = Sv$ be a block of n of them, selected by a 0/1 matrix S . Suppose the model contains rows

$$g(Sv) + Av \leq b \tag{11}$$

with $g : D \rightarrow \mathbb{R}^m$ and affine parts Av that may involve any variables, including block variables; an equality is written as two opposite rows. Here D is a compact set that contains the block part Sv of every point to which the cuts must apply, for example the block box intersected with the affine rows that involve only block variables.

Proposition 5.1 (Aggregated support cut). *Let $a \in \mathbb{R}^n$, $\lambda \in \mathbb{R}_{\geq 0}^m$ and*

$$\beta \leq \varphi(a, \lambda) := \inf_{u \in D} \{a^\top u + \lambda^\top g(u)\}. \tag{12}$$

Then every point that satisfies (11) and $Sv \in D$ satisfies

$$(S^\top a - A^\top \lambda)^\top v \geq \beta - \lambda^\top b. \tag{13}$$

Proof. At such a point, $\lambda \geq 0$ gives $\lambda^\top g(Sv) \leq \lambda^\top (b - Av)$. Substitute this into $a^\top Sv + \lambda^\top g(Sv) \geq \beta$. $\square$

This is weak Lagrangian duality [23, Sections 5.1.3 and 5.2.2]. It needs neither convexity nor optimal multipliers. The same inequality is the Lagrangian cut of a block in the extended block-separable reformulation of Nowak [106, Section 7.1], after the block's auxiliary variables are eliminated with the rows (11); related cuts and range reductions appear in Tawarmalani and Sahinidis [119], Karuppiah and Grossmann [73] and Domes and Neumaier [42], and cuts in the original space obtained by projecting lifted quadratic relaxations are studied by Saxena et al. [111]. Its strength comes from minimizing the nonlinear functions jointly over their common domain before

they are eliminated. Nonnegative combinations of linear rows that are already in the relaxation cannot cut off any point of it.

An objective is handled through its epigraph. For minimization of $g_0(x) + \ell_0(v) + c_0$ with objective variable t , the row $g_0(x) + \ell_0(v) - t \leq -c_0$ is one of the rows (11); for maximization the hypograph row $-g_0(x) - \ell_0(v) + t \leq c_0$ is used. The two multipliers of the rows of an equality may be combined into one free multiplier. Validity uses only $Sv \in D$ and the selected rows. Integrality, the other constraints and the objective play no role, so the cut is valid for the mixed-integer feasible set and for its closed convex hull. If D is built from bounds valid only at a node of the search tree, or from an incumbent cutoff, the cut is valid only there; our implementation uses only original bounds and rows, so its cuts are global.

The choice of D matters: affine rows that involve only block variables belong in D , not only in the LP. On the triangle $\{(x, y) : x, y \geq 0, x + y \leq 1\}$ every point satisfies $x^2 + 2xy + y^2 \leq x + y$ and $xy \leq \frac{1}{4}$, because $s = x + y \in [0, 1]$ gives $s^2 \leq s$ and $4xy \leq s^2$. In the coordinates $(m_x, m_y, s_x, s_y, p_{xy})$ these are linear inequalities valid for the hull of the quadratic graph over the triangle. The point $(\frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2})$, the average of the graph points at $(0, 0)$ and $(1, 1)$ of the box $[0, 1]^2$, lies in the hull over the box and satisfies $m_x + m_y \leq 1$, but violates both inequalities. Intersecting the hull over the box with the row is therefore strictly weaker than convexifying over the row-constrained domain, as is known for factorable relaxations [2, 128, 130]. Here the first inequality is the sum of the RLT products $(1 - x - y)x \geq 0$ and $(1 - x - y)y \geq 0$, so the example shows the weakness of the box hull, not an advantage over RLT or semidefinite relaxations; for a single bilinear term, SCIP's projection-based envelopes [93] exploit linear rows in the same spirit.

5.2 What the cuts can express

Let $K = \text{conv}\{(u, g(u)) : u \in D\}$ and $Q = \{0\} \times \mathbb{R}_{\geq 0}^m$. The affine map $T(v) = (Sv, b - Av)$ places a model point in the lifted space of the block. The next theorem identifies the set described by all cuts (13). It is the set version of the classical fact that the Lagrangian dual bound of a nonconvex problem equals the optimal value of its convexification in the joint space of variables and constraint values [46, 48, 53, 82]; the separation argument is the one that Chen and Luedtke [29, Theorem 3] use for Lagrangian cuts in two-stage integer programs. We include the short proof because our setting has nonlinear rows and affine parts on arbitrary variables, and because no constraint qualification is needed.

Theorem 5.2 (Closure of the aggregated cuts). *Let $D \subset \mathbb{R}^n$ be nonempty and compact and let g be continuous on D . A point v satisfies every inequality (13) with $\beta = \varphi(a, \lambda)$, over all $a \in \mathbb{R}^n$ and $\lambda \in \mathbb{R}_{\geq 0}^m$, if and only if*

$$v \in \mathcal{C} := T^{-1}(K + Q) = \{v : \exists y \text{ with } (Sv, y) \in K \text{ and } y \leq b - Av\}. \quad (14)$$

Proof. The set K is compact and convex, and Q is a closed convex cone, so $U = K + Q$ is closed and convex. A linear functional (a, λ) has finite infimum on U only if $\lambda \geq 0$, and then the infimum is $\varphi(a, \lambda)$. Evaluating the valid inequality $a^\top u + \lambda^\top y \geq \varphi(a, \lambda)$ of U at $T(v)$ gives exactly (13), so every point of $\mathcal{C}$ satisfies all cuts. Conversely, if $T(v) \notin U$, a linear functional strictly separates $T(v)$ from U ; its infimum on U is finite, so $\lambda \geq 0$, and the corresponding cut is violated at v . $\square$

If auxiliary variables $w = g(x)$ were added, the joint-graph relaxation of the block would be $(x, w) \in K$, $w + Av \leq b$, and its projection onto v is $\mathcal{C}$. The direct cuts can therefore express exactly what this lifted relaxation expresses about the original variables, without adding variables or equalities to the model, and no more. Relaxations that keep the rows inside the block, such as

the convex-hull relaxations of block-decomposition methods [106, Section 3], [128], can be strictly stronger (Example D.2). There is also a cost: with the same linearization points, outer approximation of a lifted formulation can be tighter than outer approximation in the original space [120], and a finite set of direct cuts may need more members than a lifted description. For one row, $\mathcal{C}$ is the familiar envelope relaxation. Write $\text{vex } g$ and $\text{cav } g$ for the convex and concave envelopes of a scalar function g over D , so that $\text{vex } g(x) = \min\{y : (x, y) \in K\}$.

Corollary 5.3 (One row). *If $m = 1$, then $\mathcal{C} = \{v : Sv \in \text{conv } D, \text{vex } g(Sv) + Av \leq b\}$. For an equality $g(Sv) + Av = b$, written as two rows, $\mathcal{C}$ is the set of v with $Sv \in \text{conv } D$ and $\text{vex } g(Sv) \leq b - Av \leq \text{cav } g(Sv)$.*

Proof. The projection of K onto x is $\text{conv } D$, and for each $x \in \text{conv } D$ the fiber $\{y : (x, y) \in K\}$ is the interval $[\text{vex } g(x), \text{cav } g(x)]$. A suitable y in (14) exists if and only if $\text{vex } g(Sv) \leq b - Av$. For the equality, the lifted point of the two rows is $(y, -y)$ with y in the same fiber, and the two conditions $y \leq b - Av$ and $-y \leq -(b - Av)$ force $y = b - Av$. $\square$

Joint treatment can only help when rows share block variables or when D couples them: if D is a product $D^1 \times D^2$ and the rows split into two groups whose nonlinear functions depend only on x^1 and only on x^2 , then $\mathcal{C}$ is the intersection of the closures of the two groups, because the convex hull of a product of sets is the product of their convex hulls. When the rows do share variables, the gain can be strict.

Example 5.4 (Two rows sharing one variable). Let $D = [-1, 1]$ and consider the rows $x^2 - v_1 \leq 0$ and $-x^3 + v_2 \leq 0$. The convex envelope of x^2 is x^2 , and the convex envelope of $-x^3$ on $[-1, 1]$ takes the value $-1/4$ at $x = 0$ (on $[-1/2, 1]$ it is the line through $(1, -1)$ that is tangent to $-x^3$ at $x = -1/2$). By Corollary 5.3, the intersection of the two one-row closures contains $(x, v_1, v_2) = (0, 0, 1/4)$. The aggregated cut with $a = 0$ and $\lambda = (1, 1)$ uses $\min_{u \in [-1, 1]}(u^2 - u^3) = 0$ and reads $v_1 - v_2 \geq 0$, which this point violates. In terms of (14), the only probability measure on $[-1, 1]$ with mean 0 and second moment 0 is the point mass at 0, whose third moment 0 is smaller than the required $1/4$.

That joint relaxations dominate separate envelopes is well known [106, Lemma 3.4], [83, 118]; the example makes the effect concrete for the cut family used here.

5.3 Relation to the hull of the feasible set

The closure $\mathcal{C}$ is a relaxation of the set $\Sigma = \{v : Sv \in D, v \text{ satisfies (11)}\}$, so $\text{conv } \Sigma \subseteq \mathcal{C}$. Equality holds when every row has its own free direction in the affine parts, as in the setting of Liers et al. [83] (Proposition D.1 in Appendix D). Without such free directions the gap can be strict, and it is then not a failure of the direction search. For $D = [0, 1]$ and the equality row $x^2 = 1/4$, whose feasible set is $\{1/2\}$, Corollary 5.3 gives $\mathcal{C} = \{x : x^2 \leq 1/4 \leq x\} = [1/4, 1/2]$: the point $x = 1/4$ is the mean of the measure with mass $3/4$ at 0 and $1/4$ at 1, which satisfies the row in expectation, so even the complete family of cuts accepts it. This is a Lagrangian duality gap [53, 106]. It is closed by convexifying aggregated sets instead of aggregated functions, as in the aggregation closure of Dey et al. [40], or by placing the row inside D , which changes the support problem; our implementation keeps nonlinear rows out of D so that the support problems stay in the classes of Section 4. Appendix D places $\mathcal{C}$ in a hierarchy with the surrogate relaxations of Müller et al. [94] and shows that both inclusions can be strict.

6 Certified cuts and their replay

A separation routine has two parts. One part chooses a direction; the other proves a lower bound for that direction. Only the second affects validity. This section states what must be checked for a cut in the original variables to be valid as the solver stores it, and how a record of the check is replayed.

6.1 The stored row

Floating-point LP solvers store binary64 coefficients. Every finite binary64 number is a dyadic rational, so the row that the solver uses is a rational inequality, and validity can be stated for that row. For a cut of the form (13), four things must hold.

- (C1) The direction (a, λ) is taken as the exact rational value of its binary64 representation, and a lower bound $L \leq \varphi(a, \lambda)$ over the whole domain D is certified.
- (C2) The cut is assembled exactly: $c = S^T a - A^T \lambda$ and $r = L - \lambda^T b$, with repeated occurrences of a variable combined.
- (C3) The rational row $c^T v \geq r$ is rounded to binary64 with a correction that keeps it valid (Proposition 6.1).
- (C4) The row that the solver stores is the exported row.

By Proposition 5.1, (C1) and (C2) give a valid rational row, and (C3) and (C4) carry its validity to the stored row. The direction itself may come from an LP over samples, from a heuristic or from a previous iteration; none of these needs to be correct. The principle is established: a numerically chosen affine bound needs no verification of how it was found, only a rigorous final shift [51, Section 5], [50]; cut parameters may be chosen in floating point and the derivation then repeated rigorously [102, p. 294]; and safe linear estimators are defined in terms of the machine coefficients themselves [21]. The hazard at the last step is known to SCIP’s developers. Outside its exact mode, SCIP 10 rounds a row coefficient that is integral within its tolerance to that integer and drops coefficients below its epsilon, without adjusting the sides (functions `rowAddCoef`, `rowChgCoefPos` and `rowMerge` in `src/scip/lp.c` of 122). Its nonlinear constraint handler therefore relaxes its own cuts with variable bounds before storing them, in floating point, and rounds without correction when a needed bound is infinite [12, Section 4.2.10]; cuts added by another separator receive no such treatment. Two points are specific to support cuts: the bound must be certified for the final direction on the whole block domain, which is a global nonconvex minimization, and the row that SCIP stores must be compared with the certified row.

Proposition 6.1 (Safe export). *Let $c^T v \geq r$ be valid, let $\hat{c}$ be binary64 coefficients and $\Delta = \hat{c} - c$, and let $\ell_j \leq v_j \leq u_j$ be valid bounds, where ℓ_j is finite whenever $\Delta_j > 0$ and u_j is finite whenever $\Delta_j < 0$. Put $E = \sum_{j:\Delta_j>0} \Delta_j \ell_j + \sum_{j:\Delta_j<0} \Delta_j u_j$. Then $\hat{c}^T v \geq \hat{r}$ is valid for every $\hat{r} \leq r + E$.*

Proof. $\hat{c}^T v = c^T v + \Delta^T v \geq r + E$. □

This is the bound-corrected rounding of safe LP bounds and safe cuts [34, 45, 102], applied to the eliminated row. A coefficient that is exactly representable needs no bound, even on an unbounded variable such as the objective variable. The correction must be applied after every transformation that changes the row, including merging, scaling and the removal of tiny coefficients; $\hat{r}$ is rounded downward and compared with $r + E$ in exact arithmetic. If a needed bound is infinite or the correction destroys the violation, the row is discarded; a fixed safety shift does not replace the

calculation (Section 8.2 gives an example). The bounds must be global implications of the model (Section 7.1). Our exporter is more conservative than the proposition: it requires both bounds of a variable to be finite whenever the coefficient of that variable changes. For support inequalities in lifted coordinates (x, w) , as in our first implementation, bounds on w need not be available, and the direct form is used instead: if $L \leq \min_{x \in D} \{\hat{a}^\top x + \hat{b}^\top F(x)\}$ for the exact values of the stored coefficients $(\hat{a}, \hat{b})$ and $\hat{r} \leq L$, then $\hat{a}^\top x + \hat{b}^\top w \geq \hat{r}$ is valid for $H_D(F)$ by Proposition 2.1.

6.2 Lower bounds on the whole domain

For quadratic blocks, (C1) is computed exactly by the oracles of Section 4. For polynomial *remainders* (the nonlinear part of a row, Section 7.2) in one or two variables we use Bernstein enclosures on a subdivision of the block box, computed in exact rational arithmetic [50, 52, 95]; for univariate elementary expressions, outward-rounded ball arithmetic [70] on interval cells, strengthened by a chord correction from a verified bound on the second derivative (Appendix E). A certificate depends on two hypotheses. The cells must cover the box exactly, starting from the trusted box and not from a list supplied by the generator, and a bound proved on a smaller domain is never reused on a larger one. Domain requirements, such as a positive argument of log or a nonzero denominator, are proved on every cell; they are collected from the raw source expression before any symbolic simplification (Section 7.1), and every original feature is enclosed on the whole cell before features are combined, so that cancellation between features cannot hide a requirement. The chord correction is used only on cells where the expression is twice continuously differentiable and a finite bound on its second derivative is verified. A bound computation returns one of four outcomes: a certified bound with its partition; *empty*, when every cell is excluded by a domain row, so that $D = \emptyset$; *unsupported*, for an operation outside the grammar; or *incomplete*, when a work limit is reached. Only the first produces a cut. An incomplete computation says nothing about the point being separated; in particular, it is not evidence that the point lies in the hull.

6.3 Records and replay

Each cut is accepted through the same sequence: certify the support bound for the binary64 direction (C1); eliminate the nonlinear rows exactly and merge coefficients (C2); apply the correction of Proposition 6.1 (C3); and read back the row that SCIP has created from the cut, with its columns, coefficients, constant, sides and scope, before it is passed to SCIP’s separation storage, and compare it with the certified row after mapping original to transformed variables (C4). Any difference, such as a coefficient rounded by SCIP or a variable aggregated or removed by presolve, rejects the cut. Accepted rows are passed to SCIP as global, removable cuts with `forcecut` set. SCIP’s later handling of the row, such as converting a cut in a single variable into a bound change, is outside this check.

The record of a cut contains the serialized source model, the signed source rows and multipliers, the proven bounds with their derivations, the support domain, the direction, the support certificate (the complete candidate enumeration, the star partition or the subdivision with its bounds), the eliminated row, the rounding correction and the stored SCIP row. Replay runs in a fresh process. It verifies the hashes of the archived sources and model, re-executes the model builder (which constructs, but does not solve, a SCIP model through PySCIPOpt) and the bound derivations, rebuilds every signed source side from the archived model with its own affine splitter, repeats the exact support calculation for the recorded binary64 direction on the rebuilt features and domain, recomputes elimination and rounding, and compares the result with the stored row as recorded by the separator; the comparison with SCIP’s row itself is made only when the cut is generated.




Figure 2: Data flow of the separator. Shaded steps are part of the certificate and are repeated by replay; dashed steps only choose what to certify and need not be correct. SCIP’s own computations are outside the certificate.

For one recorded cut in each part and cut mode, fourteen single-field corruptions check that replay rejects altered inputs: the model hash, the support normal, the right-hand side and identity of a source side, a feature, the box, a row right-hand side, the scope, the rounding correction, the support identity, a stored coefficient, a stored side, the column mapping and the bound proof.

Replay re-executes the generator’s exact code: the model builder, the bound replay, the support kernels and the elimination and export routine. It does not run the direction search, the LP solver or SCIP’s solving process, and it has its own affine splitter. It is therefore a re-execution in a fresh process, not an independent checker. The trusted base consists of Python’s rational arithmetic; SymPy’s expression construction (including its automatic canonicalization), expansion, polynomial conversion and differentiation; Arb through python-flint; the OSiL reader, both its parsing and its conversion of decimal tokens to binary64; PySCIPOpt’s translation of the checked expression objects and row sides into SCIP constraints, and its accessors for stored rows; and the code shared by generation and replay. None of these is formally verified. The certificates say nothing about SCIP’s presolve, its LP bounds, its branching or its final optimality claim; SCIP’s exact mode, which solves MILPs without tolerances and can write VIPR certificates of complete solves for independent checking [31, 44], is limited to MILP [68].

7 Implementation in SCIP

We implemented the cuts as a separator for SCIP 10.0.2 [68], written in Python on top of PySCIPOpt 6.2.1 [87], with exact arithmetic in Python fractions and SymPy, ball arithmetic in Arb through python-flint, and HiGHS for the direction LPs. The separator adds global cuts at the root node. It does not change SCIP’s constraint handlers, presolve, propagation or branching, and every native nonlinear constraint stays in the model. Figure 2 shows the data flow. This section describes the parts on which validity depends and the choices that shape the computational results; Appendix H lists the work limits.

7.1 A faithful model

A cut certificate refers to the model that the solver solves. Two ordinary steps of model loading, assembling expressions and simplifying them, can change that model without notice. We guard against both in every mode, including the native baseline, so that all modes solve the same model. We also describe the bounds that certificates use.

Arithmetic. Every finite binary64 number in the source file is interpreted as its exact rational value. Assembling a native expression can still change values: in $(0.1x)^2$ the source coefficient is the exact square of the binary64 number 0.1, whereas a coefficient computed in floating point is that square rounded again, and a constraint constructor may move a constant to the other side with a rounded subtraction. After assembly, the builder reads back the PySCIPOpt expression object that it passes to SCIP and compares its exact value with the source expression by symbolic subtraction and expansion, with no tolerance. If they differ, it rebuilds the expression as an explicit expression graph with separate constant nodes and compares again; a remaining difference rejects the model. Constants are kept inside the expression, so the constructor cannot round them into the row sides, which are passed unchanged. This is translation validation [99, 108] of the model import. It addresses the observation of Neumaier [101, Section 20] that modeling systems introduce uncontrolled rounding, and the requirement of Schichl and Neumaier [112, Section 3.1] that simplification of an expression graph introduce no roundoff. Its boundary is the expression object; PySCIPOpt’s translation of it into SCIP’s internal expression, and SCIP’s later simplification, presolve and floating-point LP, are outside it.

Domains. Simplification can silently enlarge a domain: $0 \cdot \log x$, x/x , $(\sqrt{x})^2$ and $(\log x)^0$ are defined only for $x > 0$, $x \neq 0$, $x \geq 0$ and $x > 0$, respectively. The builder visits every node of the source expression before simplification and proves each domain requirement from bounds and from affine rows; for example, $x - y = 1$ proves that $\log(x - y)$ is defined, with no bounds on x and y . Requirements that cannot be proved are enforced by existential constraints in a new variable u :

$$d \neq 0 \iff \exists u : du = 1, \quad d > 0 \iff \exists u \geq 0 : du = 1, \quad d \geq 0 \text{ is imposed directly.}$$

The first is the device of Rabinowitsch [109]. Over the reals these constraints describe strict domains exactly, without a positive tolerance such as the one used in common benchmark setups [13]; SCIP then enforces them, like every constraint, with its feasibility tolerance. A variable power $b(x)^{e(x)}$ is admitted as $\exp(e(x) \log b(x))$ only when $b(x) > 0$ is proved [compare 8]; otherwise the model is rejected, because imposing $b > 0$ could remove points at which the power is defined for other reasons.

Bounds. Support domains need finite bounds. Original affine rows imply bounds that may be missing from the declared box: for a row $\sum_i a_i x_i \leq b$ with $a_j \neq 0$, $a_j x_j \leq b - \sum_{i \neq j} \min\{a_i \ell_i, a_i u_i\}$. Any finite sequence of such steps, with integer rounding for integer variables, yields bounds that hold at every feasible point. Each step is recorded with its source row, and replay repeats the sequence. The bounds are used only as premises of certificates; the model keeps its declared bounds. If the propagation derives a contradiction in exact arithmetic, the model is refused: under the exact reading of its binary64 data it has no feasible point. This happened for 8 of the 322 models of the pool of Part 4D, all with a known MINLPLib solution: on two, an affine row is violated by about 10^{-16} at every point of the box; on six, the rounding of bounds of integer variables produces a row that no point satisfies.

7.2 Blocks, directions and activation

Discovery runs at the first separator call, so a model solved during presolve pays nothing. A ranged row $\ell \leq g(v) \leq u$ has two *sides*, $g(v) \leq u$ and $-g(v) \leq -\ell$. Discovery writes each finite side of each original row, including the objective epigraph row, exactly as a nonlinear *remainder* plus an affine part. A side is admissible if its remainder involves one to four variables with finite proven

bounds and is either a polynomial of total degree at most eight or a nonpolynomial expression in one variable. Blocks are the variable sets of single sides and greedy unions of overlapping sides, with at most four variables, six nonlinear sides and sixteen signed affine domain sides; at most 32 blocks are kept, the largest first and ties broken by the indices of their variables. Unsupported remainders are declined as a whole, and a remainder is never altered. Affine rows that involve only block variables, together with the proven bounds, form the support domain D . The certification kernels apply to the whole block: exact enumeration (Theorem 4.1) when the combined function is quadratic and the number of row subsets is at most 2,000, the star oracle (the simpler implementation described at the end of Section 4.2) when the enumeration budget is exceeded and the block is a star, Bernstein bounds for polynomials in a block of one or two variables (at most 1,024 tensor coefficients), and ball arithmetic in a block of one variable. With at most sixteen affine sides the enumeration budget is never exceeded in dimension three or less; for four variables it allows at most seven signed affine sides. A direction whose combined function falls outside these cases yields no cut.

At an LP solution the separator tries, for each block and each of its source sides j , the *remainder direction* $a = 0$, $\lambda = e_j$. It bounds only the side’s nonlinear remainder; the side’s affine terms in the block variables enter the cut only through the elimination. Two variants try other directions first. The *whole-row direction* of side j has $\lambda = e_j$ and a equal to the side’s coefficients of the block variables, so that its support value is the minimum of the whole side over D . The *block direction* has $\lambda_j = 1$ for every source side of the block and a equal to the sum of their coefficients of the block variables; its support value is the minimum of the sum of the block’s rows. After these directions the separator solves up to three LPs that separate the current point, after coordinate scaling, from the convex hull of sampled graph points plus the cone of the side coordinates; an LP direction is used only if its scaled violation exceeds 10^{-5} . Samples are grid points of the box (33 points in one variable, 7×7 in two, 3^d in $d = 3, 4$), plus, for quadratic blocks, the exact vertices of the domain polytope and their centroid when the enumeration budget allows. Each direction is tried at most once per block and run. The binary64 values of a and λ are then certified. For a non-quadratic block, the subdivision is refined only until the bound makes the row violated by the threshold at the current point; if it does not, the call counts as a failed certification. When an exact support computation returns a minimizer that the samples missed, the minimizer is added and the LP is solved again.

Two modes add cuts. Mode **all** tries the admitted blocks in their fixed order until a cap is reached. Mode **auto** admits only quadratic blocks that contain a side not proved convex and that have either an affine domain row that is not a bound or at least two distinct nonlinear rows; it requires a larger violation and stops trying a block after two failed certifications. Both modes run in at most three root callbacks, cap the number of cuts and support calls (*default limits*, Appendix H), and give all separator work a cumulative allowance of $\min\{1s, 0.05T\}$ for a total budget T , checked between operations; a single symbolic operation can overrun it. A cut is added only if its violation divided by $\max\{1, \|\hat{c}\|_1\}$ exceeds 10^{-5} in **all** and $5 \cdot 10^{-4}$ in **auto**. The separator runs before SCIP’s own separators in each root round, and its rows are added with **forcecut**, so they bypass SCIP’s cut selection. Mode **baseline** uses the same model builder and adds nothing. The experiments also use *raised limits* (128 blocks, 200 cuts) on MINLPLib models and, on the constructed families of Section 8.5, limits proportional to the number n of copies: *base limits* ($4n$ cuts, n per callback) and *wide limits* ($16n$ cuts, $4n$ per callback); all of these allow ten root callbacks, and the larger cut caps of Section 8.5 forty (Table 7). For the star experiment of Section 8.6, discovery can also form *star blocks*: a center variable together with up to sixteen variables that share two-variable sides with it, all these sides, and the affine rows among these variables. Star blocks precede the other blocks; for blocks of more than four variables the separator tries only the block direction and the whole-row directions, without sampling, and the star oracle certifies them.

8 Computational study

The experiments address four questions: whether the stored cuts are valid in practice, and what certification prevents; whether the cuts change SCIP’s results on benchmark models, including models that SCIP finds hard; what the separator does and what it costs; and how much of the root gap the cuts close on instances with the structure of Sections 3 and 4, compared with SCIP’s own alternatives and with Gurobi. We report five campaigns. Each was run with solver code, limits and protocol fixed before its first run; we call these results *prospective*, and we label analyses designed after results were seen as *post hoc*. Campaigns 1 and 2 used earlier implementations and are summarized in Appendix G. Campaign 3 tests the final implementation on MINLPLib models (Parts 3A and 3B) and on a constructed family (Part 3C); a post hoc diagnostic of its results is Part 3P. Campaign 4 was designed after campaign 3 to answer the questions that its results raised: it adds comparators, fresh instances, larger models, a variant of the family in which the block minima do not give the optimum, and an ablation of the certificate; the job builder and the reference computation of Part 4C4 were written during campaign 4, as its protocol specifies, before any run of that part. Campaign 5 was designed after campaign 4: it runs the star algorithm in the solver, adds the remaining comparators on the fresh instances, tests larger cut limits and compares the certificate with a numerical global solve. Table 1 lists the parts.

8.1 Setup

All runs use SCIP 10.0.2 through PySCIPOpt 6.2.1 and Python 3.13, with one SCIP thread and one BLAS thread per process, on a Linux host (WSL2) with 18 cores and 36 hardware threads that was shared with unrelated jobs. The one-minute load average was 14 to 21 during the prospective runs of campaign 3 (2 to 21 during Part 3P), 2 to 9 during campaign 4 (5th to 95th percentile; at most 15, in Part 4C4) and 1 to 20 during campaign 5, apart from a spike to 126 during Part 5S (Section 8.6). Within a part, at most six runs were active at a time, and the modes of one instance ran back to back in rotated order; the Part 3P runs on the path family overlapped the last 28 minutes of Part 3A, when up to ten runs were active. We compare times only within these paired runs of one part. Part 4C2 has no baseline of its own and is compared with the campaign-3 runs of the same instances; root bounds of identical runs repeated exactly across campaigns (all 100 baseline root bounds rerun in Part 3P), whereas solved counts and times are not paired across campaigns, and we say so where it matters. Timings are descriptive. Times are in-process wall times charged to the time budget (model build, discovery, separation and solve), excluding interpreter start-up and the primal check; *SCIP time excluding the callback* is SCIP’s solving time minus the time spent in the separator callback. Full runs have a budget of 300s; root runs have a node limit of one and a budget of 60s (Parts 3A, 3B, 4B2) or 120s (path family, Parts 4D and 5S). In Parts 3A and 3B a seed is SCIP’s random seed shift; constructed instances are identified by their generator seed and run with SCIP’s seed 0.

We write *SCIP* for SCIP with default settings, *SCIP-nolocks* for SCIP with its implicit-discreteness presolve step switched off (`constraints/nonlinear/checkvarlocks = d`), and *SCIP-extra* for SCIP with its disabled-by-default nonconvex separators switched on (Section 2.1); all three use our model builder and add no certified cut. Gurobi 13.0.3 [61] solves the same model with one thread, its nonconvex mode (`NonConvex=2`) and the same gap limit. The separator modes are listed in Table 7. Models come from MINLPLib [28, 125] in OSiL format, the XML instance format of Optimization Services (Appendix H), stratified by the library’s metadata into convex models, nonconvex models without integer variables, and nonconvex models with integer variables.

A model counts as *solved* if SCIP (or Gurobi) reports optimality or the relative gap limit 10^{-4} ,

Table 1: Parts of campaigns 3 to 5. Full runs have a budget of 300s; root runs have a node limit of one. All parts except 3P are prospective. Modes are defined in Appendix H (Table 7).

Part	Instances	Modes	Runs
3A	30 MINLPLib models of campaign 2; full runs seeds 0–2, root runs seed 0	SCIP, all, auto; root runs also all-diag	full, root
3B	30 structure-selected models; full runs seeds 0–1	SCIP, all, auto; root runs also all-diag	full, root
3C	20 path-family instances	SCIP, all-diag-mech	full, root
3P	3A, 3B (root runs); 3C (post hoc)	whole-row variant of the cut modes; 3C also wide limits	root; full, root
4C2	the 20 instances of 3C	SCIP-nolocks, SCIP-extra, frozen-wide, Gurobi	full, root
4C3	20 fresh path-family instances	SCIP, all-diag-mech, rowdir-wide, SCIP-extra, Gurobi	full, root
4C4	the instances of 4C3 with a binding row	SCIP, frozen-wide, rowdir-wide, SCIP-extra, Gurobi	full, root
4B2	the 30 models of 3B	no aggregation: SCIP, all, all-diag, all-diag-rowdir; SCIP-extra	root
4D	20 larger models that SCIP finds hard	full: SCIP, all, auto, SCIP-extra; root: SCIP, all, all-diag, all-diag-rowdir, SCIP-extra	full, root
4U	recorded cuts of campaigns 3 and 4	uncertified constants (samples, local solver)	offline
4S	random stars with 10 to 10^4 leaves	two star algorithms	offline
5S	30 star-family instances	SCIP, SCIP-nolocks, SCIP-extra, rowdir-star4, agg-star4, agg-star, Gurobi	full, root
5C-a	the instances of 4C3 and 4C4	SCIP-nolocks	full, root
5C-b	the instances of 4C4	larger cut caps (32n, 64n), both direction rules	root
5U	the cuts of 4U	numerical global solve (Gurobi)	offline

and the returned incumbent passes an independent check of all original rows, bounds, domains, integrality restrictions and the objective at a scaled tolerance of 10^{-5} ; the checker evaluates the source expressions without the model builder. Times are compared by shifted geometric means (shift 1 s) and by the median of per-run ratios, over runs solved in every compared mode. Dual bounds z and z' are compared with the tolerance $10^{-4} \max\{1, |z|, |z'|\}$, the gap limit. The *root gap closed* by a mode is $(z_{\text{mode}} - z_{\text{base}})/(z^* - z_{\text{base}})$ for minimization, where z denotes root dual bounds and z^* the optimal or best known value; for a run that ends at the root without reporting a root bound, because the root node was pruned, z is its final dual bound. Every recorded cut is replayed in a fresh process (Section 6.3). A run whose worker failed has no complete cut log; its cut count is reported as unknown, not as zero. The protocols also prespecify statistics that we do not report here, such as means per seed, bound comparisons at the tolerance 10^{-6} and scan counts per stratum; they are in the supplementary archive and do not change the conclusions.

8.2 Validity, and what certification prevents

Every recorded cut of campaigns 3 to 5 passed replay: 7,498 cuts in the prospective parts of campaign 3, 30,998 in Part 3P, 104,771 in campaign 4 and 131,222 in campaign 5, 274,489 records in total. These records contain 92,222 distinct exported rows: in the constructed families the separator runs only at the root, so the root and full runs of an instance and mode record the same cuts. In every part and cut mode, each of the fourteen corrupted records of Section 6.3 was rejected. Every incumbent returned in campaigns 3 to 5 passed the original-model check, with one exception in a run without certified cuts: in a Part 4D root run of SCIP-extra, the incumbent satisfied all constraints, but SCIP’s reported objective value differed by 34% from the value recomputed on the

Table 2: What a constant without a certificate would have done. For the recorded binary64 direction of each cut with an exact certificate, U1 is the minimum over the separator’s sample set, U2 the better of U1 and SLSQP local minimizations, and U3 the best bound of a numerical global solve by Gurobi with default tolerances. “Wrong” counts constants that exceed the exact support value by more than $10^{-6} \max\{1, |\text{value}|\}$; “cuts off” counts the resulting rows that cut off an incumbent of the same model that passed the original-model check, and “optimum” those that cut off the known optimal solution, which is known only for the path family. MINLPLib: all cuts of Parts 3A, 3B, 3P, 4B2 and 4D with an exact certificate. Path family: a random sample of 1,000 of the 138,000 cuts recorded in campaigns 3 and 4 (995 distinct).

Cuts	Constant	Wrong (models)	Cuts off (models)	Optimum
MINLPLib, 5,116 cuts, 47 models	U1	1,969 (12)	405 (7)	
	U2	230 (4)	132 (2)	
	U3	4 (1)	0	
Path family, 1,000 cuts, 60 instances	U1	976 (60)	506 (58)	404
	U2	73 (37)	27 (19)	13
	U3	0	0	0

source model.

The checks also changed what reached SCIP. The stored-row check (C4) rejected certified, violated rows in every part, because presolve had aggregated or fixed a block variable or because SCIP had dropped or rounded a coefficient (Section 8.4). Binary64 rounding changed at least one coefficient in 13.5% of the exported rows on the MINLPLib models and in 92.5% on the path family; the bound correction of Proposition 6.1 was at most $8.5 \cdot 10^{-16}$ in absolute value and never destroyed a violation.

To measure what the certificate itself buys (Part 4U), we recomputed, for the recorded binary64 direction of each cut, the constant that an uncertified pipeline would use. U1 is the minimum over the separator’s own sample set, which is what the direction LP sees. U2 is the better of U1 and local minimizations by SLSQP [79] from the three best samples, on the block box and the affine domain rows. Up to binary64 rounding, both are upper bounds on the true support value. The MINLPLib set comprises all cuts whose certified value is the exact support value; the 151 cuts certified by Bernstein or ball-arithmetic lower bounds, all in Part 3A, are excluded, because a larger uncertified constant would not show an error. Table 2 compares the constants with the certified values. On the MINLPLib models the sample minimum was wrong for 38% of the cuts, and 405 of the resulting rows would have cut off a feasible point that SCIP had returned for the same model. Local minimization reduced the share of wrong constants to 4.5%, but not all errors disappeared. On `nvs02`, SLSQP stopped at its starting point after one iteration, because the directional derivative along its first step was below its default tolerance; on `kall_ellipsoids_tc02b` it started at stationary points that are not local minima. On the path family, where the blocks are nonconvex by construction, the sample minimum was wrong for almost every cut, 40% of the uncertified rows would have cut off the known optimal solution, and local minimization still left 7.3% wrong constants, some of whose rows cut off the optimum, because SLSQP converged to non-global local minima. Every wrong constant is refuted by an explicit point: the certified minimizer is exactly feasible and attains a smaller value. The counts change little if the removal tolerance is raised tenfold.

A numerical global solve is a much better substitute on these small blocks (Part 5U). Gurobi’s best bound with default tolerances was materially wrong for only one distinct cut, recorded four times on `nvs02`: Gurobi reported optimality with zero gap for a constant $1.8 \cdot 10^{-4}$ above the true minimum: a slope of $-8.9 \cdot 10^{-7}$ in one variable fell inside its optimality tolerance, and over that

Table 3: Full runs on MINLPLib models (300s). Times are shifted geometric means (shift 1 s) over the runs solved by every mode of the part; “SCIP excl.” is SCIP’s solving time minus the separator callback; “ratio” is the median per-run ratio of the total time to SCIP. “Bounds” counts final dual bounds better/worse than SCIP at relative tolerance 10^{-4} over all paired runs. In Part 4D only two models were solved, and times are not comparable.

Part	Mode	Solved	Cuts (runs)	Time (s)	SCIP excl. (s)	Ratio	Bounds
3A (30 models, 3 seeds)	SCIP	80/90	—	1.05	0.99	—	—
	all	80/90	151 (33)	1.27	0.99	1.34	2/4
	auto	80/90	37 (15)	1.21	0.99	1.29	3/3
3B (30 models, 2 seeds)	SCIP	52/60	—	1.33	1.27	—	—
	all	52/60	179 (40)	1.87	1.25	2.66	2/0
	auto	52/60	145 (38)	1.82	1.23	2.55	2/0
4D (20 models, 1 seed)	SCIP	2/20	—				—
	all	2/20	14 (2)				5/1
	auto	2/20	14 (2)				4/0
	SCIP-extra	1/20	—				7/7

variable’s range of 200 this lost $1.8 \cdot 10^{-4}$. A safety shift of 10^{-6} would not have covered this error, although none of these rows cut off a recorded feasible point; Gurobi’s objective value instead of its bound was materially wrong for three of the sampled path cuts. The certified rows themselves were violated at recorded incumbents only within the feasibility tolerance of these incumbents: for example, incumbents of **st_e22** with objective value -85.0000017 , which satisfy the constraints only up to a scaled violation of 10^{-8} , violate the certified objective bound $t \geq -85$, the true optimum, by $1.7 \cdot 10^{-6}$.

8.3 MINLPLib models

On MINLPLib models the cuts were neutral for SCIP: they changed neither which models it solved nor its own solving time. Part 3A repeats the 30 models of campaign 2 (ten per stratum, selected by a hash ordering from 422 eligible models with at most 120 variables) with a 300-second budget, seeds 0, 1 and 2, and root runs. Part 3B uses a sample selected by structure: from the 392 eligible models not used in campaign 2, the 111 on which discovery finds a block that the **auto** rule admits, ranked by a hash, of which we take the first 30, with seeds 0 and 1. The root runs of both parts include mode **all-diag**, whose raised limits test whether the default limits, rather than the cuts, bound their effect. SCIP solves most of these models in about a second, so they can show validity and cost but hardly a benefit. Part 4D addresses this. Its pool consists of the 322 MINLPLib models with 121 to 1,000 variables that are not convex by the library’s metadata, contain a quadratic or polynomial function and were not used before. Of these, 85 have a block that **auto** admits, and 58 of those are not solved by SCIP within 60 seconds; Part 4D takes the first 20 of the 58 in a fixed hash order (Appendix H). Table 3 summarizes the full runs; Tables 9, 10 and 11 in Appendix H give the results per model.

Parts 3A and 3B. In every part and seed, all modes solved the same models. In Part 3A, cuts were added on 11 of the 30 models in mode **all** and on 5 in mode **auto**. The final dual bounds of the 90 paired runs were better in 2 and worse in 4 for **all**, and better in 3 and worse in 3 for **auto**; between seeds, the baseline itself differed in 3 of 30 comparisons, so these differences are within run-to-run variation. In Part 3B, cuts were added on 20 of the 30 models; the final bounds were

better in 2 of 60 comparisons in each cut mode, both on `kall_circles_c8a` at the time limit, and never worse. At the root the cuts had a visible effect on four models of Part 3B and none of Part 3A. They closed the root gap of `pooling_bental4tp` in every cut mode and that of `pooling_bental4pq` in modes `auto` and `all-diag`, which were then solved at the root, and they improved the root bounds of `ex3_1_4` (from -6 to -5.79 , and to -5.69 with raised limits; the optimum is -4) and of `pointpack04` (a maximization model, for which a smaller bound is stronger). In two root runs, `auto` ended with a worse root bound, on `pointpack04` and on `pooling_haverly2pq`, where the bound fell from -617.7 to -857.1 (the optimum is -600). Valid cuts cannot weaken an LP relaxation, but they change SCIP’s later separation and cut selection, and the bound at the end of the root node need not improve. Raising the limits at the root (`all-diag`) added more cuts, 403 in Part 3A and 358 in Part 3B, but improved the root bound beyond the tolerance on only one more model in each part (`cvxnonsep_psig20r` and `pooling_bental4pq`). In the post hoc root runs of Part 3P, whole-row directions changed the number of cuts in 10 of 180 cut-mode runs and no root bound. None of these effects changed a full solve. The cut modes were 15–20% slower than SCIP in shifted geometric mean on Part 3A and 37–41% on Part 3B, and slower per run by median factors of 1.3 and 2.5–2.7; for comparison, baseline runs of Part 3A that differ only in SCIP’s seed differed per run by a median factor of 1.03 (range 0.61 to 1.92). SCIP’s own solving time, without the separator callback, was unchanged (Table 3), and the node counts summed over the commonly solved runs differed by less than 3%. The cuts are therefore neutral for SCIP on these models, and the slowdown is the cost of the separator (Section 8.4).

Presolve and the stored-row check (Part 4B2). In campaign 3 the stored-row check discarded 20% (`all`) and 36% (`all-diag`) of the certified and violated rows of the Part 3B root runs. Reruns that recorded the cause showed that presolve had aggregated or fixed a block variable in 74–85% of these cases, so that SCIP stored the cut in other columns. Part 4B2 repeats the root runs with SCIP’s aggregation and multi-aggregation disabled in every mode. The rejected share fell to 5% and 9.5%; the remaining rejections come from variables that presolve fixed and from SCIP’s own coefficient handling (integer snapping and coefficients below 10^{-9} dropped). Against SCIP with the same presolve setting, mode `all` improved the root bound on 3 models and worsened it on none, and `all-diag` improved it on 6 models (including `pooling_bental4tp`, `pooling_bental4pq`, `ex3_1_4`, `sep1` and `pointpack04`) and worsened it on one (`pooling_haverly1tp`). Disabling aggregation, however, weakened SCIP’s own root bound on 6 of the 30 models, so the setting is not a free remedy. On `kall_circlespolygons_c1p12`, where 127 of 152 violated rows had been rejected before, 104 cuts now reached the LP, and the root bound stayed at 0 against an optimum of 0.34. Trying the support of each whole row first (`all-diag-rowdir`) gave the same cuts and bounds on every model.

SCIP’s disabled separators. SCIP-extra, with SCIP’s disabled-by-default nonconvex separators switched on at the root, improved the root bound on 12 of the 30 Part 3B models and worsened it on one (`sep1`). It closed the root gap completely on five pooling and small nonconvex models and by more than 95% on two more, and these seven include the four models on which the certified cuts had an effect; it solved 13 models at the root against 8. Intersection cuts and interminor cuts were productive on 26 and 21 models, edge-concave cuts on none. On these models, SCIP’s own alternatives are clearly stronger than the certified cuts.

Larger models (Part 4D). In mode `all` the separator stopped during block discovery on 12 of the 20 models, exactly those on which discovery had taken more than the one-second allowance in the scan (1.0 to 7.7 s). On six of the other eight it completed discovery and added no cut: every

certified support either fell below the violation threshold or, on three models, was rejected by the stored-row check because presolve had fixed or aggregated a block variable (11 of the 25 violated rows). Discovery also found more blocks than the cap of 32 on 6 of these 8 models, and on 10 of the 20 with the cap of 128 in mode `all-diag`; the blocks tried were then those with the smallest variable indices. With raised limits at the root (`all-diag`), the separator added 891 cuts on 13 models and improved the root bound on one model (`multiplants_mtg1c`, closing 4% of its root gap), while the root bound was worse on five, by at most 3% of the gap; the variant with whole-row directions behaved the same. In the full runs the cut modes solved the same two models as the baseline. All final-bound differences in Table 3 occurred on models that received no cut, so they reflect the changed timing of a time-limited run, not the cuts. SCIP-extra improved the root bound on 8 of the 20 models, worsened it on 3, and in full runs lost one of the two solved models. On models of this size the separator with default limits stops during discovery on most models, and where it gets past discovery its cuts are not stronger than SCIP’s relaxation.

8.4 What the separator does and what it costs

Table 8 in Appendix H follows the certified rows from the support calls to the cuts that reach SCIP. Three observations explain most of the results of Section 8.3.

First, most certified supports are not violated at the LP point: on the MINLPLib models, 65 to 86% of the certified rows fell below the violation threshold, because SCIP’s relaxation already implied them up to the threshold. Certification failed in few support calls. In the campaign-3 reruns that recorded the outcome, every failure was *incomplete* (a Bernstein or ball-arithmetic bound reached its depth budget, or did not certify the required value) or *unsupported* (a quadratic block whose number of row subsets exceeded the enumeration budget and which the fallback bounds do not cover); none was an empty domain. In Part 4D, 307 of the 313 failures came from one model.

Second, the stored-row check lost a substantial share of the violated rows on MINLPLib models, mainly to presolve aggregation (Section 8.3), and a small share on the constructed families, where SCIP dropped coefficients below its epsilon of 10^{-9} in at most 3% of the violated rows. Each such loss would be a silent change of a certified row if the check were absent. A separator that declares its variables non-aggregatable before presolve avoids most of the presolve losses; we did this only as an experiment (Part 4B2), because it changes SCIP’s own relaxation. A separator that works in the presolved space would instead have to treat SCIP’s reductions, which are made in floating point, as part of the model that the certificate refers to. The losses to dropped and snapped coefficients could be avoided by applying SCIP’s coefficient handling in the exporter, followed by the correction of Proposition 6.1.

Third, on the MINLPLib models the whole slowdown is the separator’s own time. SCIP’s solving time without the callback was unchanged in every part: its shifted geometric mean differed from SCIP’s by at most 4% in Parts 3A and 3B, and the median per-run ratio was between 0.96 and 1.01 in Parts 4B2 and 4D. The callback cost about half a second per run in Part 3B, more than SCIP’s whole solving time on most of these models (median 0.16 s). In campaign 3, exact certification took 70% of the callback time on Part 3B, the direction search and discovery the rest; on the larger models of Part 4D, discovery took 85% of the callback time in mode `all`. All of this is Python code; a compiled implementation would reduce it, but on these models it would at best break even, because the cuts did not reduce SCIP’s work. On the constructed families, in contrast, the cuts reduced SCIP’s own time sharply (Section 8.5), and the callback was the dominant remaining cost.

8.5 A family with the path structure

The benchmark models do not tell whether the cuts can supply strength that SCIP lacks where the theory predicts a gap. We therefore built a family with the structure of Section 3.2. An instance with n copies has variables $x_i, y_i, z_i \in [0, 1]$ and $t_i \in [-10, 10]$, the rows $\Phi^{(i)}(x_i, y_i, z_i) - t_i \leq 0$, where $\Phi^{(i)}$ is the function (6) with $\omega_A = \omega_C = 1$, one coupling row $\sum_i y_i \leq 0.8n$, and the objective $\min \sum_i t_i$. For each copy, four distinct values are drawn from $\{k/64 : k = 0, \dots, 48\}$ and split into two interleaving two-point sets A_i and C_i with random orientation. All coefficients are dyadic and exactly representable. Every optimal y_i is a midpoint of two values of at most $3/4$, so the coupling row does not bind, and by Theorem 3.2 the optimum is $\sum_i \delta_i^2/2$. Glued exact pair hulls give the bound 0 for every copy: the chords of A_i and C_i cross at a point with $m_{y_i} \leq \max A_i \leq 3/4$, so the zero points of all copies together satisfy the coupling row. The part of the gap that only the joint block closes is therefore the whole optimum. Part 3C used five instances for each $n \in \{10, 20, 40, 80\}$ (seeds 0–4); Part 4C3 adds five fresh instances per n (seeds 5–9), which no earlier run had used. The randomized family was adopted after a pilot with SCIP alone had shown that the deterministic family, with every copy equal to the example of Proposition 3.1, is easy for SCIP: it solved 160 copies in 19 s.

Because the objective is a sum of block functions and the coupling row does not bind, the n cuts $t_i \geq \min \Phi^{(i)}$, which are the exact supports of the blocks in their whole-row directions, make the root bound equal to the optimum. This is the most favorable case for block cuts. Part 4C4 removes it: it replaces the coupling row of each fresh instance by $\sum_i y_i \leq c$, where c is half the sum of the smallest unconstrained minimizers y_i^* , rounded down to a multiple of $1/64$. The row then binds, the block minima no longer give the optimum, and the cuts must be found in other directions. For Part 4C4 we know two reference values to high accuracy. The first is the optimum, from a convex mixed-integer reformulation (one binary variable per choice of nearest points) solved by Gurobi, re-solved exactly for the returned choice and certified by an exact branch and bound to within 10^{-58} . The second is the *block closure*: the bound of all aggregated cuts on the blocks (x_i, y_i, z_i) together with the linear rows. Each block has the single nonlinear row $\Phi^{(i)} - t_i \leq 0$ and the domain $[0, 1]^3$, because the coupling row involves other blocks and, since $c \geq 1$, implies no tighter bounds. By Theorem 5.2 and Corollary 5.3 the closure of the block is $t_i \geq \text{vex } \Phi^{(i)}(x_i, y_i, z_i)$; minimizing over x_i and z_i commutes with convexification, because the projection of a convex hull is the convex hull of the projection. With $d_i(y) = \min_{x, z \in [0, 1]} \Phi^{(i)}(x, y, z) = \text{dist}(y, A_i)^2 + \text{dist}(y, C_i)^2$ the block closure is therefore $\min\{\sum_i \text{vex } d_i(y_i) : \sum_i y_i \leq c, 0 \leq y_i \leq 1\}$, which is also the Lagrangian dual bound of the coupling row. We computed it with the exact convex envelopes of the piecewise quadratic functions d_i and a bisection on the multiplier, to within $2 \cdot 10^{-28}$. Table 4 summarizes the results; Tables 12 and 14 in Appendix H give them per instance.

SCIP’s root bound and the pair-hull level. SCIP’s root bound is weak: on the instances of Part 3C it lies between -0.035 and -0.013 per copy, while each copy contributes at least $1/8192$ to the optimum. The cause is a presolve step: $\Phi^{(i)}$ is concave in x_i and in z_i , so SCIP’s implicit-discreteness detection (Section 2.1) makes these variables binary, and SCIP then closes the gap by branching. SCIP-nolocks, with this step switched off, raises the root bound to just below 0, the bound of glued exact pair hulls, on every instance of Part 3C; SCIP’s semidefinite minor cuts were productive in all of these runs and in none of the default runs of campaign 4. On the fresh instances of Part 4C3, SCIP-nolocks (Part 5C-a) again stopped just below the pair-hull level, with a median closure of 0.92. SCIP-extra, of whose extra separators only the intersection cuts were productive, closed about 60% of SCIP’s root gap. None of the three SCIP settings we ran closed any part of the gap between the pair-hull level and the optimum, which is the gap that Section 3 attributes to joint

Table 4: The path family: median fraction of SCIP’s root gap closed, $(z_{\text{mode}} - z_{\text{base}})/(z^* - z_{\text{base}})$, for each n (five instances each), and instances solved within 300 s (of 20). “Remainder” and “whole row” are the first directions of the separator, “base” and “wide” its limits (Section 7.2, Table 7). “Glued pair hulls” is the analytic bound 0 of Theorem 3.2 and “block closure” the bound of all cuts on the blocks; no run computes them. SCIP closes 0 by definition, and Gurobi had no root runs. Rows of different parts ran under different host loads, so their solved counts are not paired (Section 8.1). The binding-row block shows three decimals; its Part 5C-b rows are root runs only.

Instances	Part	Mode	Root gap closed (median)				Solved
			$n = 10$	$n = 20$	$n = 40$	$n = 80$	
seeds 0–4	3C	SCIP					9
	4C2	SCIP-nolocks	0.95	0.89	0.93	0.91	10
	4C2	SCIP-extra	0.64	0.58	0.66	0.59	10
	4C2	Gurobi					7
	—	glued pair hulls	0.95	0.89	0.93	0.92	
	3C	remainder, base	0.35	0.25	0.19	0.22	9
	4C2	remainder, wide	0.86	0.87	0.89	0.90	13
	3P	whole row, base	0.44	0.44	0.43	0.46	10
	3P	whole row, wide	1.00	1.00	1.00	1.00	20
seeds 5–9	4C3	SCIP					9
	5C-a	SCIP-nolocks	0.92	0.92	0.92	0.92	10
	4C3	SCIP-extra	0.57	0.64	0.61	0.66	9
	4C3	Gurobi					5
	—	glued pair hulls	0.93	0.92	0.93	0.92	
	4C3	remainder, base	0.32	0.31	0.24	0.20	10
	4C3	whole row, wide	1.00	1.00	1.00	1.00	20
binding row (seeds 5–9)	4C4	SCIP					10
	5C-a	SCIP-nolocks	0.980	0.980	0.972	0.959	20
	4C4	SCIP-extra	0.287	0.600	0.685	0.655	13
	4C4	Gurobi					6
	—	block closure	1.000	1.000	1.000	1.000	
	4C4	remainder, wide	0.983	0.960	0.954	0.933	20
	4C4	whole row, wide	0.997	0.989	0.981	0.982	20
	5C-b	remainder, cap $64n$	1.000	0.999	1.000	0.999	
	5C-b	whole row, cap $64n$	1.000	0.999	1.000	0.999	

blocks. Gurobi 13 solved the instances of Part 4C2 with $n = 10$ and two with $n = 20$; on $n = 40$ and $n = 80$ its gap after 300 s was 61 to 113% (53 to 121% on the instances of Part 4C3).

The prospective separator and the attribution. The separator of Part 3C (remainder directions, base limits) improved the root bound on all 20 instances but closed only a median of 19–35% of the root gap, less than the pair-hull level. It solved the same 9 instances as SCIP, with fewer nodes and a shifted geometric mean time of 6.7 s instead of 10.3 s. After seeing these results we identified two limitations. First, the remainder directions bound only the nonlinear remainder of a row, whereas the support of the whole row is the block minimum. Second, the first blocks of a callback take up to four cuts each, so with n cuts per callback the later blocks receive none. A post hoc diagnostic (Part 3P) tried whole-row directions first, once with base limits and once with wide limits; we had tried the wide limits by hand on two of the 20 instances before. With base limits, whole-row directions raised the median closure to 43–46% and solved 10 instances; with wide limits they closed the root gap on all 20 instances and solved all 20, 18 of them at the root node. Part 4C2

ran the remaining combination prospectively: remainder directions with wide limits closed 86–90%, close to the pair-hull level and above it on 7 of the 20 instances, and solved 13. Every run of these four configurations stopped at its cut cap, so their root bounds do not depend on the host load, which was lower in campaign 4. The wide limits therefore account for most of the closure; on this family the whole-row directions supply the rest and raise the solved count from 13 to 20.

Prospective test on fresh instances. On the 20 instances of Part 4C3, the whole-row variant with wide limits, fixed before these instances were generated, closed the root gap on every instance: its root bound was within $2 \cdot 10^{-7}$ of the optimum. It solved all 20 in 3 to 33 s, with a median of one node and at most 38, against 9 for SCIP, 9 for SCIP-extra, 10 for SCIP-nolocks, 10 for the separator with remainder directions and base limits, and 5 for Gurobi. On the nine instances that SCIP solved, the cut mode was slower on four of the five with $n = 10$, by factors of up to 3.4 because of the separator’s fixed cost, and 4 to 20 times faster on the four with $n = 20$; SCIP’s own time fell to a median of 4% of its time in the baseline. Over all 20 runs, 96% of the time was spent in the Python separator. The separator with remainder directions and base limits closed 20–32% of the root gap, as in Part 3C.

A binding coupling row (Part 4C4). When the coupling row binds, the block minima no longer give the optimum, and the cuts must have the right slopes in y_i . Both separator configurations with wide limits solved all 20 instances, in 3 to 134 s, against 10 for SCIP, 13 for SCIP-extra and 6 for Gurobi; on the ten instances that SCIP solved, all with $n \leq 20$, the cut modes were slower, by median factors of 3.5 (remainder directions) and 3.2 (whole-row directions), because of the separator’s time. SCIP-nolocks, however, also solved all 20 instances, in 0.1 to 169 s and faster than the cut modes (shifted geometric means 2.0 s against 12 to 13 s), with a root bound that closed 96 to 98% of SCIP’s root gap. With a binding row, most of the gap is thus within reach of SCIP’s own relaxation once implicit discreteness is switched off, and the part that only the joint blocks close is small. At the root they closed a median of 95% and 99% of SCIP’s root gap. The block closure equals the optimum on 13 of the 20 instances and lies at most $6.5 \cdot 10^{-4}$ below it on the others. No root run reached it: every cut-mode run stopped at its cap of $16n$ cuts, and with whole-row directions the root bound stayed between 0.016 and 0.029 below the block closure for $n = 80$. With larger cut caps and up to 40 callbacks (Part 5C-b), separation ended by itself after 14 to 27 callbacks, with about 26 to 30 cuts per block, and the root bound came within $8 \cdot 10^{-4}$ of the block closure on every instance; the caps of $32n$ and $64n$ gave the same results. The distance left with the $16n$ cap was therefore due to the cap, not to the direction search. The cuts that did the work came from the LP direction search: 99.9% (remainder directions) and 94% (whole-row directions) of the cuts were LP directions, almost all with a nonzero coefficient of y_i , mostly with the negative slope that the binding row requires; none of them is an envelope cut in (y_i, t_i) alone. The whole-row cuts alone give only $\sum_i \min d_i$, which is below SCIP’s root bound on half of the instances.

Both path families have a convex mixed-integer reformulation, the one we used to compute the optima, which Gurobi solves in under a second on 19 of the 20 instances of Part 4C4; the experiment therefore tests whether a solver recovers the block structure from the nonconvex formulation. With a single coupling row, the block closure is close to the optimum for a known reason: by the Shapley–Folkman estimate of Aubin and Ekeland [4], the duality gap of a separable problem with one coupling constraint is at most twice the largest nonconvexity $\max_y (d_i(y) - \text{vex } d_i(y))$ of a single block. These families were constructed to have the structure of Section 3, and their objectives are sums of block functions, which favors block cuts. On the family with a non-binding row, exact block support supplied root strength that none of the three SCIP settings obtained, and the separator

solved instances that neither SCIP nor Gurobi solved within 300 s; this strength reached the solver only when the separator could add enough cuts per block in useful directions. With a binding row, SCIP without implicit discreteness was as successful and faster. The families are not evidence about other models.

8.6 Stars

The blocks of Sections 8.3–8.5 have at most four variables, so the star algorithm of Theorem 4.4 was not needed there. We test it in two ways: offline at large scale, and inside the solver on a family whose merged blocks are stars with up to 16 leaves.

The algorithm at scale (Part 4S). On random constrained stars with $k \in \{10, 100, 1000, 10000\}$ leaves (five instances each), $2k$ center–leaf rows, one to three rows in the center alone and small rational coefficients of both signs, we compared an implementation of the sweep of the proof with the simpler star oracle of our separator (Section 4.2). Where both ran ($k \leq 1000$), the two minima were equal as rational numbers, and on all 20 instances the sweep’s minimizer was exactly feasible and attained the reported value. The sweep took a median of 1.3 ms, 12 ms, 0.17 s and 2.1 s of CPU time for the four values of k , close to linear growth, whereas the simpler oracle took 6 ms, 0.43 s and 29 s for $k \leq 1000$, about 70 times more per tenfold increase in k ; at $k = 10^4$ it was not run, because its extrapolated time exceeded ten minutes per instance. The minima had numerators and denominators of at most 94 and 78 bits, with no steady growth in k , in line with the size bounds of Appendix C.

Stars in the solver (Part 5S). An instance has $n \in \{10, 20\}$ stars with $k \in \{4, 8, 16\}$ leaves each (five instances per pair, 30 in total). Star i has a center $y_i \in [0, 1]$ and leaves $x_{ij} \in [0, 1]$ with rows $\Phi_{ij}(x_{ij}, y_i) - t_{ij} \leq 0$, where Φ_{ij} is the function Φ_A of (6) with $\omega_A = 1$ for a random two-point set $A_{ij} \subset \{m/64 : m = 0, \dots, 48\}$, and one center–leaf row $x_{ij} - y_i \leq d_{ij}$ with $d_{ij} \in \{1/8, 1/4, 1/2, 1\}$; a coupling row $\sum_i y_i \leq 0.8n$ does not bind, and the objective is $\min \sum_{ij} t_{ij}$. The optimum of each star is the minimum over y of a sum of k piecewise quadratic functions, which we computed exactly and checked against an independent implementation of the sweep. Pairs and blocks of four variables miss the joint minimum of a star whenever the leaves’ preferred centers have no common point (Proposition 3.5), so the family tests whether the whole star is worth relaxing jointly. We compared SCIP, SCIP-nolocks, SCIP-extra and Gurobi with three separator modes that differ only in their blocks and first directions (Table 7): blocks of at most four variables with whole-row directions or with the block direction first, and star blocks with the block direction first.

Table 5 gives the results. With star blocks, the root bound was within 10^{-6} (relative) of the optimum on all 30 instances, whereas SCIP’s root bound reached a median of 95%, 82% and 66% of the optimum for $k = 4, 8$ and 16 , and SCIP-extra 94%, 86% and 74%. Here SCIP’s implicit-discreteness presolve helps: SCIP-nolocks reached only 53 to 58%. Blocks of four variables added less than two percentage points to SCIP’s root bound, with either direction rule: a block that contains three leaves of a star cannot see the joint minimum of all k . The star oracle certified all 4,650 of its support calls in the root runs, on blocks of 5, 9 and 17 variables, with a median of 7.8 ms and a maximum of 0.2 s per call, whereas the enumeration of Theorem 4.1 took a median of 79 to 114 ms on the blocks of four variables. The star mode solved 19 instances, more than any SCIP setting (at most 15) and than the separator with smaller blocks (at most 14), but fewer than Gurobi (22). Where it did not solve an instance, its bound was already optimal, but SCIP had not found a near-optimal solution: on the largest instances its incumbent was 7 to 10% above the optimum after 300 s, while Gurobi’s incumbent was within $2 \cdot 10^{-4}$ (relative) of the optimum and its bound up to

Table 5: The star family (Part 5S): median root bound as a fraction of the optimum, and instances solved within 300 s, for each number k of leaves (ten instances each: $n = 10$ and $n = 20$ stars). Gurobi had no root runs. Eight full runs stopped at the hard process limit during a load spike count as unsolved (see the text).

Mode	Root bound / optimum			Solved			
	$k = 4$	$k = 8$	$k = 16$	$k = 4$	$k = 8$	$k = 16$	total
SCIP	0.954	0.815	0.655	7	3	0	10
SCIP-nolocks	0.526	0.580	0.547	10	5	0	15
SCIP-extra	0.943	0.858	0.738	10	3	0	13
Gurobi				9	8	5	22
whole row, ≤ 4 variables	0.956	0.821	0.651	10	2	0	12
block direction, ≤ 4 variables	0.963	0.834	0.664	10	4	0	14
block direction, stars	1.000	1.000	1.000	9	9	1	19

1.4% below it. The separator also used its whole allowance of 60 s, mostly on the enumeration for the blocks of four variables that it tried after the star blocks. Eight full runs (three of SCIP, one of Gurobi and four of the separator modes) were stopped at the hard process limit during a 13-minute period in which an unrelated process raised the host load to 126; we count them as unsolved. A rerun of these eight runs at normal load, outside the protocol, solved three of them, two of SCIP and one of the separator with whole-row directions and blocks of four variables; with them, SCIP would solve 12 instances and that mode 13, and the comparison above would not change.

8.7 Summary and limitations

Every one of the 274,489 recorded cuts passed replay against the source model and the row stored by SCIP. Constants taken from the separator’s samples would have cut off feasible points on seven MINLPLib models and on 58 of the 60 sampled path-family instances, and constants improved by local minimization still did so on two models and 19 instances; a numerical global solve was almost always right, but not always. On MINLPLib models the cuts were neutral for SCIP: they changed neither which models it solved in campaigns 2 to 4 nor its own solving time, the Python separator made the runs slower, and SCIP’s disabled-by-default separators improved the root bound more often than the certified cuts did. On the constructed families, exact block support closed the root gap that SCIP leaves. With a non-binding coupling row, a separator configuration chosen after the campaign-3 diagnostic and fixed before the fresh instances were generated solved all 20 of them, against at most 10 for each SCIP setting and 5 for Gurobi. With a binding row, both direction rules solved all instances, but so did SCIP without implicit discreteness, faster. On the star family, star blocks certified by the star oracle made the root bound exact, which neither blocks of four variables nor any SCIP setting achieved; the solved count was then limited by SCIP’s search for good solutions, and Gurobi solved more instances. We therefore do not recommend enabling the separator by default; its value is on models whose joint block structure SCIP’s relaxation misses, and an activation rule that recognizes such models is not available.

Limitations. The MINLPLib parts use models with at most 1,000 variables, and their blocks have at most four variables; star blocks occur only in the constructed star family. The implementation is in Python, so the measured cost is an upper bound on that of a compiled separator. The host was shared, with different loads in different campaigns and one load spike during Part 5S, so we compare times only within paired runs of one part, and time differences of a few percent are not

meaningful [85]. Runs used one seed, except Parts 3A and 3B, and one thread. Gurobi was run with one thread and default settings apart from its nonconvex mode. The constructed families favor block cuts; the post hoc diagnostic of campaign 3 was designed after its results were seen, and its finding was then tested prospectively on fresh instances.

9 Discussion and conclusions

Two points of the theory matter for solver design. First, pair hulls glued on finitely many moments of a shared variable accept distributions that no joint distribution reproduces, and they give no positive bound exactly when the local zero sets alternate often enough. On three-variable paths, dense moment relaxations with the missing product close these gaps, so there the advantage of joint blocks over their pairs is one of representation, which support cuts obtain in the model’s own coordinates. For larger blocks this is no longer true: on a four-variable path, and with a relative gap of nearly 100% on the full-gap instances of Proposition 3.4 with five and seven variables, an exact support cut of the block is strictly stronger than the dense first-level relaxation with all products. Second, for stars, the merged blocks of Section 3, the support problem remains exactly solvable in nearly linear time when rows couple the center with one leaf at a time, and, as in safe-cut practice, certification does not constrain how directions are found: only the final binary64 row, its bound over the whole block domain and the row that the solver stores need to be checked, and all of this can be replayed.

The computations add two findings to the theory. First, certification matters in practice. The sample minimum would have given wrong constants for 38% of the MINLPLib cuts and for almost all sampled path-family cuts, and local minimization still left wrong constants that cut off feasible solutions, including known optima. A numerical global solve was a much better substitute on these small blocks, but on one model it reported optimality for a constant $1.8 \cdot 10^{-4}$ above the true minimum, which a fixed safety shift of 10^{-6} does not cover. Second, the stored-row check exposed an obstacle that the theory does not see: presolve aggregates or fixes the variables to which a block refers, so a separator stated in the original variables loses cuts unless it works in the presolved space or prevents these reductions, and preventing them weakens SCIP’s own relaxation.

On MINLPLib models the cuts did not help SCIP: they improved its root bound on few models and changed no full solve, while SCIP’s disabled-by-default separators achieved as much or more. On the constructed families, where the theory predicts a gap, the picture is different. With a non-binding coupling row, native SCIP stopped at the level of glued pair relaxations in all three settings we ran, and the certified cuts went beyond it when the separator could add enough cuts per block and tried the exact support of each whole row first; this configuration solved every instance, against at most half of them for each SCIP setting and for Gurobi. With a binding row, both direction rules solved every instance, and the cuts that mattered came from the LP direction search; but SCIP with its implicit-discreteness presolve step switched off solved every instance as well, and faster, because its own relaxation already captured most of the gap. On the star family, star blocks certified by the star oracle made the root bound exact, whereas blocks of four variables added almost nothing to SCIP’s bound; the solved count was then limited by SCIP’s search for good solutions, where Gurobi did better. Joint blocks larger than the separator’s default of four variables thus matter, and the near-linear star algorithm makes them cheap to certify.

Several questions remain open. Exact support for trees deeper than stars with edge-local rows would extend the star algorithm to the structures of network models; Remark 4.5 shows that the rational partition argument does not extend directly. Proposition 3.5(ii) reduces the gap of the glued relaxation in any direction to a concave maximization over re-splittings, but which directions

to merge, and when merging pays off, is not answered by it. A direction search with the guarantees of Appendix F and the speed of a solver callback is not available. Finally, an activation rule that predicts when joint cuts pay off would be needed before any default use; our structural rule did not achieve this.

Code and data availability

The supplementary archive of this article contains the separator and the model builder, the replay checker, the frozen code snapshots of every campaign with their SHA-256 manifests, the protocols, the job lists and instance selections with the instance generators and reference values, the run records and logs, the replay outputs, the scripts that produce every table and number of Section 8 from the records, the independent recomputations of these numbers, and the verification scripts for the theorems and examples (exact rational and symbolic checks; the two semidefinite values of Section 3.4 are computed numerically). The experiments used SCIP 10.0.2 through PySCIPOpt 6.2.1, Python 3.13.11, SymPy 1.14.0, python-flint 0.9.0 (Arb), NumPy 2.5.3, SciPy 1.18.1 (HiGHS [69] for the direction LPs and SLSQP [79] for the ablation), mpmath 1.3.0, Gurobi 13.0.3, and CVXPY 1.7.3 with Clarabel 0.11.1 for the semidefinite checks, on an Intel Xeon w5-2565X (18 cores) under Linux (WSL2), and a local copy of MINLPLib [28, 125] whose files are archived with their hashes. Replay needs Python with SymPy, python-flint, PySCIPOpt, NumPy and SciPy; it constructs, but does not solve, a SCIP model. The Gurobi runs and the reference optima of Part 4C4 require a Gurobi licence; replay and all other checks do not.

References

- [1] Claire S. Adjiman, Stefan Dallwig, Christodoulos A. Floudas, and Arnold Neumaier. A global optimization method, α BB, for general twice-differentiable constrained NLPs — I. theoretical advances. *Computers & Chemical Engineering*, 22(9):1137–1158, 1998. doi: 10.1016/S0098-1354(98)00027-1.
- [2] Kurt M. Anstreicher. On convex relaxations for quadratically constrained quadratic programming. *Mathematical Programming*, 136(2):233–251, 2012. doi: 10.1007/s10107-012-0602-3.
- [3] Kurt M. Anstreicher and Samuel Burer. Computable representations for convex hulls of low-dimensional quadratic forms. *Mathematical Programming*, 124(1–2):33–43, 2010. doi: 10.1007/s10107-010-0355-9.
- [4] J. P. Aubin and I. Ekeland. Estimates of the duality gap in nonconvex optimization. *Mathematics of Operations Research*, 1(3):225–245, 1976. doi: 10.1287/moor.1.3.225.
- [5] Martin Ballerstein. *Convex Relaxations for Mixed-Integer Nonlinear Programs*. PhD thesis, ETH Zurich, 2013. Diss. ETH No. 21024, doi: 10.3929/ethz-a-009959194.
- [6] Xiaowei Bao, Nikolaos V. Sahinidis, and Mohit Tawarmalani. Multiterm polyhedral relaxations for nonconvex, quadratically constrained quadratic programs. *Optimization Methods and Software*, 24(4–5):485–504, 2009. doi: 10.1080/10556780902883184.
- [7] Xiaowei Bao, Aida Khajavirad, Nikolaos V. Sahinidis, and Mohit Tawarmalani. Global optimization of nonconvex problems with multilinear intermediates. *Mathematical Programming Computation*, 7(1):1–37, 2015. doi: 10.1007/s12532-014-0073-z.

- [8] Pietro Belotti, Jon Lee, Leo Liberti, François Margot, and Andreas Wächter. Branching and bounds tightening techniques for non-convex MINLP. *Optimization Methods and Software*, 24(4–5):597–634, 2009. doi: 10.1080/10556780903087124.
- [9] Alberto Bemporad, Manfred Morari, Vivek Dua, and Efstratios N. Pistikopoulos. The explicit linear quadratic regulator for constrained systems. *Automatica*, 38(1):3–20, 2002. doi: 10.1016/S0005-1098(01)00174-1.
- [10] Michael Ben-Or. Lower bounds for algebraic computation trees. In *Proceedings of the Fifteenth Annual ACM Symposium on Theory of Computing (STOC 1983)*, pages 80–86, 1983. doi: 10.1145/800061.808735.
- [11] Umberto Bertelè and Francesco Brioschi. *Nonserial Dynamic Programming*, volume 91 of *Mathematics in Science and Engineering*. Academic Press, New York, 1972.
- [12] Ksenia Bestuzheva, Mathieu Besançon, Wei-Kun Chen, Antonia Chmiela, Tim Donkiewicz, Jasper van Doornmalen, Leon Eifler, Oliver Gaul, Gerald Gamrath, Ambros Gleixner, Leona Gottwald, Christoph Graczyk, Katrin Halbig, Alexander Hoen, Christopher Hojny, Rolf van der Hulst, Thorsten Koch, Marco Lübbecke, Stephen J. Maher, Frederic Matter, Erik Mühmer, Benjamin Müller, Marc E. Pfetsch, Daniel Rehfeldt, Steffan Schlein, Franziska Schlösser, Felipe Serrano, Yuji Shinano, Boro Sofranac, Mark Turner, Stefan Vigerske, Fabian Wegscheider, Philipp Wellner, Dieter Weninger, and Jakob Witzig. The SCIP Optimization Suite 8.0. ZIB-Report 21-41, Zuse Institute Berlin, 2021. arXiv:2112.08872.
- [13] Ksenia Bestuzheva, Antonia Chmiela, Benjamin Müller, Felipe Serrano, Stefan Vigerske, and Fabian Wegscheider. Global optimization of mixed-integer nonlinear programs with SCIP 8. *Journal of Global Optimization*, 91(2):287–310, 2025. doi: 10.1007/s10898-023-01345-1.
- [14] Ksenia Bestuzheva, Ambros Gleixner, and Tobias Achterberg. Efficient separation of RLT cuts for implicit and explicit bilinear terms. *Mathematical Programming*, 210(1–2):47–74, 2025. doi: 10.1007/s10107-024-02104-0.
- [15] Aaresh Bhathena, Salar Fattahi, Andrés Gómez, and Simge Küçükyavuz. A parametric approach for solving convex quadratic optimization with indicators over trees. *Mathematical Programming*, 218(1–2):291–336, 2026. doi: 10.1007/s10107-025-02222-3.
- [16] Daniel Bienstock and Gonzalo Muñoz. LP formulations for polynomial optimization problems. *SIAM Journal on Optimization*, 28(2):1121–1150, 2018. doi: 10.1137/15M1054079.
- [17] Daniel Bienstock, Chen Chen, and Gonzalo Muñoz. Outer-product-free sets for polynomial optimization and oracle-based cuts. *Mathematical Programming*, 183(1–2):105–148, 2020. doi: 10.1007/s10107-020-01484-3.
- [18] Grigoriy Blekherman, Santanu S. Dey, and Shengding Sun. Aggregations of quadratic inequalities and hidden hyperplane convexity. *SIAM Journal on Optimization*, 34(1):98–126, 2024. doi: 10.1137/22M1528215.
- [19] Natasha Boland, Santanu S. Dey, Thomas Kalinowski, Marco Molinaro, and Fabian Rigterink. Bounding the gap between the McCormick relaxation and the convex hull for bilinear functions. *Mathematical Programming*, 162(1–2):523–535, 2017. doi: 10.1007/s10107-016-1031-5.

- [20] Dominik Bongartz and Alexander Mitsos. Deterministic global optimization of process flowsheets in a reduced space using McCormick relaxations. *Journal of Global Optimization*, 69(4):761–796, 2017. doi: 10.1007/s10898-017-0547-4.
- [21] Glencora Borradaile and Pascal Van Hentenryck. Safe and tight linear estimators for global optimization. *Mathematical Programming*, 102(3):495–517, 2005. doi: 10.1007/s10107-004-0533-8.
- [22] E. Andrew Boyd. Fenchel cutting planes for integer programs. *Operations Research*, 42(1): 53–64, 1994. doi: 10.1287/opre.42.1.53.
- [23] Stephen Boyd and Lieven Vandenbergh. *Convex Optimization*. Cambridge University Press, 2004. URL https://web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf.
- [24] Marilyn Breen. Primitive Radon partitions for cyclic polytopes. *Israel Journal of Mathematics*, 15(2):156–157, 1973. doi: 10.1007/BF02764601.
- [25] Stephen Brierley, Miguel Navascués, and Tamás Vértési. Convex separation from convex optimization for large-scale problems. arXiv:1609.05011 [quant-ph], 2016. URL <https://arxiv.org/abs/1609.05011>.
- [26] Samuel Burer and Adam N. Letchford. On nonconvex quadratic programming with box constraints. *SIAM Journal on Optimization*, 20(2):1073–1089, 2009. doi: 10.1137/080729529.
- [27] Samuel Burer, Karthik Natarajan, and Rick Willemsen. On the semidefinite representability of continuous quadratic submodular minimization with applications to pricing and moment problems. arXiv:2504.03996, 2025. URL <https://arxiv.org/abs/2504.03996>.
- [28] Michael R. Bussieck, Arne Stolbjerg Drud, and Alexander Meeraus. MINLPLib—a collection of test models for mixed-integer nonlinear programming. *INFORMS Journal on Computing*, 15(1):114–119, 2003. doi: 10.1287/ijoc.15.1.114.15159.
- [29] Rui Chen and James Luedtke. On generating Lagrangian cuts for two-stage stochastic integer programs. *INFORMS Journal on Computing*, 34(4):2332–2349, 2022. doi: 10.1287/ijoc.2022.1185.
- [30] E. W. Cheney and A. A. Goldstein. Newton’s method for convex programming and Tchebycheff approximation. *Numerische Mathematik*, 1(1):253–268, 1959. doi: 10.1007/BF01386389.
- [31] Kevin K. H. Cheung, Ambros Gleixner, and Daniel E. Steffy. Verifying integer programming results. In *Integer Programming and Combinatorial Optimization (IPCO 2017)*, volume 10328 of *Lecture Notes in Computer Science*, pages 148–160. Springer, 2017. doi: 10.1007/978-3-319-59250-3_13.
- [32] Antonia Chmiela, Gonzalo Muñoz, and Felipe Serrano. On the implementation and strengthening of intersection cuts for QCQPs. *Mathematical Programming*, 197(2):549–586, 2023. doi: 10.1007/s10107-022-01808-5.
- [33] Vašek Chvátal, William Cook, and Daniel Espinoza. Local cuts for mixed-integer programming. *Mathematical Programming Computation*, 5(2):171–200, 2013. doi: 10.1007/s12532-013-0052-9.
- [34] William Cook, Sanjeeb Dash, Ricardo Fukasawa, and Marcos Goycoolea. Numerically safe Gomory mixed-integer cuts. *INFORMS Journal on Computing*, 21(4):641–649, 2009. doi: 10.1287/ijoc.1090.0324.

- [35] William Cook, Thorsten Koch, Daniel E. Steffy, and Kati Wolter. A hybrid branch-and-bound approach for exact rational mixed-integer programming. *Mathematical Programming Computation*, 5(3):305–344, 2013. doi: 10.1007/s12532-013-0055-6.
- [36] Danial Davarnia, Jean-Philippe P. Richard, and Mohit Tawarmalani. Simultaneous convexification of bilinear functions over polytopes with application to network interdiction. *SIAM Journal on Optimization*, 27(3):1801–1833, 2017. doi: 10.1137/16M1066166.
- [37] Alberto Del Pia and Aida Khajavirad. Treewidth and the complexity of box-constrained quadratic programs. arXiv:2609.35595, 2026. URL <https://arxiv.org/abs/2609.35595>.
- [38] Alberto Del Pia, Santanu S. Dey, and Marco Molinaro. Mixed-integer quadratic programming is in NP. *Mathematical Programming*, 162(1–2):225–240, 2017. doi: 10.1007/s10107-016-1036-0.
- [39] Santanu S. Dey and Aida Khajavirad. A second-order cone representable class of nonconvex quadratic programs. *Mathematical Programming*, 2026. doi: 10.1007/s10107-026-02364-y.
- [40] Santanu S. Dey, Gonzalo Muñoz, and Felipe Serrano. On obtaining the convex hull of quadratic inequalities via aggregations. *SIAM Journal on Optimization*, 32(2):659–686, 2022. doi: 10.1137/21M1428583.
- [41] Ferenc Domes and Arnold Neumaier. Rigorous filtering using linear relaxations. *Journal of Global Optimization*, 53(3):441–473, 2012. doi: 10.1007/s10898-011-9722-1.
- [42] Ferenc Domes and Arnold Neumaier. Constraint aggregation for rigorous global optimization. *Mathematical Programming*, 155(1–2):375–401, 2016. doi: 10.1007/s10107-014-0851-4.
- [43] Yoel Drori and Marc Teboulle. An optimal variant of Kelley’s cutting-plane method. *Mathematical Programming*, 160(1–2):321–351, 2016. doi: 10.1007/s10107-016-0985-7.
- [44] Leon Eifler and Ambros Gleixner. A computational status update for exact rational mixed integer programming. *Mathematical Programming*, 197(2):793–812, 2023. doi: 10.1007/s10107-021-01749-5.
- [45] Leon Eifler and Ambros Gleixner. Safe and verified Gomory mixed-integer cuts in a rational mixed-integer program framework. *SIAM Journal on Optimization*, 34(1):742–763, 2024. doi: 10.1137/23M156046X.
- [46] James E. Falk. Lagrange multipliers and nonconvex programs. *SIAM Journal on Control*, 7(4):534–545, 1969. doi: 10.1137/0307039.
- [47] Giovanni Fantuzzi and Federico Fuentes. Finite convergence and minimizer extraction in moment relaxations with correlative sparsity. arXiv:2502.01410, 2025. URL <https://arxiv.org/abs/2502.01410>.
- [48] Stefan Feltenmark and Krzysztof C. Kiwiel. Dual applications of proximal bundle methods, including Lagrangian relaxation of nonconvex problems. *SIAM Journal on Optimization*, 10(3):697–721, 2000. doi: 10.1137/S1052623498332336.
- [49] Michael R. Garey, David S. Johnson, and Larry Stockmeyer. Some simplified NP-complete graph problems. *Theoretical Computer Science*, 1(3):237–267, 1976. doi: 10.1016/0304-3975(76)90059-1.

- [50] Jürgen Garloff and Andrew P. Smith. Rigorous affine lower bound functions for multivariate polynomials and their use in global optimisation. In *Proceedings of the 1st International Conference on Applied Operational Research*, volume 1 of *Lecture Notes in Management Science*, pages 199–211, 2008. URL <https://www-home.htwg-konstanz.de/~garloff/rigorous.pdf>.
- [51] Jürgen Garloff, Christian Jansson, and Andrew P. Smith. Lower bound functions for polynomials. *Journal of Computational and Applied Mathematics*, 157(1):207–225, 2003. doi: 10.1016/S0377-0427(03)00422-9.
- [52] Jürgen Garloff, Christian Jansson, and Andrew Paul Smith. Inclusion isotonicity of convex-concave extensions for polynomials based on Bernstein expansion. *Computing*, 70(2):111–119, 2003. doi: 10.1007/s00607-003-1471-7.
- [53] Arthur M. Geoffrion. Lagrangean relaxation for integer programming. In *Approaches to Integer Programming*, volume 2 of *Mathematical Programming Studies*, pages 82–114. Springer, 1974. doi: 10.1007/BFb0120690.
- [54] Elmer G. Gilbert. An iterative procedure for computing the minimum of a quadratic form on a convex set. *SIAM Journal on Control*, 4(1):61–80, 1966. doi: 10.1137/0304007.
- [55] Ambros M. Gleixner, Timo Berthold, Benjamin Müller, and Stefan Weltge. Three enhancements for optimization-based bound tightening. *Journal of Global Optimization*, 67(4):731–757, 2017. doi: 10.1007/s10898-016-0450-4.
- [56] David Grimm, Tim Netzer, and Markus Schweighofer. A note on the representation of positive polynomials with structured sparsity. *Archiv der Mathematik*, 89(5):399–403, 2007. doi: 10.1007/s00013-007-2234-z.
- [57] Robert Grone, Charles R. Johnson, Eduardo M. Sá, and Henry Wolkowicz. Positive definite completions of partial Hermitian matrices. *Linear Algebra and its Applications*, 58:109–124, 1984. doi: 10.1016/0024-3795(84)90207-6.
- [58] Martin Grötschel, László Lovász, and Alexander Schrijver. The ellipsoid method and its consequences in combinatorial optimization. *Combinatorica*, 1(2):169–197, 1981. doi: 10.1007/BF02579273.
- [59] Martin Grötschel, László Lovász, and Alexander Schrijver. *Geometric Algorithms and Combinatorial Optimization*, volume 2 of *Algorithms and Combinatorics*. Springer, Berlin, 1988. doi: 10.1007/978-3-642-97881-4.
- [60] Monique Guignard and Siwhan Kim. Lagrangean decomposition: a model yielding stronger Lagrangean bounds. *Mathematical Programming*, 39(2):215–228, 1987. doi: 10.1007/BF02592954.
- [61] Gurobi Optimization, LLC. *Gurobi Optimizer Reference Manual, Version 13.0*, 2026. URL <https://docs.gurobi.com/projects/optimizer/en/13.0/>.
- [62] Katrin Halbig, Lukas Hümb, Florian Rösel, Lars Schewe, and Dieter Weninger. Computing optimality certificates for convex mixed-integer nonlinear problems. *INFORMS Journal on Computing*, 36(6):1579–1610, 2024. doi: 10.1287/ijoc.2022.0099.
- [63] Pierre Hansen, Brigitte Jaumard, Michèle Ruiz, and Junjie Xiong. Global minimization of indefinite quadratic functions subject to box constraints. *Naval Research Logistics*, 40(3):373–392, 1993. doi: 10.1002/1520-6750(199304)40:3<373::AID-NAV3220400307>3.0.CO;2-A.

- [64] Taotao He and Mohit Tawarmalani. A new framework to relax composite functions in nonlinear programs. *Mathematical Programming*, 190(1–2):427–466, 2021. doi: 10.1007/s10107-020-01541-x.
- [65] Taotao He and Mohit Tawarmalani. Tractable relaxations of composite functions. *Mathematics of Operations Research*, 47(2):1110–1140, 2022. doi: 10.1287/moor.2021.1162.
- [66] Taotao He and Mohit Tawarmalani. MIP relaxations in factorable programming. *SIAM Journal on Optimization*, 34(3):2856–2882, 2024. doi: 10.1137/22M1515537.
- [67] Taotao He, Siyue Liu, and Mohit Tawarmalani. Convexification techniques for fractional programs. *Mathematical Programming*, 213(1–2):107–149, 2025. doi: 10.1007/s10107-024-02131-x.
- [68] Christopher Hojny, Mathieu Besançon, Ksenia Bestuzheva, Sander Borst, João Dionísio, Johannes Ehls, Leon Eifler, Mohammed Ghannam, Ambros Gleixner, Adrian Göß, Alexander Hoen, Jacob von Holly-Ponientzietz, Rolf van der Hulst, Dominik Kamp, Thorsten Koch, Kevin Kofler, Jurgen Lentz, Marco Lübbecke, Stephen J. Maher, Paul Matti Meinhold, Gioni Mexi, Til Mohr, Erik Mühmer, Krunal Kishor Patel, Marc E. Pfetsch, Sebastian Pokutta, Chantal Reinartz Groba, Felipe Serrano, Yuji Shinano, Mark Turner, Stefan Vigerske, Matthias Walter, Dieter Weninger, and Liding Xu. The SCIP Optimization Suite 10.0. arXiv:2511.18580, 2025. URL <https://arxiv.org/abs/2511.18580>.
- [69] Q. Huangfu and J. A. J. Hall. Parallelizing the dual revised simplex method. *Mathematical Programming Computation*, 10(1):119–142, 2018. doi: 10.1007/s12532-017-0130-5.
- [70] Fredrik Johansson. Arb: Efficient arbitrary-precision midpoint-radius interval arithmetic. *IEEE Transactions on Computers*, 66(8):1281–1292, 2017. doi: 10.1109/TC.2017.2690633.
- [71] Samuel Karlin and William J. Studden. *Tchebycheff Systems: With Applications in Analysis and Statistics*. Interscience, New York, 1966.
- [72] Richard M. Karp. Reducibility among combinatorial problems. In Raymond E. Miller and James W. Thatcher, editors, *Complexity of Computer Computations*, pages 85–103. Plenum Press, 1972. doi: 10.1007/978-1-4684-2001-2_9.
- [73] Ramkumar Karuppiah and Ignacio E. Grossmann. A Lagrangean based branch-and-cut algorithm for global optimization of nonconvex mixed-integer nonlinear programs with decomposable structures. *Journal of Global Optimization*, 41(2):163–186, 2008. doi: 10.1007/s10898-007-9203-8.
- [74] R. Baker Kearfott. Interval computations, rigour and non-rigour in deterministic continuous global optimization. *Optimization Methods and Software*, 26(2):259–279, 2011. doi: 10.1080/10556781003636851.
- [75] J. E. Kelley, Jr. The cutting-plane method for solving convex programs. *Journal of the Society for Industrial and Applied Mathematics*, 8(4):703–712, 1960. doi: 10.1137/0108053.
- [76] Aida Khajavirad. Tight semidefinite programming relaxations for sparse box-constrained quadratic programs. arXiv:2601.18545, 2026. URL <https://arxiv.org/abs/2601.18545>.

- [77] Sunyoung Kim, Masakazu Kojima, and Kim-Chuan Toh. Doubly nonnegative relaxations are equivalent to completely positive reformulations of quadratic optimization problems with block-clique graph structures. *Journal of Global Optimization*, 77(3):513–541, 2020. doi: 10.1007/s10898-020-00879-y.
- [78] Masakazu Kojima, Sunyoung Kim, and Naohiko Arima. Local-to-global exactness of SDP relaxations for sparse QCQPs. arXiv:2606.21823, 2026. URL <https://arxiv.org/abs/2606.21823>.
- [79] Dieter Kraft. A software package for sequential quadratic programming. Technical Report DFVLR-FB 88-28, DFVLR, Institut für Dynamik der Flugsysteme, Oberpfaffenhofen, 1988.
- [80] Jean B. Lasserre. Convergent SDP-relaxations in polynomial optimization with sparsity. *SIAM Journal on Optimization*, 17(3):822–843, 2006. doi: 10.1137/05064504X.
- [81] Monique Laurent. Sums of squares, moment matrices and optimization over polynomials. In Mihai Putinar and Seth Sullivant, editors, *Emerging Applications of Algebraic Geometry*, volume 149 of *The IMA Volumes in Mathematics and its Applications*, pages 157–270. Springer, New York, 2009. doi: 10.1007/978-0-387-09686-5_7.
- [82] Claude Lemaréchal and Arnaud Renaud. A geometric study of duality gaps, with applications. *Mathematical Programming*, 90(3):399–427, 2001. doi: 10.1007/PL00011429.
- [83] Frauke Liers, Alexander Martin, Maximilian Merkert, Nick Mertens, and Dennis Michaels. Solving mixed-integer nonlinear optimization problems using simultaneous convexification: a case study for gas networks. *Journal of Global Optimization*, 80(2):307–340, 2021. doi: 10.1007/s10898-020-00974-0.
- [84] Marco Locatelli and Fabio Schoen. On convex envelopes for bivariate functions over polytopes. *Mathematical Programming*, 144(1–2):65–91, 2014. doi: 10.1007/s10107-012-0616-x.
- [85] Andrea Lodi and Andrea Tramontani. Performance variability in mixed-integer programming. In *Theory Driven by Influential Applications*, INFORMS TutORials in Operations Research, pages 1–12. INFORMS, 2013. doi: 10.1287/educ.2013.0112.
- [86] James Luedtke, Mahdi Namazifar, and Jeff Linderoth. Some results on the strength of relaxations of multilinear functions. *Mathematical Programming*, 136(2):325–351, 2012. doi: 10.1007/s10107-012-0606-z.
- [87] Stephen Maher, Matthias Miltenberger, João Pedro Pedroso, Daniel Rehfeldt, Robert Schwarz, and Felipe Serrano. PySCIPOpt: Mathematical programming in Python with the SCIP Optimization Suite. In *Mathematical Software – ICMS 2016*, volume 9725 of *Lecture Notes in Computer Science*, pages 301–307. Springer, 2016. doi: 10.1007/978-3-319-42432-3_37.
- [88] O. L. Mangasarian. Arbitrary-norm separating plane. *Operations Research Letters*, 24(1–2): 15–23, 1999. doi: 10.1016/S0167-6377(98)00049-2.
- [89] François Margot. Testing cut generators for mixed-integer linear programming. *Mathematical Programming Computation*, 1(1):69–95, 2009. doi: 10.1007/s12532-009-0003-7.
- [90] Garth P. McCormick. Computability of global solutions to factorable nonconvex programs: Part I – Convex underestimating problems. *Mathematical Programming*, 10(1):147–175, 1976. doi: 10.1007/BF01580665.

- [91] Ruth Misener and Christodoulos A. Floudas. Global optimization of mixed-integer quadratically-constrained quadratic programs (MIQCQP) through piecewise-linear and edge-concave relaxations. *Mathematical Programming*, 136(1):155–182, 2012. doi: 10.1007/s10107-012-0555-6.
- [92] Jorge J. Moré and Stephen A. Vavasis. On the solution of concave knapsack problems. *Mathematical Programming*, 49(1–3):397–411, 1990. doi: 10.1007/BF01588800.
- [93] Benjamin Müller, Felipe Serrano, and Ambros Gleixner. Using two-dimensional projections for stronger separation and propagation of bilinear terms. *SIAM Journal on Optimization*, 30(2):1339–1365, 2020. doi: 10.1137/19M1249825.
- [94] Benjamin Müller, Gonzalo Muñoz, Maxime Gasse, Ambros Gleixner, Andrea Lodi, and Felipe Serrano. On generalized surrogate duality in mixed-integer nonlinear programming. *Mathematical Programming*, 192(1–2):89–118, 2022. doi: 10.1007/s10107-021-01691-6.
- [95] César Muñoz and Anthony Narkawicz. Formalization of Bernstein polynomials and applications to global optimization. *Journal of Automated Reasoning*, 51(2):151–196, 2013. doi: 10.1007/s10817-012-9256-3.
- [96] Gonzalo Muñoz and Felipe Serrano. Maximal quadratic-free sets. *Mathematical Programming*, 192(1–2):229–270, 2022. doi: 10.1007/s10107-021-01738-8.
- [97] Katta G. Murty. *Linear Complementarity, Linear and Nonlinear Programming*. University of Michigan, internet edition, 1997. URL https://public.websites.umich.edu/~murty/books/linear_complementarity_webbook/. Internet edition; first published by Heldermann, Berlin, 1988.
- [98] Jaromil Najman, Dominik Bongartz, and Alexander Mitsos. Linearization of McCormick relaxations and hybridization with the auxiliary variable method. *Journal of Global Optimization*, 80(4):731–756, 2021. doi: 10.1007/s10898-020-00977-x.
- [99] George C. Necula. Translation validation for an optimizing compiler. In *Proceedings of PLDI 2000*, pages 83–94, 2000. doi: 10.1145/349299.349314.
- [100] G. L. Nemhauser and L. E. Trotter. Properties of vertex packing and independence system polyhedra. *Mathematical Programming*, 6(1):48–61, 1974. doi: 10.1007/BF01580222.
- [101] Arnold Neumaier. Complete search in continuous global optimization and constraint satisfaction. *Acta Numerica*, 13:271–369, 2004. doi: 10.1017/S0962492904000194.
- [102] Arnold Neumaier and Oleg Shcherbina. Safe bounds in linear and mixed-integer linear programming. *Mathematical Programming*, 99(2):283–296, 2004. doi: 10.1007/s10107-003-0433-3.
- [103] Jiawang Nie and James Demmel. Sparse SOS relaxations for minimizing functions that are summations of small polynomials. *SIAM Journal on Optimization*, 19(4):1534–1558, 2009. doi: 10.1137/060668791.
- [104] Jiawang Nie, Zheng Qu, Xindong Tang, and Linghao Zhang. A characterization for tightness of the sparse moment-SOS hierarchy. *Mathematical Programming*, 215(1–2):369–405, 2026. doi: 10.1007/s10107-025-02223-2.

- [105] Jordan Ninin, Frédéric Messine, and Pierre Hansen. A reliable affine relaxation method for global optimization. *4OR*, 13(3):247–277, 2015. doi: 10.1007/s10288-014-0269-0.
- [106] Ivo Nowak. *Relaxation and Decomposition Methods for Mixed Integer Nonlinear Programming*, volume 152 of *International Series of Numerical Mathematics*. Birkhäuser, Basel, 2005. ISBN 978-3-7643-7238-5. doi: 10.1007/3-7643-7374-1.
- [107] Manfred Padberg. The boolean quadric polytope: some characteristics, facets and relatives. *Mathematical Programming*, 45(1–3):139–172, 1989. doi: 10.1007/BF01589101.
- [108] Amir Pnueli, Michael Siegel, and Eli Singerman. Translation validation. In *Tools and Algorithms for the Construction and Analysis of Systems (TACAS 1998)*, volume 1384 of *Lecture Notes in Computer Science*, pages 151–166, 1998. doi: 10.1007/BFb0054170.
- [109] J. L. Rabinowitsch. Zum Hilbertschen Nullstellensatz. *Mathematische Annalen*, 102(1):520, 1930. doi: 10.1007/BF01782361.
- [110] Anatoliy D. Rikun. A convex envelope formula for multilinear functions. *Journal of Global Optimization*, 10(4):425–437, 1997. doi: 10.1023/A:1008217604285.
- [111] Anureet Saxena, Pierre Bonami, and Jon Lee. Convex relaxations of non-convex mixed integer quadratically constrained programs: projected formulations. *Mathematical Programming*, 130(2):359–413, 2011. doi: 10.1007/s10107-010-0340-3.
- [112] Hermann Schichl and Arnold Neumaier. Interval analysis on directed acyclic graphs for global optimization. *Journal of Global Optimization*, 33(4):541–562, 2005. doi: 10.1007/s10898-005-0937-x.
- [113] Hanif D. Sherali and Warren P. Adams. A hierarchy of relaxations between the continuous and convex hull representations for zero-one programming problems. *SIAM Journal on Discrete Mathematics*, 3(3):411–430, 1990. doi: 10.1137/0403036.
- [114] Hanif D. Sherali and Cihan H. Tuncbilek. A global optimization algorithm for polynomial programming problems using a Reformulation-Linearization Technique. *Journal of Global Optimization*, 2(1):101–112, 1992. doi: 10.1007/BF00121304.
- [115] N. Z. Shor. Quadratic optimization problems. *Soviet Journal of Computer and Systems Sciences*, 25:1–11, 1987.
- [116] Maurice Sion. On general minimax theorems. *Pacific Journal of Mathematics*, 8(1):171–176, 1958. doi: 10.2140/pjm.1958.8.171.
- [117] Edward M. B. Smith and Constantinos C. Pantelides. A symbolic reformulation/spatial branch-and-bound algorithm for the global optimisation of nonconvex MINLPs. *Computers & Chemical Engineering*, 23(4–5):457–478, 1999. doi: 10.1016/S0098-1354(98)00286-5.
- [118] Mohit Tawarmalani. Inclusion certificates and simultaneous convexification of functions. Optimization Online preprint 2722, 2010. URL <https://optimization-online.org/wp-content/uploads/2010/09/2722.pdf>.
- [119] Mohit Tawarmalani and Nikolaos V. Sahinidis. Global optimization of mixed-integer nonlinear programs: A theoretical and computational study. *Mathematical Programming*, 99(3):563–591, 2004. doi: 10.1007/s10107-003-0467-6.

- [120] Mohit Tawarmalani and Nikolaos V. Sahinidis. A polyhedral branch-and-cut approach to global optimization. *Mathematical Programming*, 103(2):225–249, 2005. doi: 10.1007/s10107-005-0581-8.
- [121] Mohit Tawarmalani, Jean-Philippe P. Richard, and Chuanhui Xiong. Explicit convex and concave envelopes through polyhedral subdivisions. *Mathematical Programming*, 138(1–2): 531–577, 2013. doi: 10.1007/s10107-012-0581-4.
- [122] The SCIP Optimization Suite developers. SCIP source code, version 10.0.2, file src/scip/lp.c, 2026. URL <https://github.com/scipopt/scip/tree/v10.0.2>.
- [123] Angelos Tsoukalas and Alexander Mitsos. Multivariate McCormick relaxations. *Journal of Global Optimization*, 59(2–3):633–662, 2014. doi: 10.1007/s10898-014-0176-0.
- [124] Stephen A. Vavasis. Quadratic programming is in NP. *Information Processing Letters*, 36(2): 73–77, 1990. doi: 10.1016/0020-0190(90)90100-C.
- [125] Stefan Vigerske. MINLPLib: A library of mixed-integer and continuous nonlinear programming instances, 2026. URL <https://www.minlplib.org>. Accessed 2026-10.
- [126] Nikolai N. Vorob’ev. Consistent families of measures and their extensions. *Theory of Probability and Its Applications*, 7(2):147–163, 1962. doi: 10.1137/1107014.
- [127] Hayato Waki, Sunyoung Kim, Masakazu Kojima, and Masakazu Muramatsu. Sums of squares and semidefinite program relaxations for polynomial optimization problems with structured sparsity. *SIAM Journal on Optimization*, 17(1):218–242, 2006. doi: 10.1137/050623802.
- [128] Ouyang Wu, Pavlo Muts, Ivo Nowak, and Eligius M. T. Hendrix. On the use of overlapping convex hull relaxations to solve nonconvex MINLPs. *Journal of Global Optimization*, 91(2): 415–436, 2025. doi: 10.1007/s10898-024-01376-2.
- [129] Liding Xu and Sebastian Pokutta. Joint-range inequalities for nonconvex QCQPs. arXiv:2608.03318, 2026. URL <https://arxiv.org/abs/2608.03318>.
- [130] Haisheng Zhu, Taotao He, and Mohit Tawarmalani. Axis-aligned relaxations for mixed-integer nonlinear programming. arXiv:2603.18458, 2026. URL <https://arxiv.org/abs/2603.18458>.

A Proof details for Section 3

A.1 Lower bounds for Theorem 3.2

Let ν_A and ν_C be probability measures on $[\ell, u]$ with equal moments of orders one and two. For any polynomial p of degree at most two, $\int \text{dist}(y, A)^2 d\nu_A + \int \text{dist}(y, C)^2 d\nu_C$ equals $\int (\text{dist}(y, A)^2 - p) d\nu_A + \int (\text{dist}(y, C)^2 + p) d\nu_C$, so it suffices to find p with $\text{dist}(\cdot, A)^2 - p \geq \eta_A$ and $\text{dist}(\cdot, C)^2 + p \geq \eta_C$ on $\mathbb{R}$ and $\eta_A + \eta_C = \delta^2/2$. Because $\text{dist}(y, S)^2 = \min_{s \in S} (y - s)^2$, each inequality can be checked one point s at a time. We use $\inf_y [\alpha(y - s)^2 + \beta(y - t)^2] = \alpha\beta(s - t)^2/(\alpha + \beta)$ for $\alpha + \beta > 0$.

Separated sets. By symmetry let $\max A < \min C$, put $a_0 = \max A$, $c_0 = \min C$, $\delta = c_0 - a_0$, $t_0 = (a_0 + c_0)/2$ and $p(y) = \delta(y - t_0)$. For $s \in A$, $\inf_y [(y - s)^2 - \delta(y - t_0)] = -\delta^2/4 - \delta(s - t_0) \geq \delta^2/4$, and for $t \in C$, $\inf_y [(y - t)^2 + \delta(y - t_0)] = -\delta^2/4 + \delta(t - t_0) \geq \delta^2/4$.

Nested sets. By symmetry let C lie strictly between the two points $a_- < a_+$ of A . Put $t_0 = (\min C + \max C)/2$, $h = (\max C - \min C)/2$ and $\rho = \delta + h$; then $|s - t_0| \geq \rho$ for $s \in A$ and

$|t - t_0| \leq h$ for $t \in C$. If $h > 0$, let $\gamma = (\rho - h)/(\rho + h) \in (0, 1)$ and $p(y) = -\gamma(y - t_0)^2$. Then $\inf_y[(y - s)^2 + \gamma(y - t_0)^2] = \gamma(s - t_0)^2/(1 + \gamma) \geq \gamma\rho^2/(1 + \gamma)$ and $\inf_y[(y - t)^2 - \gamma(y - t_0)^2] = -\gamma(t - t_0)^2/(1 - \gamma) \geq -\gamma h^2/(1 - \gamma)$. With $1 + \gamma = 2\rho/(\rho + h)$ and $1 - \gamma = 2h/(\rho + h)$ the two bounds add up to $(\rho - h)^2/2 = \delta^2/2$. If $h = 0$, take $p(y) = -(y - t_0)^2$; then $\text{dist}(\cdot, C)^2 + p \equiv 0$ and $\text{dist}(\cdot, A)^2 - p \geq \delta^2/2$.

A.2 A separating polynomial for Theorem 3.3

Let A and C be disjoint, nonempty and compact, let $T = A \cup C$, and define the label $\epsilon(t) = +1$ for $t \in A$ and $\epsilon(t) = -1$ for $t \in C$. Let $\varepsilon = \text{dist}(A, C) > 0$. Starting from $s_0 = \min T$, let s_{i+1} be the smallest point of T larger than s_i with the opposite label, as long as one exists. Such a point exceeds s_i by at least ε , so the process stops at some s_r , and every point of T in $[s_i, s_{i+1})$ has the label of s_i . Every alternating chain visits these runs in increasing order, so $\text{alt}(A, C) = r$. Let e_i be the largest point of T in $[s_{i-1}, s_i)$ and choose $\rho_i \in (e_i, s_i)$. Then $p(t) = \epsilon(s_0) \prod_{i=1}^r (\rho_i - t)$ has degree r , and a point of the i -th run lies between ρ_i and ρ_{i+1} , so $\text{sign } p(t) = \epsilon(t)$: $p > 0$ on A and $p < 0$ on C . If $r \leq \kappa$, then $\int p d\mu_A = \int p d\mu_C$ for any probability measures μ_A on A and μ_C on C with equal moments up to order κ , which is impossible.

A.3 Dense moment matrices for the path family

We use the notation of Theorem 3.2, assume that A and C interleave, and let $v \in R$ be a point with $\text{lin } \Phi(v) = 0$. Let M be the moment matrix of $(1, x, y, z)$ with the entries of v and an unspecified entry p_{xz} .

The dense moment matrix alone does not exclude v . The specified entries of M form a chordal pattern whose two maximal blocks, the moment matrices of $(1, x, y)$ and $(1, y, z)$, are positive semidefinite because the two parts of v are moment vectors. Hence M has a positive semidefinite completion [57], and v together with this completion satisfies the dense moment constraint and all RLT inequalities of the edges and squares.

The McCormick inequalities of xz alone do not exclude v . They require $\max\{0, m_x + m_z - 1\} \leq p_{xz} \leq \min\{m_x, m_z\}$, an interval that is nonempty for $m_x, m_z \in [0, 1]$.

Together they exclude it. Replacing x by $1 - x$ exchanges a_1 and a_2 and maps the four McCormick inequalities of xz to each other, and similarly for z , so we may assume $a_1 < a_2$ and $c_1 < c_2$. The (x, y) part of v is the moment vector of a measure supported on the zero set of Φ_A , which is $\{(0, a_1), (1, a_2)\}$: at a point with $0 < x < 1$, concavity or affinity in x gives $\Phi_A(x, y) \geq (1 - x)(y - a_1)^2 + x(y - a_2)^2 > 0$. Hence $x = \alpha + \beta y$ on this support, with $\beta = 1/(a_2 - a_1)$ and $\alpha = -a_1\beta$, and the vector $w = (\alpha, -1, \beta, 0)$ satisfies $w^\top M w = 0$. If $M \succeq 0$, then $M w = 0$, and the last row gives $p_{xz} = \alpha m_z + \beta p_{yz}$; so the completion is unique. In the same way $z = (y - c_1)/(c_2 - c_1)$ on the (y, z) support, which gives $m_z = (m_y - c_1)/(c_2 - c_1)$ and $p_{yz} = (s_y - c_1 m_y)/(c_2 - c_1)$. Combining,

$$p_{xz} = \frac{s_y - (a_1 + c_1)m_y + a_1 c_1}{(a_2 - a_1)(c_2 - c_1)}.$$

The numerator is the expectation of $(y - a_1)(y - c_1)$ under any measure with moments (m_y, s_y) . Evaluating it with the (x, y) measure, which puts mass m_x on $y = a_2$, gives $p_{xz} = m_x(a_2 - c_1)/(c_2 - c_1)$; evaluating it with the (y, z) measure gives $p_{xz} = m_z(c_2 - a_1)/(a_2 - a_1)$. Since the chords meet at a point interior to both, $0 < m_x < 1$ and $0 < m_z < 1$. If $a_1 < c_1 < a_2 < c_2$, then $(c_2 - a_1)/(a_2 - a_1) > 1$ and $p_{xz} > m_z$. If $c_1 < a_1 < c_2 < a_2$, then $(a_2 - c_1)/(c_2 - c_1) > 1$ and $p_{xz} > m_x$. Either way a McCormick inequality is violated. More strongly, the two ingredients prove $\Phi \geq \delta^2/2$ through (7):

write $\xi_A^2 + \xi_C^2 = \frac{1}{2}(\xi_A + \xi_C)^2 + \frac{1}{2}(\xi_A - \xi_C)^2$; the function $g = \frac{1}{2}(\xi_A - \xi_C)^2 + \frac{1}{2}(a_2 - a_1)^2 x(1 - x) + \frac{1}{2}(c_2 - c_1)^2 z(1 - z)$ has no x^2 or z^2 terms, so it is bilinear in (x, z) and equals its interpolation $\sum_{ij} g(i, j) \ell_{ij}$ from the four vertices, where $g(i, j) = \frac{1}{2}(c_{j+1} - a_{i+1})^2$; and $\sum_{ij} \ell_{ij} = 1$.

A.4 Examples for Proposition 3.5

With $q_1 = \Phi_A$ and $q_2 = \Phi_C$ for the nested sets $A = \{0, 3\}$ and $C = \{1, 2\}$ on $y \in [0, 3]$, both pair minima are zero and $\Delta = \beta^* = 1/2$; the re-splitting $q_1 + \frac{1}{2}(y - \frac{3}{2})^2$ and $q_2 - \frac{1}{2}(y - \frac{3}{2})^2$ has pair minima $\frac{3}{4}$ and $-\frac{1}{4}$, so $\Delta' = 0$, in agreement with the zero gap of Theorem 3.2. The polynomials p of Appendix A.1 are optimal re-splittings of this kind. Pairwise intersections of the projections in Proposition 3.5(i) do not suffice when $k \geq 3$: for three leaves of the form Φ_A with $A_1 = \{0, 1\}$, $A_2 = \{1, 2\}$ and $A_3 = \{0, 2\}$ on $y \in [0, 2]$, the projections meet pairwise but have no common point, and $\Delta = 2/3$.

B Exact quadratic support on polytopes

B.1 Proof of Theorem 4.1

Every candidate lies in P , so an empty P has none. Assume $P \neq \emptyset$. Among the global minimizers, which exist by compactness, choose x^* such that the face G of P with $x^* \in \text{relint } G$ has the smallest dimension, and let V be the linear space parallel to $\text{aff } G$. Because x^* minimizes q over G and lies in its relative interior, $v^\top (Hx^* + c) = 0$ and $v^\top H v \geq 0$ for all $v \in V$.

Suppose $v \in V \setminus \{0\}$ satisfies $v^\top H v = 0$. Then $q(x^* + tv) = q(x^*)$ for all t . The line $x^* + tv$ meets the bounded face G in a segment that contains x^* in its relative interior, and an endpoint of the segment is a global minimizer on a proper face of G , which contradicts the choice of x^* . Hence H is positive definite on V , or $V = \{0\}$.

Let J be the set of rows active at x^* , so that $\text{aff } G = \{x : A_J x = b_J\}$ and $V = \ker A_J$, and let $I \subseteq J$ index a basis of the row space of A_J . Then $|I| \leq d$ and $\ker A_I = V$. Since $Hx^* + c$ is orthogonal to V , it lies in the row space of A_I , so (x^*, μ) solves (8) for some μ . If $M_I(v, \nu) = 0$, then $A_I v = 0$, so $v \in V$, and $v^\top H v = -v^\top A_I^\top \nu = 0$; hence $v = 0$, and $A_I^\top \nu = 0$ gives $\nu = 0$ by independence. So M_I is nonsingular, its unique solution contains x^* , and x^* is a candidate. All data are rational, so the candidates and their values are rational. A set of dependent rows never gives a nonsingular M_I , so enumerating independent sets suffices. $\square$

B.2 Encoding lengths

Multiply each row of $Ax \leq b$ by the product of the denominators of its entries, and multiply q by the product σ of the denominators of its coefficients. Positive row scalings and $\sigma > 0$ change neither P , nor the nonsingularity of M_I , nor the x -part of its solution. Afterwards all row entries are integers of absolute value at most $2^{\langle A \rangle}$ and all entries of σH and σc are integers of absolute value at most $2^{\langle q \rangle + 1}$; write $\alpha = \langle A \rangle$ and $\eta = \langle q \rangle + 1$.

Let $|I| = s \leq d$ and $n = d + s$. Scale the last s rows and the last s columns of the integer matrix $[M_I \mid (-\sigma c, b_I)]$ by $2^{-\alpha}$. In the scaled matrix, the first d rows have entries of absolute value at most 2^η and the others at most 1. Hadamard's inequality applied to any square submatrix of order at most n , and the undoing of the scaling, give

$$|\det N| \leq n^{n/2} 2^{d\eta + 2s\alpha} \leq B := (2d)^d 2^{d(\eta + 2\alpha)}$$

for every square submatrix N of $[M_I \mid (-\sigma c, b_I)]$. By Cramer's rule a candidate has coordinates $x_j = N_j/\Lambda$ with $|N_j|, |\Lambda| \leq B$. Writing $\sigma q = \tilde{c}_0 + \sum_j \tilde{c}_j x_j + \sum_{i \leq j} \tilde{Q}_{ij} x_i x_j$ with integer coefficients bounded by 2^η ,

$$\sigma q(x) = \frac{\Lambda^2 \tilde{c}_0 + \Lambda \sum_j \tilde{c}_j N_j + \sum_{i \leq j} \tilde{Q}_{ij} N_i N_j}{\Lambda^2},$$

a numerator with at most $(d+1)^2$ terms, each bounded by $2^\eta B^2$. So $\min_P q$ is a fraction whose numerator and denominator are bounded by $(d+1)^2 2^\eta B^2$ and $2^\eta B^2$. Hence every candidate and the minimum have bit length $O(d(\langle q \rangle + \langle A \rangle + \log d))$, for every d .

Bareiss's fraction-free elimination keeps every intermediate number equal to a minor of the matrix above, and the products formed before each exact division are bounded by $2B^2$. Feasibility tests and value comparisons are integer comparisons of numbers bounded by $(d+1)2^\alpha B$ and $(d+1)^2 2^\eta B^4$. The total work is $O(N(m, d)(d^3 + md))$ arithmetic operations on integers of bit length $O(d(\langle q \rangle + \langle A \rangle + \log d))$. For fixed d , $N(m, d) \leq (m+1)^d$, the number of arithmetic operations does not depend on the sizes of the numbers, and the algorithm is strongly polynomial.

B.3 Proof of Proposition 4.3

For a graph $G = (V, E)$ let $q(x) = -\sum_{\{i,j\} \in E} (x_i + x_j - 2x_i x_j)$. Since q is affine in each coordinate, moving coordinates one at a time to a minimizing endpoint does not increase q , so some minimizer is binary. At a binary point the summand of an edge equals one exactly when the edge crosses the cut defined by x , so $\min q$ is minus the maximum cut size, and unweighted MAX CUT is NP-hard [49, 72]. The coefficients are integers of absolute value at most n . For the second statement, membership in coNP follows because a violating point can be taken to be a candidate of Theorem 4.1, which has polynomial size. For the hardness, $\min q$ is an integer for the quadratic q of the reduction, because q has integer coefficients and a binary minimizer; hence $\min q \leq -W$ holds if and only if $q \geq -W + \frac{1}{2}$ fails on $[0, 1]^n$. $\square$

C Sizes, a lower bound and a depth-two example for stars

For a rational $r = p/q$ in lowest terms let $\text{ht}(r) = \log_2 \max\{|p|, |q|\}$. Then $\text{ht}(rs), \text{ht}(r/s) \leq \text{ht}(r) + \text{ht}(s)$ and $\text{ht}(r_1 + \dots + r_N) \leq \log_2 N + \sum_j \text{ht}(r_j)$. Let τ_i bound ht over the data of leaf i (its coefficients, bounds and rows), τ_0 over the center data (its coefficients, bounds and the rows in y alone), and $\tau = \max_{0 \leq i \leq k} \tau_i$.

Sizes. The coefficients of the affine functions defining L_i and U_i have $\text{ht} \leq 2\tau_i$, and v_i has $\text{ht} \leq 2\tau_i + 1$. A breakpoint is the intersection of two such functions or the zero of the affine factor in (10); in both cases $\text{ht} \leq 12\tau + 4$, independently of k . The endpoints of I are bounds of y , zeros of rows in y alone ($\text{ht} \leq 2\tau_0$) or intersections of two lines of one leaf, so they obey the same bound. On a piece, the contribution of leaf i to V is a quadratic in y whose coefficients have $\text{ht} \leq 11\tau_i + 7$, so the coefficients of V , which are sums of $k+1$ such terms, have $\text{ht} \leq C := \log_2(k+1) + \tau_0 + \sum_i (11\tau_i + 7)$, which is linear in the input length. Evaluating at a breakpoint, at an endpoint of I or at a stationary point $-c_1/(2c_2)$ gives a minimum with $\text{ht} \leq 4C + 36\tau + 14$. Because the sweep updates the coefficients exactly, every running sum equals the coefficient of the current piece and obeys the same bound. The algorithm is therefore strongly polynomial.

Lower bound. Ben-Or's theorem states that an algebraic computation tree deciding membership in a set $W \subseteq \mathbb{R}^n$ with N connected components has depth $\Omega(\log N - n)$ [10]. For $r \geq 1$ let

$W = \{z \in (0, 2r)^r : |z_i - z_j| \geq 1 \text{ for } i \neq j\}$, which has one component for each of the $r!$ orderings. For $z \in (0, 2r)^r$ build a star with center $y \in [0, 2r]$ and, for each j , three leaves: $x \in [0, 1]$ with objective $(y - z_j + 1)x$; $x \in [0, 1]$ with objective $(z_j + 1 - y)x$; and $x \in [0, 2r + 1]$ with the row $x \geq y - z_j$ and objective $2x$; add the center terms $-ry + \sum_j z_j - r$. With $t = y - z_j$, the identity $\min\{0, t + 1\} + \min\{0, 1 - t\} + 2\max\{0, t\} - t - 1 = -\max\{0, 1 - |t|\}$ shows that $V(y) = -\sum_j \max\{0, 1 - |y - z_j|\}$, a sum of negative tents of width two. If $z \in W$, at most two tents are positive at any y and their sum is at most one, so $\min V \geq -1$; if $|z_i - z_j| < 1$, then $V < -1$ at the midpoint of z_i and z_j . The instance has $3r$ leaves and r leaf rows and is computed from z with $O(r)$ operations, so deciding $\min V \geq -1$ needs depth $\Omega(r \log r)$.

Depth two. For $f = (x_2 - \frac{1}{8})x_1 + x_2^2 - x_2 + x_2x_3 + x_3^2 - x_3$ on $[0, 1]^3$, minimizing over x_1 first gives $\min\{0, x_2 - \frac{1}{8}\} + x_2^2 + (t - 1)x_2$ with $t = x_3$. On $[0, \frac{1}{8}]$ its minimum is $-\frac{1}{8}$ at $x_2 = 0$; on $[\frac{1}{8}, 1]$ it is $-\frac{1}{4}(1 - t)^2$ at $x_2 = (1 - t)/2$ when $t \leq \frac{3}{4}$, and larger than $-\frac{1}{8}$ otherwise. Hence the conditional value is $\min\{-\frac{1}{8}, -\frac{1}{4}(1 - t)^2\}$, whose two pieces meet at $t = 1 - \sqrt{2}/2$, a root of $2t^2 - 4t + 1$. Adding $t^2 - t$ and minimizing over $t \in [0, 1]$ gives $\min f = -\frac{3}{8}$ at $(1, 0, \frac{1}{2})$.

D The closure and the hull of the feasible set

The closure $\mathcal{C}$ is a relaxation of the set $\Sigma = \{v : Sv \in D, v \text{ satisfies (11)}\}$, so $\text{conv } \Sigma \subseteq \mathcal{C}$. Equality holds when every row has its own free direction in the affine parts.

Proposition D.1 (Free affine columns). *Write $v = (x, z)$ with x the block, and let M be the matrix formed by the columns of A that belong to z , with one row for each inequality and one row for each equality (of the two opposite rows of an equality, M keeps one). If M has full row rank, then $\mathcal{C} = \text{conv } \Sigma$.*

Proof. Let $v = (x, z) \in \mathcal{C}$. By Carathéodory's theorem there are $u_j \in D$ and weights $\theta_j > 0$ summing to one with $\sum_j \theta_j u_j = x$ and $\bar{y} := \sum_j \theta_j g(u_j) \leq b - Av$, with equality in the rows of each equality. Let A_x denote the block columns of A , restricted like M , let M^+ be a right inverse of M , put $\rho_j = \bar{y} - g(u_j) + A_x(x - u_j)$ in the same rows, and $z_j = z + M^+ \rho_j$. A direct substitution shows that each row of (u_j, z_j) has the same left-hand side as $\bar{y} + Av$, so $(u_j, z_j) \in \Sigma$. Since $\sum_j \theta_j \rho_j = 0$, the convex combination of the points (u_j, z_j) is v . $\square$

The objective row always contains $-t$, so for the objective alone the closure is the hull of the epigraph over D . The setting of Liers et al. [83], rows $z_j = g_j(x)$ with a separate variable z_j for each row, is also covered by Proposition D.1. Without free affine columns the gap can be strict.

Example D.2 (Gap to the feasible hull). Let $D = [0, 1]$ and take the equality row $x^2 = 1/4$, whose feasible set is $\{1/2\}$. By Corollary 5.3, $\mathcal{C} = \{x : x^2 \leq 1/4 \leq x\} = [1/4, 1/2]$, and the two cuts $x \geq 1/4$ and $x \leq 1/2$ describe it. The point $x = 1/4$ is the mean of the measure with mass $3/4$ at 0 and $1/4$ at 1, which satisfies the row in expectation.

This is a Lagrangian duality gap [53, 106]: the rows are imposed on averages, after convexification. Even the complete family of cuts accepts $x = 1/4$, so the gap is not a failure of direction search. It is closed by convexifying aggregated sets instead of aggregated functions, as in the aggregation closure of Dey et al. [40]: here $\text{conv}\{x \in [0, 1] : x^2 \leq 1/4\} = [0, 1/2]$ and $\text{conv}\{x \in [0, 1] : x^2 \geq 1/4\} = [1/2, 1]$ intersect in $\{1/2\}$. It is also closed by placing the row inside D , which changes the support problem;

our implementation keeps nonlinear rows out of D so that the support problems stay in the classes of Section 4. With $\Sigma_\lambda = \{v : Sv \in D, \lambda^\top(g(Sv) + Av - b) \leq 0\}$ for $\lambda \geq 0$, the hierarchy is

$$\text{conv } \Sigma \subseteq \bigcap_{\lambda \geq 0} \text{conv } \Sigma_\lambda \subseteq \mathcal{C}.$$

The middle set is the convexified closure of the surrogate relaxations of Müller et al. [94], and the gap to $\mathcal{C}$ is a surrogate–Lagrangian gap. Example D.2 shows that the second inclusion can be strict. The first can be strict as well, even when D is exactly the projection of Σ : with $D = [0, 2]$, a free z and the rows $z \leq (x - 2)^2/2$ and $z \leq x^2/2$, the set Σ satisfies $z \leq \min\{x/2, 1 - x/2\}$ on its hull, so $z \leq 1/2$ at $x = 1$. Every aggregated row reads $z \leq q_\lambda(x)$ with $q_\lambda = [\lambda_1(x - 2)^2 + \lambda_2 x^2]/(2(\lambda_1 + \lambda_2))$ convex and $q_\lambda(0) + q_\lambda(2) = 2$, so the chord gives $(1, 1) \in \text{conv } \Sigma_\lambda$ for every λ , and $(1, 1)$, the mean of the point masses at $x = 0$ and $x = 2$, also lies in $\mathcal{C}$.

E Whole-domain bounds for non-quadratic blocks

Polynomials. Let f be a polynomial with rational coefficients on a rational box $B = \prod_{i=1}^n [\ell_i, u_i]$. After the affine change of variables $x_i = \ell_i + (u_i - \ell_i)t_i$ and removal of fixed coordinates, write

$$f(x(t)) = \sum_{\beta \leq d} f_\beta t^\beta = \sum_{\theta \leq d} b_\theta B_\theta^d(t), \quad B_\theta^d(t) = \prod_i \binom{d_i}{\theta_i} t_i^{\theta_i} (1 - t_i)^{d_i - \theta_i},$$

where d is a componentwise degree bound. The Bernstein coefficients are

$$b_\theta = \sum_{\beta \leq \theta} f_\beta \prod_i \binom{\theta_i}{\beta_i} / \binom{d_i}{\beta_i}. \quad (15)$$

Lemma E.1 (Bernstein enclosure; see, e.g., 50 and the references therein). *For every $x \in B$, $\min_\theta b_\theta \leq f(x) \leq \max_\theta b_\theta$.*

Proof. On $[0, 1]^n$ the basis polynomials are nonnegative and sum to one, so $f(x(t))$ is a convex combination of the coefficients. Formula (15) follows from the univariate identity $t^j = \sum_{i \geq j} \binom{i}{j} \binom{d}{j}^{-1} B_i^d(t)$ applied in each coordinate. $\square$

All quantities in (15) are rational, so the enclosure is computed exactly. A weak enclosure is improved by subdivision, and the enclosures shrink monotonically with the box [52]; exact rational Bernstein subdivision has been formalized in a proof assistant [95].

Proposition E.2 (Partition certificate). *Let $B_1, \dots, B_N$ be closed boxes with $\bigcup_j B_j = B$, and let $L_j \leq \inf_{B_j} f$ for each j . Then $\min_j L_j \leq \inf_B f$. Let $D = \{x \in B : g_r^\top x \leq \gamma_r, r \in \mathcal{R}\}$ for rational rows, and let O be the set of cells with $\min_{x \in B_j} g_r^\top x > \gamma_r$ for some r . Then $\min_{j \notin O} L_j \leq \inf_D f$, and if O contains every cell, $D = \emptyset$.*

The proof is immediate. Its hypothesis is what a checker must verify: the cells must cover B exactly, starting from the trusted box and not from a list supplied by the generator. A missing end interval or a child box generated from a different parent would invalidate the bound. For the same reason, a stored bound may be reused only for the same expression, direction and domain.

Univariate elementary expressions. For one variable we also admit sums, products and rational powers of $\exp$, $\log$, $\sin$, $\cos$, $\sinh$, $\cosh$, $\tanh$, $\arctan$ and $|\cdot|$, with rational constants and π and e . On every cell of an interval subdivision, outward-rounded ball arithmetic [70] gives an enclosure of the support objective $t \mapsto at + b^\top F(t)$. Domain requirements are part of the expression: $\log$ needs a positive argument, a quotient needs a denominator interval that excludes zero, and fractional powers use the nonnegative-base real convention. If an enclosure is not finite or a domain requirement cannot be proved on a cell, that cell is subdivided further, and when the depth budget is exhausted the computation is incomplete.

A curvature bound often gives a much better lower bound on a cell than a natural interval extension. The following bound is the classical remainder of linear interpolation; with $f'' \leq M$, the quantity $M(\beta - \alpha)^2/8$ is also the maximal separation of the α BB underestimator of $-f$ with parameter $M/2$ [1].

Lemma E.3 (Chord correction). *Let f be twice continuously differentiable on $[\alpha, \beta]$ with $f'' \leq M$. Then*

$$\min_{[\alpha, \beta]} f \geq \min\{f(\alpha), f(\beta)\} - \max\{0, M\}(\beta - \alpha)^2/8.$$

Proof. The function $f(t) - Mt^2/2$ is concave, so it lies above its chord. Rearranging gives $f(t) \geq \ell_f(t) - \frac{M}{2}(t - \alpha)(\beta - t)$, where ℓ_f is the chord of f . If $M \geq 0$, use $(t - \alpha)(\beta - t) \leq (\beta - \alpha)^2/4$ and $\ell_f(t) \geq \min\{f(\alpha), f(\beta)\}$. If $M < 0$, the correction term is nonnegative and $f \geq \ell_f$. $\square$

Verified enclosures of $f(\alpha)$, $f(\beta)$ and of f'' on the cell are enough, provided that f is twice continuously differentiable on the cell. On a cell that contains a nondifferentiable point of the expression, such as a zero of the argument of $|\cdot|$, only the natural enclosure is used. The implementation takes the larger of the two bounds on each cell.

F Complete separation to a positive tolerance

An exact support oracle evaluates one direction; a separation routine must also find a direction, and the bounded searches used inside a solver can miss one. Given an exact or approximate support oracle, the search can be completed in the following weak sense: for every query and every positive tolerance, a finite exact procedure either returns a valid cut that the query violates or certifies that the query is within the tolerance of the hull. The tools are classical; we state the procedure because its answer contract is what a certified implementation needs. The SCIP separator of Section 7 does not run it; it uses a fixed budget and never interprets an unsuccessful search as membership.

Setting. In this appendix $F : \mathbb{R}^d \rightarrow \mathbb{R}^r$ lists *all* lifted coordinates, including the block variables themselves, so that the joint graph is $F(P)$. Let $P \subset \mathbb{R}^d$ be a nonempty bounded rational polytope, let F be continuous with $F(P \cap \mathbb{Q}^d) \subseteq \mathbb{Q}^r$ (for example, polynomial with rational coefficients), let $\bar{z} \in \mathbb{Q}^r$ be the query and let $\epsilon > 0$ be rational. To cover the aggregated cuts of Section 5, a set $J \subseteq \{1, \dots, r\}$ of coordinates may be one-sided: we separate from

$$U_J = \text{conv } F(P) + \text{cone}\{e_j : j \in J\}.$$

For $F = (x, g(x))$, the query $(S\bar{v}, b - A\bar{v})$ and J equal to the coordinates of g , this is the set $K + Q$ of Theorem 5.2. Directions are taken from $C_J = \{c \in [-1, 1]^r : c_j \geq 0 \ (j \in J)\}$; we write $h(c) = \min_{x \in P} c^\top F(x)$, $\psi(c) = h(c) - c^\top \bar{z}$ and $\phi_J(w) = \sum_{j \notin J} |w_j| + \sum_{j \in J} \max\{w_j, 0\}$, and we let $\rho \geq \max_{x \in P} \|F(x) - \bar{z}\|_1$ be a rational bound, obtained for example from interval enclosures of the F_j .

Lemma F.1 (Distance duality). $\text{dist}_1(\bar{z}, U_J) = \min_{w \in \text{conv } F(P)} \phi_J(w - \bar{z}) = \max_{c \in C_J} \psi(c)$, and both extrema are attained. Moreover $|\psi(c) - \psi(c')| \leq \rho \|c - c'\|_\infty$ for all $c, c' \in C_J$.

Proof. The set $\text{conv } F(P)$ is compact, so U_J is closed and convex. Adding a cone vector reduces the distance in a coordinate $j \in J$ to $\max\{w_j - \bar{z}_j, 0\}$, which gives the first expression. For the second, $c^\top(z - \bar{z}) \geq \psi(c)$ for every $z \in U_J$ and $c \in C_J$, while $c^\top(z - \bar{z}) \leq \|z - \bar{z}\|_1$; hence $\psi(c) \leq \text{dist}_1(\bar{z}, U_J)$. If the distance ϱ is positive, a linear functional separates U_J from the open ℓ_1 ball of radius ϱ around $\bar{z}$; scaled to ℓ_∞ norm one, it lies in C_J because its infimum on U_J is finite, and it has $\psi(c) \geq \varrho$. The Lipschitz bound holds pointwise for every $x \in P$ and survives taking minima. $\square$

An *oracle of accuracy* $\theta \in [0, \epsilon)$ returns, for a rational $c \in C_J$, a rational lower bound $L_c \leq h(c)$ and a rational point $p_c \in P \cap \mathbb{Q}^d$ with $c^\top F(p_c) \leq L_c + \theta$. For quadratic features, Theorem 4.1 is such an oracle with $\theta = 0$; for polynomials of higher degree, Lemma F.3 gives one with any prescribed $\theta > 0$.

Column generation. Start with $W = \{F(p_0) - \bar{z}\}$ for a rational $p_0 \in P$, for example a vertex; if $F(p_0) = \bar{z}$, stop with a one-point certificate. Then repeat: solve the *master LP*

$$\text{UB} = \max_{c \in C_J} \min_{w \in W} c^\top w = \min \left\{ \phi_J \left(\sum_{w \in W} \lambda_w w \right) : \lambda \geq 0, \sum_w \lambda_w = 1 \right\} \quad (16)$$

in exact rational arithmetic, obtaining an optimal direction $\bar{c}$ and weights λ ; if $\text{UB} \leq \epsilon$, stop with the convex combination as a *within-tolerance* certificate; otherwise call the oracle at $\bar{c}$; if $L_{\bar{c}} > \bar{c}^\top \bar{z}$, return the cut $\bar{c}^\top z \geq L_{\bar{c}}$; otherwise add $F(p_{\bar{c}}) - \bar{z}$ to W . The two problems in (16) are an LP and its dual. This is the separation scheme of Fenchel cuts and local cuts [22, 33], in our norm and with cone coordinates; in direction space it is Kelley's cutting-plane method [30, 75] for maximizing the concave Lipschitz function ψ over C_J , and the termination proof below is its standard convergence argument.

Theorem F.2 (Complete weak separation). *If $\rho = 0$, then $F \equiv \bar{z}$ on P and the procedure stops without an oracle call. Otherwise it terminates after at most $2r[2\rho/(\epsilon - \theta)]^{r-1}$ oracle calls. A returned cut is valid for U_J and violated at $\bar{z}$; it is returned whenever $\text{dist}_1(\bar{z}, U_J) > \epsilon$. A within-tolerance answer consists of at most $r + 1$ rational points of P and rational weights whose mixture w satisfies $\phi_J(w - \bar{z}) \leq \epsilon$, so $\text{dist}_1(\bar{z}, U_J) \leq \epsilon$; it is returned whenever $\bar{z} \in U_J$.*

Proof. Validity of a returned cut follows from $L_{\bar{c}} \leq h(\bar{c})$ and $\bar{c}_j \geq 0$ on the cone coordinates. The master value UB is the value ϕ_J of a mixture of graph points, hence at least the distance; a basic optimal dual solution has at most $r + 1$ positive weights. If the distance exceeds ϵ , the procedure can therefore stop only with a cut.

For termination, let $\bar{c}_1, \bar{c}_2, \dots$ be the directions at which the oracle is called and w_s the point added after call s . Since no cut was returned at call s , $\bar{c}_s^\top w_s \leq L_{\bar{c}_s} + \theta - \bar{c}_s^\top \bar{z} \leq \theta$. At a later call t , $w_s \in W$ and $\text{UB} > \epsilon$, so $\bar{c}_t^\top w_s > \epsilon$. Hence $(\bar{c}_t - \bar{c}_s)^\top w_s > \epsilon - \theta$, and since $\|w_s\|_1 \leq \rho$, $\|\bar{c}_t - \bar{c}_s\|_\infty > (\epsilon - \theta)/\rho$. Because $\text{UB} > 0$ and the master objective is positively homogeneous, every $\bar{c}_t$ has a coordinate equal to ± 1 (equal to 1 on J), so it lies on one of at most $2r$ faces of C_J , each a product of $r - 1$ intervals of length at most two. Covering each face by $[2\rho/(\epsilon - \theta)]^{r-1}$ boxes of side at most $(\epsilon - \theta)/\rho$ bounds the number of calls. $\square$

A returned cut is only guaranteed to be violated. If a cut is returned only when $L_{\bar{c}} - \bar{c}^\top \bar{z} > \eta$ for a fixed $\eta \in [0, \epsilon - \theta)$, the same proof, with $\bar{c}_s^\top w_s \leq \eta + \theta$, gives termination after at most $2r[2\rho/(\epsilon - \theta - \eta)]^{r-1}$ calls, and every returned cut is violated by more than η ; this matters for

binary64 export, where a very small violation can be lost. The bound is exponential in r , as is the known worst case of Kelley-type cutting-plane methods [see the discussion in 43]; it is not a property of the problem. The set of directions C_J is a full-dimensional box even when $\text{conv } F(P)$ is lower-dimensional, so the central-cut ellipsoid method can localize a good direction in it. Using the oracle answers $-w/\|w\|_\infty$ as cuts in direction space and finishing with the master LP (16) over all recorded points gives the same two outcomes after $O(r^2 \log(r\rho/(\epsilon - \theta)))$ oracle calls (see below). For quadratic features on a polytope of fixed dimension, the oracle of Theorem 4.1 is polynomial, so weak separation with an exactly valid cut is then polynomial in the input length and $\log(1/\epsilon)$. For the Euclidean distance, Brierley et al. [25] state and solve the same contract for arbitrary convex sets by Gilbert’s nearest-point method [54], with a number of oracle calls that is polynomial in $1/\epsilon$ and does not depend on r .

A positive tolerance is needed even for certification. For $P = [-3, 3]$, $F(t) = (t, (t^2 - 2)^2)$ and $\bar{z} = (0, 0)$, the query $\bar{z} = \frac{1}{2}F(\sqrt{2}) + \frac{1}{2}F(-\sqrt{2})$ lies in the hull, but the valid inequality $z_2 \geq 0$ is tight at $\bar{z}$ and holds with equality on the graph only at $t = \pm\sqrt{2}$, so every representing mixture uses only irrational points.

Relation to the ellipsoid method. In the oracle framework of Grötschel et al. [59], weak separation is obtained from weak optimization for every convex body, without radius information. Convex bodies are full-dimensional by definition there, however, and the weak separation problem asks only for a hyperplane that is almost valid on the points deep inside the body. For a lower-dimensional hull this requirement is vacuous, and an almost valid hyperplane is not a cut. The contract above is different: the cut side is exact, only the membership side is approximate, and lower-dimensional hulls are allowed; the same contract appears in Brierley et al. [25].

An oracle-polynomial variant. Assume $\rho > 0$, put $\tau = (\epsilon - \theta)/2$, $\ell_0 = \min\{\tau/\rho, 1\}$ and $\nu = \ell_0^r/2$, and run the central-cut ellipsoid method of Grötschel et al. [59, Theorem 3.2.1] in $\mathbb{R}^r$ with outer radius r and volume bound ν , answering a query y as follows. If $y \notin C_J$, return the coordinate vector of a violated bound. Otherwise call the oracle at y , record $w_y = F(p_y) - \bar{z}$, keep the largest violation $\text{LB} = \max\{0, \max_y(L_y - y^\top \bar{z})\}$ found so far together with its cut, stop if $w_y = 0$, and return $-w_y/\|w_y\|_\infty$. After $O(r^2 \log(r\rho/(\epsilon - \theta)))$ calls, solve the master LP (16) over all recorded points. Its value satisfies $\text{UB} \leq \text{LB} + \epsilon$, so that either a cut is available ($\text{LB} > 0$) or the master LP gives a within-tolerance certificate. Indeed, let LB^* be the final value of LB and let $K^* = \{c \in C_J : w^\top c \geq \text{LB}^* + \theta + \tau \text{ for all recorded } w\}$. Every answer given during the run is a valid separation answer for K^* : coordinate answers separate y from $C_J \supseteq K^*$, and for a recorded w_y and $c \in K^*$ we have $w_y^\top c \geq \text{LB}^* + \theta + \tau > (L_y - y^\top \bar{z}) + \theta \geq w_y^\top y$. Since the method never receives a membership answer, it ends with an ellipsoid of volume at most ν that contains K^* . If $\text{UB} > \text{LB}^* + \epsilon$ at an optimal $\bar{c}$, then every $c \in C_J$ with $\|c - \bar{c}\|_\infty \leq \tau/\rho$ satisfies $w^\top c \geq \text{UB} - \tau > \text{LB}^* + \theta + \tau$ for all recorded w , so K^* contains a box of side ℓ_0 and has volume at least $\ell_0^r > \nu$, a contradiction. The queries have polynomial encoding length, and the final certificate is checked independently of the ellipsoid computation.

Polynomials of higher degree. For polynomial features, an exactly feasible sample net turns a support problem into a finite one.

Lemma F.3 (Rational domain net). *Let $v_1, \dots, v_s$ be the vertices of P , let $\Delta_\infty = \max_i(\max_P x_i -$*

$\min_P x_i$), let F be continuously differentiable on a box $B \supseteq P$, and let

$$\Lambda \geq \sum_{j=1}^r \sum_{i=1}^d \max_{x \in B} |\partial_i F_j(x)|.$$

For an integer $M \geq 1$ let $N_M = \{\sum_i (n_i/M)v_i : n_i \in \mathbb{Z}_{\geq 0}, \sum_i n_i = M\} \subset P$. Then for every $c \in [-1, 1]^r$,

$$\min_{x \in N_M} c^\top F(x) - \frac{\Lambda t \Delta_\infty}{M} \leq \min_{x \in P} c^\top F(x) \leq \min_{x \in N_M} c^\top F(x), \quad t = \min\{s-1, \dim P\}.$$

Proof. By Carathéodory's theorem in the affine hull of P , a point $x \in P$ is a convex combination $\sum_{i \leq t+1} \alpha_i v_{a_i}$ of at most $t+1$ vertices. Put $n_{a_i} = \lfloor M\alpha_i \rfloor$ for $i \leq t$ and give the remaining mass to $v_{a_{t+1}}$. The resulting $\hat{x} \in N_M$ satisfies $x - \hat{x} = \sum_{i \leq t} (\alpha_i - n_{a_i}/M)(v_{a_i} - v_{a_{t+1}})$, so $\|x - \hat{x}\|_\infty \leq t\Delta_\infty/M$. By the mean value theorem, $c^\top F$ is Λ -Lipschitz for the ℓ_∞ norm on B . Apply this at a minimizer over P ; the upper bound holds because $N_M \subset P$. $\square$

All points of N_M are exactly feasible, including on lower-dimensional polytopes, where a rectangular floating-point grid would almost never meet an equality exactly. With $M \geq 2\Lambda t \Delta_\infty / \epsilon$ the net is an oracle of accuracy $\epsilon/2$. The net has at most $\binom{M+s-1}{s-1}$ points, which is far too many for routine use; this is consistent with the strong NP-hardness of box-constrained quartic minimization on paths [37].

Outcomes and a reference implementation. The procedure returns *cut* with an exact rational row, *within tolerance* with a certificate, or *empty domain* when exact vertex enumeration of P finds no vertex. A bounded variant, which caps the number of oracle calls and net points, may in addition return *unresolved*, which carries no geometric information. A cut needs a lower bound over the whole domain, whereas a within-tolerance answer needs only feasible points and exact arithmetic. A rational cut can be exported in binary64 with the correction of Proposition 6.1, using enclosures of the features; rounding to nearest preserves the signs of the cone coefficients. Our reference implementation solves the master LP numerically, makes the returned weights nonnegative and renormalizes them in exact arithmetic, checks the resulting mixture exactly, and certifies every proposed direction with the exact oracle, so an inaccurate LP can waste work but cannot produce a wrong answer. Completeness is guaranteed by a grid fallback over C_J , with values $-1 + 2i/N$ in the free coordinates, i/N in the cone coordinates, $i = 0, \dots, N$, and $N = \lceil 2\rho/(\epsilon - \theta) \rceil$; every direction is within ℓ_∞ distance $1/N$ of this grid, so if no grid direction yields a cut, Lemma F.1 bounds $\max_{C_J} \psi$ by $\theta + \rho/N \leq \epsilon$, and the within-tolerance answer is then one exactly feasible point per grid direction. An exact version with a rational simplex method for (16) is part of the verification scripts (see Code and data availability); on the test instances it needed at most six oracle calls, while the bound of Theorem F.2 ranged from 10^2 to 10^{14} . For aggregated cuts, a returned rational row maps to an original-variable cut through Proposition 5.1; a within-tolerance answer bounds a distance in the lifted coordinates (x, g) , not a distance in the original variables to the set $\mathcal{C}$ of Theorem 5.2.

G The first two campaigns

Both campaigns were run with code, limits and protocol fixed in advance. Their records were audited independently from the raw run records (418 checks of the original reports, all passed), and the statistics below were recomputed from the same records.

Campaign 1: lifted cuts with auxiliary variables. The first implementation added an auxiliary variable $w = f(x)$ for each selected expression of one or two variables and separated support cuts in (x, w) (Section 6.1). Its oracles were the polytope oracle of Theorem 4.1 in dimension two, the star oracle of Section 4.2, Bernstein bounds and ball arithmetic. Because the auxiliary variables change the model, it was compared with a *control* mode that adds them without cuts. The automatic mode also used an inexpensive screening certificate, a dual-norm bound on the violation of any cut computed from a certified convex combination of graph points [88]; it allowed only 2 of 506 attempted skips and was dropped afterwards. The campaign used 24 MINLPLib models selected by a fixed hash ordering (Appendix H), 13 constructed mechanism cases (small models with shared expressions and linking rows, built so that joint support cuts close a gap that separate envelopes leave), a historical suite of four models from earlier development, a six-second budget and a two-second root budget.

Of the 24 models, 20 passed the model import in every mode; three were rejected because a domain restriction could not be proved on the declared box, and one failed with an unsupported variable power. Baseline and control solved 19 of the 20, both cut modes 18; the lost model, `genpooling_lee2`, was solved by the baseline in 5.6 s and by the control in 4.9 s, and reached the 6 s limit in both cut modes. On the 18 models solved in every mode, the shifted geometric mean time was 0.26 s for the baseline, 0.41 s for the control, 0.57 s for `all` and 0.46 s for `auto`. On the mechanism cases, the control and both cut modes solved all 13 and the baseline 12, so the additional solve is explained by the reformulation, not by the cuts. On the historical model `waterno2_06` the final dual bound after 6 s was 26.59 for the baseline and 21.17 for the control, against 7.05 (`all`) and 2.41 (`auto`) for this minimization model. The 1,082 recorded cuts (450 from the two-dimensional polytope oracle, 125 from the star oracle, 460 from Bernstein bounds and 47 from ball arithmetic) all passed replay, with 12 of 12 corrupted records rejected, and all 270 returned incumbents passed the original-model check. The auxiliary variables thus cost time on their own, and the cuts did not recover it; this motivated the original-variable cuts of Section 5.

Campaign 2: cuts in the original variables. The second implementation is, up to the repairs described below, the one of Section 7. Its campaign used 30 new MINLPLib models, ten per stratum, selected by a fixed hash ordering from the 422 eligible models that had not been used before, with a 30-second budget, root runs with one node, 13 mechanism cases, ten diagnostic models from the first campaign, and seed-one repeats on six models (282 runs). All 30 models were admitted, and every mode solved the same 25. The five unsolved models received no cut in any mode. Cuts were added on 7 models in mode `all` (38 cuts) and on 3 in mode `auto` (10 cuts), all of which the baseline also solved; four of them were solved at the root in every mode. At the relative tolerance 10^{-4} no final or root dual bound was better in a cut mode than in the baseline; four time-limited models had worse final bounds in both cut modes, and `graphpart_clique-40` had a worse root bound in both cut modes. At a tolerance of 10^{-6} two models appeared better and up to seven worse, differences of the same size as the variation between identical baseline reruns. On the 25 commonly solved models the shifted geometric mean time rose from 0.55 s to 0.99 s (`all`) and 0.96 s (`auto`), and each cut mode was more than 10% slower on 21 of the 25 (both on 19). On the 30 full runs of mode `all`, discovery accounted for 24.4 of the 27.6 callback seconds; direction LPs took 0.7 s and certification 2.4 s. On the diagnostic model `btest14` the final bound was -72.30 for the baseline and -115.60 in both cut modes.

The campaign exposed two defects of the frozen code. On two diagnostic models, symbolic polynomial conversion over all model symbols crashed during discovery (four worker errors, with unknown cut logs), and discovery did not check the time allowance between rows, so it exceeded

the allowance in 46 cut-mode runs, by up to 19s. Both were repaired; the repair also added two guards, rejecting rows whose side reaches SCIP’s infinity value (all 123 recorded rows of campaign 2 were below it) and tolerating a failed evaluation of an exchange sample. A matched cohort of the 25 affected model and seed groups, 75 runs, was rerun with all three modes. It had no worker errors, every mode solved the same models, and the largest remaining overrun of the separator allowance was 0.03s. The worse cut-mode root bound of `graphpart_clique-40` disappeared, but the cut modes still had worse final bounds on two time-limited models, `waterno2_06` (76.75 in the cut modes against 81.64 for the baseline, with no cut added) and `waterx`, which reflects the separator’s time, not the cuts. All 123 cuts of the campaign and all 42 of the repair cohort passed replay, and all 338 returned incumbents passed the original-model check.

H Work limits, modes, test sets and detailed results

Separator limits. Table 6 lists the limits of the two implementations, and Table 7 the modes of campaigns 3 to 5. The default limits of the second implementation are those of modes `all` and `auto`; the other modes raise the limits listed in Table 7, including the number of root callbacks, and set the allowance to $\min\{s, 0.5T\}$ with the listed s .

Table 6: Work limits of the separators. “Allowance” is the cumulative separator time per run for a total budget T ; “sides” are signed row sides (an equality has two).

Setting	Campaign 1 (lifted cuts)	Campaigns 2–4 (original variables)
Blocks	expressions of 1–2 variables, stars with 2–4 leaves; at most 64 blocks	sides with 1–4 variables at most 32 blocks, largest first
Sides per block	—	6 nonlinear, 16 affine
Nonpolynomial sides	one variable	one variable
Polynomial degree in discovery	8 (one variable), 4 (pairs)	total degree 8
Root callbacks	5	3
Cuts (total / per callback)	24 / 6	12 / 4
Support calls per run	not capped	24
Directions per block and callback	LP exchange, 3 rounds	one per source side, then up to 3 LPs
Samples	65 / 13×13 / 128 Halton points plus corners	33 / 7×7 / 3^d grid points, plus exact vertices and centroid
LP time limit	0.1 s	0.05 s (HiGHS, one thread)
Exact quadratic oracle	polytope ($d = 2$); star	polytope ($\leq 2,000$ subsets); star
Subdivision	128 cells, depth 16	128 cells, depth 16
Bernstein	1–2 variables, $\leq 1,024$ coefficients	same, degree ≤ 32 per variable
Ball arithmetic	128 bits	128 bits
Allowance	$\min\{2s, 0.15T\}$	$\min\{1s, 0.05T\}$
Violation threshold (all / auto)	$10^{-5} / 5 \cdot 10^{-4}$	$10^{-5} / 5 \cdot 10^{-4}$
Cut flags	global, removable, forced	global, removable, forced
Relative gap limit	10^{-4}	10^{-4}

Selection. The models come from a copy of MINLPLib [28, 125] in OSiL format downloaded on 2026-09-21 (1,632 files), with the library metadata file of 2026-09-12; both are archived with their hashes. Campaign 1 drew its 24 models from the 279 models with at most 80 variables, 120 constraints and 150,000 bytes of OSiL, a quadratic or general nonlinear constraint and at least two nonfixed bounded variables, after excluding 123 names, in the order of the SHA-256 hash of a fixed prefix and the name. Campaigns 2 and 3 use the 422 models with at most 120 variables, 180 constraints and 250,000 bytes of OSiL, at least two nonfixed bounded variables and at least one

Table 7: Modes of campaigns 3 to 5. “First directions” are the directions tried before the LP search: remainder ($a = 0$), whole row (a equal to the side’s coefficients of the block variables) or block (all sides of the block summed), see Section 7.2. n is the number of copies of a constructed instance and k the number of leaves per star; “star” blocks have up to 16 leaves. In the text, mech. (base) limits are those of **all-diag-mech** and wide limits those of **frozen-wide**.

Mode	First directions	Blocks	Max blocks	Cuts	Per callback	Call- backs	Support calls	s
all, auto	remainder	≤ 4 var.	32	12	4	3	24	1 s
all-diag	remainder	≤ 4	128	200	50	10	500	30 s
all-diag-rowdir	whole row	≤ 4	128	200	50	10	500	30 s
all-diag-mech	remainder	≤ 4	n	$4n$	n	10	$20n$	60 s
whole row, mech. (Part 3P)	whole row	≤ 4	n	$4n$	n	10	$20n$	60 s
frozen-wide	remainder	≤ 4	n	$16n$	$4n$	10	$40n$	60 s
rowdir-wide	whole row	≤ 4	n	$16n$	$4n$	10	$40n$	60 s
frozen-, rowdir-cap32	remainder; whole row	≤ 4	n	$32n$	$8n$	40	$80n$	150 s
frozen-, rowdir-cap64	remainder; whole row	≤ 4	n	$64n$	$16n$	40	$160n$	150 s
rowdir-star4	whole row	≤ 4	nk	$4nk$	nk	10	$10nk$	60 s
agg-star4	block, whole row	≤ 4	nk	$4nk$	nk	10	$10nk$	60 s
agg-star	block, whole row	star, then ≤ 4	nk	$4nk$	nk	10	$10nk$	60 s
SCIP parameter modes (no certified cuts):								
baseline-novarlocks	constraints/nonlinear/checkvarlocks = d (implicit discreteness off)							
baseline-extra	separating/eccuts/freq = 0, separating/interminor/freq = 0, nlhdlr/quadratic/useintersectioncuts, separating/rlt/detecthidden, separating/rlt/hiddenrlt set to true							
*-noaggr	presolving/donotaggr and presolving/donotmultaggr set to true							

nonlinear function, after excluding the 160 names used or excluded by the first campaign and 7 models whose expressions the importer does not support (**tanh**, **log10**, **signpower**, **gammaFn**, **erf**); the strata contain 85, 208 and 129 models. Campaign 2 took the first ten per stratum in the order of SHA-256(**convexification-holdout-v2**: + name). The structure-selected sample of campaign 3 ranks the qualifying models of the remaining pool by SHA-256(**convexification-structure-v3**: + name) and takes the first 30.

The structure-selected models of Part B2 are those of Part B. The pool of Part D consists of the 322 models of the same MINLP Lib copy with more than 120 and at most 1,000 variables, not convex by the library’s metadata, with at least one quadratic or polynomial function and an OSiL file below 2 MB, that were not used in campaigns 1–3. A scan of import and discovery (60-second discovery deadline, 120-second process limit) admitted 301 of them; 8 were refused as infeasible in exact arithmetic (Section 7.1), 9 contained unsupported input or an unproved domain, and 4 did not finish the import within 120 seconds. Of the 301, 85 have a block that the **auto** rule admits among the first 32 blocks of discovery. SCIP’s baseline, run once for 60 seconds on each of these 85, solved 27, and Part D takes the first 20 of the remaining 58 in the order of SHA-256(**convexification-hard-v4**: + name): **blend480**, **kall_ellipsoids_tc05a**, **kall_ellipsoids_tc02b**, **pooling_sppa0stp**, **camshape100**, **multiplants_mtg1a**, **ringpack_20_2**, **mpbp_31**, **hydroenergy2**, **waternd2**, **kriging_peaks-full100**, **multiplants_mtg1c**, **blend718**, **kall_circlesrectangles_c6r1**, **pooling_sppa0pq**, **kall_circlespolygons_c1p5b**, **sonet23v4**, **ringpack_20_1**, **crudeoil_li03** and **multiplants_mtg6**. The root gap of these models is measured against the best known primal value of the library’s metadata; **mpbp_31** has none.

Path family. The instances of Section 8.5 are generated by a fixed script from the seeds $1000n + s$ with $s \in \{0, \dots, 4\}$ (campaign 3) and $s \in \{5, \dots, 9\}$ (campaign 4), with $t_i \in [-10, 10]$; the instances of Part 4C4 differ from the latter only in the right-hand side of the coupling row. The star instances

of Part 5S are generated from the seeds $100000k + 1000n + s$, $s \in \{0, \dots, 4\}$ (Section 8.6). All instances, the generators and the reference values with their certificates are part of the supplementary archive.

Table 8: Separator funnel, summed over the runs of a part and mode: support calls, failed certifications, certified rows below the violation threshold, rows rejected by the stored-row check, and cuts added; then the separator callback time and its main parts (discovery; direction search, that is, sampling and direction LPs; certification). No row was ever discarded by safe rounding. In the constructed families the separator runs only at the root, so full and root runs have the same funnel; the times shown are those of the root runs. The three timed parts do not add up to the total, which also contains untimed work.

Part	Mode	Support	Failed	Below	Stored-row	Cuts	Callback time (s)			
		calls	cert.	thresh.	rejected	added	total	discovery	directions	certif.
3A full	all	684	57	461	15	151	19.1	6.3	2.5	8.6
3B full	all	884	14	646	45	179	32.9	4.4	3.9	23.0
3B root	all-diag	1,618	19	1,039	202	358	45.1	2.3	9.6	30.4
4B2 root	all-diag , no aggr.	1,572	17	1,058	47	450	38.6	2.1	8.6	25.7
4D root	all	176	2	149	11	14	18.4	15.7	0.4	1.8
4D root	all-diag	6,352	313	5,042	106	891	188.4	53.6	22.8	69.8
3C	remainder, base	3,462	0	391	71	3,000	73.6	10.6	9.6	32.0
4C2	remainder, wide	13,502	0	1,189	313	12,000	240.4	10.3	31.1	121.3
4C3	whole row, wide	15,534	0	3,181	353	12,000	269.2	10.3	32.6	140.1
4C4	remainder, wide	14,186	0	1,939	247	12,000	261.9	10.2	36.5	129.6
4C4	whole row, wide	15,204	0	2,989	215	12,000	274.8	10.7	36.6	137.7
5C-b	whole row, cap $64n$	25,340	0	5,282	286	19,772	459.7	10.7	68.2	233.6
5S	block, ≤ 4 var.	10,576	0	2,496	65	8,015	1,700.3	32.6	27.1	1,557.1
5S	block, stars	15,187	0	6,697	48	8,442	1,696.6	32.3	26.8	1,496.6

Per-model and per-instance results. Tables 9–17 list the runs behind Sections 8.3, 8.5 and 8.6. A root bound is SCIP’s dual bound at the end of the root node; for a run that ended at the root without reporting one, because the root node was pruned, it is the final dual bound. Times are the seconds charged to the budget (Section 8.1). Numbers have four significant digits; values of at least 1000 are rounded to integers.

Table 9: Part 3A: the 30 models of campaign 2, in the order of selection. **syn05m** is a maximization model, marked (max), for which a smaller bound is stronger; all other models are minimization models. Root runs (node limit 1, 60s): root bound of the baseline (SCIP default) and of mode **all-diag**, and the number of cuts that **all-diag** added. Full runs (seed 0, 300s): status and time in seconds of the baseline and of mode **all**, and the number of cuts that **all** added. Status: o optimal, g gap limit, t time limit.

Model	Root run			Full run, seed 0		
	Baseline	all-diag	Cuts	Baseline	all	Cuts
p_ball_10b_5p_2d_m	0	0	36	o 1.464	o 1.877	3
st_glmpp_kk90	3	3	0	g 0.03213	g 0.06396	0
du-opt5	5.181	5.181	0	o 0.6685	o 0.7592	0
nvs24	-1036	-1036	0	g 4.976	g 5.231	0
nous1	0.3451	0.3451	30	g 2.787	g 4.466	2
supplychain	1837	1837	0	o 0.2283	o 0.3114	0
ex4_1_8	-16.74	-16.74	15	g 0.03344	g 0.1156	5
ex8_3_4	-10	-10	8	t 300.1	t 300.1	6
cvxnonsep_normcon20r	-21.75	-21.75	200	g 0.0776	g 0.1926	12
cvxnonsep_psig20r	35.6	35.63	33	g 0.1401	g 0.3581	10
hybridynamic_var	1.504	1.504	2	g 0.5563	g 0.6918	2
graphpart_clique-40	0.2857	0.2857	0	o 146.1	o 147.5	0
pooling_adhya4pq	-961.9	-961.9	60	o 0.7906	o 1.138	1
st_glmpp_kk92	-12	-12	0	g 0.03497	g 0.05116	0
st_testph4	-80.5	-80.5	0	o 0.03173	o 0.03361	0
du-opt	3.512	3.512	0	o 3.479	o 3.378	0
ex5_2_5	-9700	-9700	0	t 300.1	t 300.1	0
kall_congruentcircles_c41	0	0	0	g 0.1283	g 1.034	0
ex14_2_7	$-9.705 \cdot 10^{-9}$	$-9.705 \cdot 10^{-9}$	0	o 0.1273	o 0.1771	0
st_e22	-85	-85	1	g 0.03488	o 0.06303	1
kall_congruentcircles_c51	0	0	10	o 2.958	o 3.459	6
tln4	8.3	8.3	0	o 0.1384	o 0.1656	0
synthes1	6.007	6.007	0	o 0.09826	o 0.1314	0
syn05m (max)	837.7	837.7	8	o 0.04351	o 0.2115	2
nvs17	-1105	-1105	0	o 0.8779	o 0.9978	0
tln7	13.81	13.81	0	t 300.1	t 300.1	0
st_testgr3	-20.6	-20.6	0	o 0.1448	o 0.165	0
graphpart_2g-0066-0066	-2891970	-2891970	0	o 1.004	o 1.089	0
ex1223	4.579	4.579	0	g 0.06082	g 0.371	0
waterx	182.8	182.8	0	t 300.1	t 300.1	0

Table 10: Parts 3B and 4B2: the 30 structure-selected models, in the order of selection. **pointpack02** and **pointpack04** are maximization models, marked (max), for which a smaller bound is stronger; all other models are minimization models. Root bounds (node limit 1, 60s) of the baseline (SCIP default) and of mode **all-diag** in Part 3B; of the same two with presolve aggregation disabled (**baseline-noaggr**, **all-diag-noaggr**) and of SCIP with its disabled nonconvex separators switched on (**baseline-extra**) in Part 4B2. Full runs of Part 3B (seed 0, 300s): status and time in seconds of the baseline and of mode **all**, and the number of cuts that **all** added. Status: o optimal, g gap limit, t time limit.

Model	Root bound, Part 3B		Root bound, Part 4B2		Full run, Part 3B, seed 0			
	Baseline	all-diag	no aggregation		Extra	Baseline	all	Cuts
			Baseline	all-diag				
pooling_bental4tp	-450.2	-450	-497.3	-462.7	-450	o 0.09327	o 0.6675	3
wastewater04m2	71.96	71.96	71.96	71.96	80.17	g 0.294	g 0.6074	0
sep1	-522.7	-522.7	-524.5	-522.8	-533.8	o 0.1762	o 0.5253	0
kall_congruentcircles_c63	0	0	0	0	0	o 2.109	o 2.325	5
kall_circlespolygons_c1p12	0	0	0	0	0	g 39.92	g 40.31	0
nvs02	5.964	5.964	5.964	5.964	5.964	o 0.03464	o 1.045	11
kall_congruentcircles_c71	0	0	0	0	0	o 191.4	o 190.7	4
pooling_haverly2pq	-617.7	-617.7	-857.1	-857.1	-600	o 0.07421	o 0.4832	2
ex3_1_4	-6	-5.692	-6	-5.692	-4	o 0.07758	o 0.1925	4
bayes2_20	$1.376 \cdot 10^{-13}$	$1.376 \cdot 10^{-13}$	0	0	$2.58 \cdot 10^{-11}$	t 300.1	t 300.1	0
pooling_adhya4tp	-976.4	-976.4	-976.4	-976.4	-881.9	o 0.9462	o 1.346	2
prob06	1.177	1.177	1.177	1.177	1.177	g 0.03413	g 0.1695	9
tanksize	0.8432	0.8433	0.8424	0.8431	0.8473	o 1.307	o 2.531	0
kall_congruentcircles_c32	0	0	0	0	0.105	o 0.1967	o 0.7083	4
ex8_1_7	-5911	-5911	-5911	-5911	-5838	g 18.21	g 17.04	7
pooling_haverly3pq	-750	-750	-750	-750	-750	o 0.04308	o 0.4512	4
bayes2_50	$2.842 \cdot 10^{-13}$	$3.766 \cdot 10^{-13}$	0	0	$1.987 \cdot 10^{-12}$	t 300.2	t 300.1	0
pooling_bental4pq	-464.9	-450	-464.9	-450	-450	o 0.1062	o 0.4949	3
pooling_haverly3tp	-750.8	-750.8	-875	-875	-750	o 0.05725	o 0.3735	2
pointpack02 (max)	2	2	2	2	2	o 0.03801	o 0.186	2
pooling_rt2tp	-5528	-5528	-5528	-5528	-4934	o 0.3343	o 0.8944	1
pooling_bental5tp	-3500	-3500	-3500	-3500	-3500	o 0.1665	o 0.1476	0
st_e02	201.2	201.2	201.2	201.2	201.2	o 0.0224	o 0.03373	0
kall_congruentcircles_c72	0	0	0	0	0	o 8.591	o 10.71	4
kall_congruentcircles_c61	0	0	0	0	0	o 23.17	o 19.66	5
pooling_haverly1tp	-400	-400	-414.5	-422.2	-400	o 0.03718	o 0.3174	2
bayes2_30	$1.705 \cdot 10^{-13}$	$1.705 \cdot 10^{-13}$	0	0	$1.371 \cdot 10^{-12}$	t 300.2	t 300.1	0
pointpack04 (max)	1.164	1.109	1.164	1.109	1.007	o 0.1776	o 0.8163	11
kall_circles_c8a	0	0	0	0	0	t 300.1	t 300.1	4
kall_circlesrectangles_c1r11	0.1996	0.1996	0.1996	0.1996	0.1996	o 0.09139	o 0.08552	0

Table 11: Part 4D: the 20 larger models that SCIP did not solve in 60 s, in the order of selection. Best known: best known primal value of the MINLP Lib metadata (none for `mpbp_31`); sense: objective sense; for maximization models a smaller bound is stronger. Root runs (node limit 1, 120 s): root bound of the baseline (SCIP default), of mode `all-diag` and of SCIP with its disabled nonconvex separators switched on (`baseline-extra`), and the number of cuts that `all-diag` added. Full runs (seed 0, 300 s): status and final dual bound of the baseline, of mode `all` and of `baseline-extra`. Status: o optimal, g gap limit, t time limit. [†]The incumbent of the baseline-extra root run of `waternd2` failed the independent primal check: SCIP reports the objective value 3084096, the original model gives 2297367. The run is flagged, and its bound is not used in comparisons.

Model	Best known	Sense	Root run				Full run: final dual bound		
			Baseline	<code>all-diag</code>	Extra	Cuts	Baseline	<code>all</code>	Extra
<code>blend480</code>	9.227	max	10.8	10.8	10.68	0	o 9.227	o 9.227	o 9.227
<code>kall_ellipsoids_tc05a</code>	39.4	min	20.99	20.99	20.99	200	t 20.99	t 20.99	t 20.99
<code>kall_ellipsoids_tc02b</code>	32.4	min	22.5	22.5	22.5	200	t 22.5	t 22.5	t 22.5
<code>pooling_sppa0stp</code>	-35812	min	-37480	-37480	-37232	54	t -36867	t -36869	t -36827
<code>camshape100</code>	-4.284	min	-5.027	-5.027	-4.772	0	t -4.574	t -4.574	t -4.529
<code>multiplants_mtg1a</code>	391.6	max	1781	1823	1798	4	t 856.8	t 859.1	t 942.5
<code>ringpack_20_2</code>	-39.34	min	-41.72	-41.72	-41.72	0	t -41.72	t -41.72	t -41.72
<code>mpbp_31</code>	-	max	2149	2149	2149	0	t 2071	t 2071	t 2081
<code>hydroenergy2</code>	371812	max	379602	379795	379705	133	t 376618	t 376633	t 375963
<code>waternd2</code>	1062685	min	1047772	1047772	1044153 [†]	0	t 1051237	t 1051237	t 1050730
<code>kriging_peaks-full100</code>	-2.638	min	-334.3	-335.4	-279.8	3	g -2.638	g -2.638	t -2.639
<code>multiplants_mtg1c</code>	684	max	8893	8574	8647	4	t 6603	t 6586	t 7086
<code>blend718</code>	7.394	max	20.69	20.69	20.67	1	t 10.56	t 10.4	t 9.523
<code>kall_circlesrectangles_c6r1</code>	7.165	min	0	0	0	36	t 0.7848	t 0.7848	t 0
<code>pooling_sppa0pq</code>	-35812	min	-37780	-37780	-37288	42	t -36415	t -36415	t -36409
<code>kall_circlespolygons_c1p5b</code>	3.77	min	0	0	0	143	t 0	t 0	t 0
<code>sonet23v4</code>	-22748	min	-51431	-51431	-51431	0	t -39634	t -39618	t -39577
<code>ringpack_20_1</code>	-39.34	min	-41.72	-41.72	-41.72	0	t -41.72	t -41.72	t -41.72
<code>crudeoil_li03</code>	3487	max	3577	3578	3578	31	t 3561	t 3560	t 3557
<code>multiplants_mtg6</code>	5314	max	6870	6870	6869	40	t 6819	t 6819	t 6865

Table 12: The path family, seeds 0–4 (the instances of Part 3C), root runs (node limit 1, 120 s): number of copies n , seed, optimum z^* and SCIP’s root bound z_{base} (Part 3C). Root gap closed $(z - z_{\text{base}})/(z^* - z_{\text{base}})$, with z the root bound of SCIP-nolocks and SCIP-extra and of the separator with remainder or whole-row first directions and base or wide limits (Section 7.2, Table 7): remainder with base limits is **all-diag-mech**, remainder with wide limits **frozen-wide**, and the whole-row columns are the post hoc runs of Part 3P. The row *Part* gives the part of each column. Full runs are in Table 13.

<i>n</i> <i>Part</i>	Seed	Optimum	SCIP	Root gap closed						
			root							
			bound							
			3C	SCIP-nolocks	SCIP-extra	Remainder		Whole row		
				4C2	4C2	base	wide	base	wide	
						3C	4C2	3P	3P	
10	0	0.0498	−0.2492	0.8292	0.537	0.4318	0.8347	0.5913		1
10	1	0.0105	−0.2523	0.9566	0.6395	0.2843	0.7833	0.7559		1
10	2	0.01721	−0.3413	0.9484	0.6744	0.35	0.863	0.4418		1
10	3	0.01306	−0.2881	0.953	0.7326	0.3835	0.9505	0.3111		1
10	4	0.0531	−0.2362	0.8127	0.4949	0.08081	0.8909	0.2553		1
20	0	0.07166	−0.5796	0.8863	0.6327	0.2495	0.8497	0.4383		1
20	1	0.0343	−0.2612	0.8757	0.4993	0.2208	0.8419	0.3543		1
20	2	0.06104	−0.3958	0.8612	0.5805	0.3275	0.8742	0.477		1
20	3	0.06726	−0.648	0.9024	0.5438	0.2813	0.911	0.447		1
20	4	0.04834	−0.6198	0.9245	0.6277	0.08006	0.9326	0.3712		1
40	0	0.07776	−1.226	0.9364	0.6926	0.2696	0.8807	0.4332		1
40	1	0.09766	−0.9456	0.9012	0.5925	0.1871	0.8867	0.441		1
40	2	0.08752	−1.417	0.9385	0.6653	0.3301	0.9178	0.4643		1
40	3	0.08875	−1.275	0.9312	0.6589	0.148	0.8585	0.3642		1
40	4	0.09021	−0.9454	0.9084	0.5895	0.1906	0.8907	0.4157		1
80	0	0.1611	−1.749	0.9106	0.5141	0.1701	0.8758	0.4577		1
80	1	0.2155	−2.265	0.9092	0.5909	0.3286	0.9304	0.5156		1
80	2	0.1366	−1.764	0.9229	0.5381	0.2241	0.85	0.4268		1
80	3	0.1799	−2.383	0.926	0.6393	0.3316	0.8983	0.4898		1
80	4	0.2269	−1.818	0.8844	0.6342	0.1995	0.9005	0.3705		1

Table 13: The path family, seeds 0–4, full runs (300 s): status and time in seconds of SCIP, SCIP-nolocks, SCIP-extra and Gurobi, and of the separator with remainder or whole-row first directions and base or wide limits (modes as in Table 12). The row *Part* gives the part of each column. Solved: number of the runs above it in the column with status optimal or gap limit and an incumbent that passed the independent primal check (Section 8.1). The solved counts are those of Table 4. Status: o optimal, g gap limit, t time limit.

Full run: status and time (s)									
<i>n</i> <i>Part</i>	Seed					Remainder		Whole row	
		SCIP 3C	SCIP-nolocks 4C2	SCIP-extra 4C2	Gurobi 4C2	base 3C	wide 4C2	base 3P	wide 3P
10	0	o 2.589	o 1.053	o 1.025	g 45.63	o 1.928	o 3.033	o 1.225	o 3.343
10	1	o 1.271	o 0.321	o 0.6125	o 0.3251	o 1.544	o 3.132	o 1.628	o 3.403
10	2	o 2.182	o 0.3417	o 0.7333	o 0.3456	o 1.688	o 3.54	o 1.452	o 3.528
10	3	o 1.389	o 0.3309	o 1.643	g 0.2899	o 1.868	o 3.713	o 1.441	o 3.326
10	4	o 1.252	o 0.7778	o 0.9129	g 77.62	o 1.553	o 3.66	o 1.658	o 3.348
20	0	t 300.2	o 37.85	o 67.58	t 300.2	t 300.1	o 8.896	o 9.517	o 7.763
20	1	o 26.67	o 19.51	o 55.53	g 88.67	o 6.068	o 7.314	o 2.987	o 7.168
20	2	o 139.9	o 53.72	o 49.38	g 0.4913	o 49.92	o 11.87	o 7.187	o 7.072
20	3	o 30.91	o 41.85	o 26.02	t 300.2	o 9.65	o 7.377	o 4.075	o 6.89
20	4	o 168	o 17.43	o 34.38	t 300.2	o 160.8	o 7.655	o 4.897	o 7.138
40	0	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	o 26.25	t 300.1	o 15.64
40	1	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 16.56
40	2	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	o 167.5	t 300.1	o 16
40	3	t 300.2	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 15.22
40	4	t 300.1	t 300.2	t 300.2	t 300.2	t 300.2	o 38.38	t 300.1	o 15.94
80	0	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 37.1
80	1	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 37.9
80	2	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 38.59
80	3	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 37.85
80	4	t 300.1	t 300.2	t 300.2	t 300.2	t 300.1	t 300.2	t 300.1	o 36.77
Solved		9	10	10	7	9	13	10	20

Table 14: The path family, seeds 5–9, root runs: Part 4C3 with the coupling row of Part 3C, which does not bind, and Part 4C4 with a binding coupling row. Number of copies n , seed, optimum z^* ; opt.–closure: optimum minus the block closure (Section 8.5), in units of 10^{-4} ; SCIP’s root bound z_{base} . Root gap closed $(z - z_{\text{base}})/(z^* - z_{\text{base}})$, with z the root bound of SCIP-nolocks (nolocks), SCIP-extra (extra) and the separator: remainder directions with base limits (**all-diag-mech**), wide limits (**frozen-wide**) or the cut cap $64n$ (**frozen-cap64**); whole-row directions with wide limits (**rowdir-wide**) or the cut cap $64n$ (**rowdir-cap64**). Closure–root: block closure minus the root bound of the two modes with cut cap $64n$, in units of 10^{-4} . Root runs: node limit 1, 120 s; Part 5C-b: up to 40 callbacks, 300 s (Table 7). Columns without a stated part are from the part of the block. Empty cells: not run in that part. Full runs are in Table 15.

			SCIP		Root gap closed						Closure—root		
			opt.—	root	SCIP		Remainder			Whole row		cap 64n	
n	Seed	Optimum	closure	bound	nolocks	extra	base	wide	64n	wide	64n	rem.	row
Part 4C3: coupling row $\sum_i y_i \leq 0.8n$, not binding (SCIP-nolocks: Part 5C-a)													
10	5	0.01917		−0.2424	0.9214	0.5724	0.3574			1			
10	6	0.03162		−0.1446	0.8134	0.5799	0.4557			1			
10	7	0.005493		−0.1714	0.9623	0.4882	0.0151			1			
10	8	0.01416		−0.2576	0.9433	0.518	0.2162			1			
10	9	0.02332		−0.2533	0.9111	0.7324	0.3226			1			
20	5	0.06519		−0.4708	0.8736	0.5394	0.374			1			
20	6	0.047		−0.8441	0.9444	0.6713	0.3105			1			
20	7	0.04333		−0.5102	0.9173	0.6462	0.09581			1			
20	8	0.0542		−0.5584	0.9077	0.6399	0.3355			1			
20	9	0.03711		−0.4975	0.926	0.5571	0.1693			1			
40	5	0.07678		−0.9965	0.9236	0.57	0.3122			1			
40	6	0.06873		−0.962	0.9289	0.6112	0.2425			1			
40	7	0.05786		−1.031	0.9421	0.6517	0.2382			1			
40	8	0.08008		−0.7706	0.9002	0.423	0.1168			1			
40	9	0.135		−1.144	0.8906	0.6221	0.2262			1			
80	5	0.1542		−1.847	0.9183	0.6688	0.2946			1			
80	6	0.1897		−2.472	0.9249	0.6597	0.1909			1			
80	7	0.2371		−1.969	0.888	0.5742	0.2048			1			
80	8	0.23		−2.376	0.9077	0.6363	0.1577			1			
80	9	0.1406		−1.889	0.9259	0.6785	0.2242			1			
Part 4C4: coupling row $\sum_i y_i \leq c$, binding (SCIP-nolocks: Part 5C-a; cap 64n: Part 5C-b)													
10	5	0.1301	0	−0.02472	0.9627	0.2872		0.9886	0.9996	0.9969	0.9999	0.5682	0.2192
10	6	0.2547	0	0.2214	0.9909	0.4199		0.9993	0.9998	0.9997	0.9997	0.07091	0.1128
10	7	0.09535	0.1526	−0.0234	0.9801	0.2386		0.9557	0.9997	0.9948	0.9994	0.1569	0.5325
10	8	0.06988	0	−0.1117	0.98	0.1878		0.8748	0.9998	0.9807	0.9998	0.2864	0.3217
10	9	0.1561	0.6104	−0.009054	0.98	0.6403		0.9833	0.9996	0.9968	0.9994	0.07097	0.4479
20	5	0.3409	0	0.09005	0.9797	0.5995		0.9599	0.9997	0.9967	0.9998	0.8265	0.6161
20	6	0.4741	0.2713	−0.0242	0.9708	0.7421		0.9453	0.9998	0.9937	0.9998	0.8505	0.5073
20	7	0.3526	1.526	0.01972	0.9471	0.2988		0.9644	0.9994	0.9841	0.9994	0.46	0.5782
20	8	0.5402	6.473	0.2429	0.9818	0.758		0.9514	0.9975	0.989	0.9975	0.9822	1.085
20	9	0.3327	0	0.1314	0.9898	0.5294		0.9865	0.9994	0.9869	0.9991	1.207	1.769
40	5	0.63	0	0.0889	0.9734	0.7259		0.9686	0.9995	0.9925	0.9994	2.962	3.419
40	6	0.5567	0	−0.001551	0.9721	0.5943		0.9481	0.9995	0.9879	0.9995	2.629	3.002
40	7	0.8541	0	0.187	0.9697	0.6848		0.9539	0.9998	0.9796	0.9995	1.502	3.03
40	8	0.7505	0	0.205	0.9924	0.7342		0.9659	0.9998	0.9815	0.9996	1.151	1.975
40	9	0.5366	0	−0.2923	0.9677	0.5975		0.9441	0.9997	0.9768	0.9996	2.293	3.532
80	5	1.783	0	0.7687	0.9578	0.6837		0.933	0.9995	0.9739	0.9995	5.177	4.94
80	6	1.341	0	0.2171	0.9606	0.6314		0.9398	0.9995	0.9859	0.9994	6.06	6.695
80	7	1.286	2.686	−0.1272	0.9592	0.6549		0.926	0.9994	0.9821	0.9993	5.262	7.703
80	8	1.186	0	−0.1146	0.9505	0.5529		0.9435	0.9996	0.9844	0.9995	4.697	7.039
80	9	1.84	0.602	0.6497	0.965	0.6735		0.9322	0.9994	0.9758	0.9994	6.155	6.24

Table 15: The path family, seeds 5–9, full runs (300 s): status and time in seconds of SCIP, SCIP-nolocks, SCIP-extra and Gurobi, and of the separator with remainder directions and base limits (**all-diag-mech**, Part 4C3) or wide limits (**frozen-wide**, Part 4C4) and with whole-row directions and wide limits (**rowdir-wide**). Columns without a stated part are from the part of the block. Empty cells: not run in that part. Solved: number of the runs above it in the column with status optimal or gap limit and an incumbent that passed the independent primal check (Section 8.1). The solved counts are those of Table 4. Status: o optimal, g gap limit, t time limit.

Full run: status and time (s)								
n	Seed	SCIP	SCIP-nolocks	SCIP-extra	Gurobi	Remainder		Whole row
						base	wide	wide
<i>Part 4C3: coupling row $\sum_i y_i \leq 0.8n$, not binding (SCIP-nolocks: Part 5C-a)</i>								
10	5	o 1.827	o 0.3555	o 1.145	g 0.4302	o 2.808		o 3.391
10	6	o 0.9223	o 0.4247	o 0.6687	g 0.2931	o 1.56		o 3.174
10	7	o 0.9557	o 0.4608	o 1.229	o 0.249	o 1.824		o 3.256
10	8	o 1.845	o 0.3206	o 0.8447	o 0.3198	o 2.751		o 3.403
10	9	o 3.724	o 0.6022	o 1.528	g 0.3846	o 1.792		o 3.203
20	5	o 46.79	o 18.42	o 128.5	t 300.2	o 15.57		o 6.612
20	6	t 300.2	o 43.7	t 300.2	t 300.2	o 88.69		o 6.348
20	7	o 69.75	o 6.089	o 12.41	t 300.2	o 59.86		o 6.37
20	8	o 24.59	o 6.073	o 33.9	t 300.2	o 13.32		o 6.631
20	9	o 112.1	o 10.03	o 63.17	t 300.2	o 42.35		o 5.688
40	5	t 300.2	t 300	t 300.2	t 300.2	t 300.2		o 14.87
40	6	t 300.2	t 300	t 300.2	t 300.2	t 300.2		o 14.59
40	7	t 300.2	t 300	t 300.2	t 300.2	t 300.2		o 13.27
40	8	t 300.2	t 300	t 300.2	t 300.2	t 300.2		o 13.82
40	9	t 300.2	t 300	t 300.2	t 300.2	t 300.2		o 13.65
80	5	t 300.2	t 300.1	t 300.2	t 300.2	t 300.2		o 32.11
80	6	t 300.2	t 300.1	t 300.2	t 300.2	t 300.2		o 31.66
80	7	t 300.2	t 300.1	t 300.2	t 300.2	t 300.2		o 33.17
80	8	t 300.2	t 300.1	t 300.2	t 300.2	t 300.2		o 32.29
80	9	t 300.2	t 300.1	t 300.2	t 300.2	t 300.2		o 32.24
Solved		9	10	9	5	10		20
<i>Part 4C4: coupling row $\sum_i y_i \leq c$, binding (SCIP-nolocks: Part 5C-a)</i>								
10	5	o 0.8934	o 0.1038	o 0.8728	o 0.0834	g 3.387		g 3.485
10	6	o 0.4948	g 0.09662	g 0.5615	o 0.03426	g 3.294		g 3.247
10	7	o 0.7076	o 0.1359	o 1.243	o 0.1099	g 3.502		g 3.244
10	8	o 0.9382	o 0.1683	o 0.6494	o 0.1362	o 4.173		g 3.461
10	9	o 0.6253	o 0.138	o 0.4624	o 0.1106	g 3.267		g 3.269
20	5	o 3.948	o 0.2458	o 2.558	t 300.2	o 7.086		g 7.365
20	6	o 11.88	o 0.413	o 3.628	t 300.2	o 7.095		g 6.927
20	7	o 14.08	o 0.4787	o 19.81	t 300.2	o 7.792		o 7.301
20	8	o 4.731	o 0.2868	o 1.929	t 300.2	o 7.323		g 7.29
20	9	o 2.257	o 0.1674	o 1.647	o 0.42	o 7.306		g 6.336
40	5	t 300.2	o 0.8669	o 52.16	t 300.2	o 18.34		g 16.52
40	6	t 300.2	o 1.15	o 86.56	t 300.2	o 18.86		o 16.18
40	7	t 300.2	g 1.219	t 300.2	t 300.2	o 15.9		o 15.63
40	8	t 300.2	g 0.588	o 98.96	t 300.2	o 15.98		g 14.72
40	9	t 300.2	g 2.557	t 300.2	t 300.2	o 20.86		o 15.85
80	5	t 300.2	g 6.448	t 300.2	t 300.2	g 42.07		o 37.84
80	6	t 300.2	g 6.506	t 300.2	t 300.2	o 39.57		o 37.28
80	7	t 300.2	o 169	t 300.2	t 300.2	o 107.5		o 45.9
80	8	t 300.2	o 71.81	t 300.2	t 300.2	o 59.41		o 133.7
80	9	t 300.2	o 15.24	t 300.2	t 300.2	o 37.56		o 39.14
Solved		10	20	13	6	20		20

Table 16: The star family (Part 5S), root runs (node limit 1, 120 s): number k of leaves per star, number n of stars, seed, optimum z^* and the root bound z as a fraction z/z^* of the optimum for SCIP, SCIP-nolocks, SCIP-extra and the separator with blocks of at most four variables and whole-row first directions (**rowdir-star4**), with blocks of at most four variables and the block direction first (**agg-star4**) and with star blocks and the block direction first (**agg-star**); limits in Table 7. The medians per k are those of Table 5. Full runs are in Table 17.

		Root bound / optimum							
							Whole row	Block direction	
k	n	Seed	Optimum	SCIP	SCIP-nolocks	SCIP-extra	≤ 4 var.	≤ 4 var.	stars
4	10	0	0.7426	0.9607	0.5916	0.9374	0.9655	0.9952	1
4	10	1	0.9026	0.9606	0.64	0.9459	0.9742	0.9638	1
4	10	2	1.229	0.8951	0.4781	0.9156	0.9431	0.946	1
4	10	3	1.041	0.9159	0.4715	0.9223	0.9452	0.9649	1
4	10	4	1.193	0.9162	0.5549	0.9489	0.9323	0.941	1
4	20	0	1.954	0.9404	0.6432	0.9567	0.9287	0.9462	1
4	20	1	2.845	0.98	0.4651	0.9551	0.969	0.9797	1
4	20	2	2.294	0.9787	0.5809	0.9592	0.9734	0.9856	1
4	20	3	2.274	0.9477	0.4859	0.9392	0.9523	0.9477	1
4	20	4	1.769	0.9666	0.4981	0.9288	0.959	0.9626	1
8	10	0	2.522	0.8773	0.5602	0.8534	0.8748	0.8777	1
8	10	1	2.012	0.8113	0.593	0.834	0.8057	0.831	1
8	10	2	2.598	0.924	0.5042	0.9265	0.9354	0.9422	1
8	10	3	2.492	0.9041	0.6946	0.8785	0.9086	0.9323	1
8	10	4	3.086	0.9434	0.5757	0.9419	0.9316	0.9223	1
8	20	0	5.378	0.7595	0.5849	0.8581	0.7757	0.7838	1
8	20	1	5.444	0.7463	0.5559	0.8139	0.7537	0.7596	1
8	20	2	5.236	0.8191	0.6062	0.8662	0.8362	0.8362	1
8	20	3	4.692	0.7802	0.6172	0.8585	0.8032	0.8018	1
8	20	4	5.635	0.7672	0.5308	0.8482	0.7787	0.7752	1
16	10	0	6.151	0.6814	0.5332	0.8117	0.711	0.7483	1
16	10	1	5.991	0.6863	0.5357	0.796	0.724	0.7016	1
16	10	2	5.508	0.6965	0.5457	0.7677	0.7444	0.7599	1
16	10	3	6.1	0.7539	0.5191	0.8018	0.7326	0.7659	1
16	10	4	5.303	0.6219	0.5109	0.6883	0.6163	0.6279	1
16	20	0	10.58	0.6174	0.5632	0.7104	0.611	0.6256	1
16	20	1	12.62	0.5944	0.548	0.7323	0.6245	0.6223	1
16	20	2	12	0.6312	0.5868	0.7278	0.6208	0.643	1
16	20	3	11.66	0.6794	0.5908	0.7278	0.6602	0.6859	1
16	20	4	13.37	0.6232	0.5794	0.7432	0.6422	0.636	1

Table 17: The star family (Part 5S), full runs (300 s): status and time in seconds of SCIP, SCIP-nolocks, SCIP-extra, Gurobi and the three separator modes of Table 16. [†]One of the 8 runs whose worker process was stopped at the hard limit of 360s during a load spike on the host (status `process_timeout`, Section 8.6); such a run recorded no status, time or bound and counts as unsolved. Solved: number of the runs above it in the column with status optimal or gap limit and an incumbent that passed the independent primal check (Section 8.1). The solved counts are those of Table 5. Status: o optimal, g gap limit, t time limit.

Full run: status and time (s)									
k	n	Seed					Whole row	Block direction	
			SCIP	SCIP-nolocks	SCIP-extra	Gurobi	≤ 4 var.	≤ 4 var.	stars
4	10	0	o 2.151	o 0.5735	o 0.7135	g 0.1642	o 39.2	o 40.88	o 39.71
4	10	1	o 0.4967	o 0.6363	g 0.9776	g 0.2641	o 49.05	o 43.67	o 40.77
4	10	2	o 2.04	g 2.128	o 1.633	g 0.3627	o 51.62	o 49.29	o 42.25
4	10	3	o 0.5753	o 0.8428	o 0.8658	g 0.3071	o 40.72	o 40.94	o 39.79
4	10	4	o 0.7029	o 1.227	g 1.211	g 0.6012	o 39.34	o 35.97	o 42.37
4	20	0	o 3.38	o 7.799	o 3.983	g 0.5793	o 66.34	o 61.58	g 61.16
4	20	1	g 1.564	o 4.51	o 6.012	g 0.6082	g 62.17	g 61.53	t 300.1
4	20	2	t 300.1	g 2.801	o 12.56	g 0.525	o 61.3	o 61.31	o 61.19
4	20	3	— [†]	g 11.41	o 6.355	— [†]	o 64.72	o 66.37	o 60.64
4	20	4	— [†]	o 3.801	o 4.38	g 1.733	o 62.44	o 65.24	o 204.8
Solved, $k = 4$			7	10	10	9	10	10	9
8	10	0	o 36.8	o 20.28	t 300	g 0.6548	— [†]	o 255.6	o 63.33
8	10	1	t 300	o 10.67	t 300.1	g 0.5977	t 300	t 300	o 60.84
8	10	2	t 300	o 32.52	o 155.9	g 0.5119	t 300	g 130.7	o 60.9
8	10	3	o 5.405	o 7.166	g 256.6	g 0.879	g 67.33	g 63.88	o 62.48
8	10	4	o 12.1	o 59.38	o 8.661	g 0.8555	o 69.77	o 79.56	g 62.14
8	20	0	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1	g 67.1
8	20	1	t 300.1	t 300	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1
8	20	2	t 300.1	t 300.1	t 300.1	g 126.5	t 300.1	t 300.1	g 143.6
8	20	3	t 300.1	t 300.1	t 300.1	g 1.189	t 300.1	t 300.1	o 83.53
8	20	4	t 300.1	t 300.1	t 300.1	g 75.3	t 300.1	t 300.1	g 77.68
Solved, $k = 8$			3	5	3	8	2	4	9
16	10	0	t 300.1	t 300.1	t 300.1	g 0.9829	t 300.1	t 300.1	g 175.8
16	10	1	t 300.1	t 300.1	t 300.1	g 1.299	t 300.1	t 300.1	t 300.1
16	10	2	t 300	t 300	t 300.1	g 7.968	— [†]	— [†]	t 300
16	10	3	t 300	t 300	t 300	g 100.9	t 300.1	t 330.2	— [†]
16	10	4	— [†]	t 300	t 300	g 1.938	t 300	t 300.1	t 305.2
16	20	0	t 300.1	t 300.1	t 300.1	t 300.1	t 300	t 300.1	t 300.1
16	20	1	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1	t 300	t 300.1
16	20	2	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1	t 300
16	20	3	t 300.1	t 300.1	t 300.1	t 300	t 300.1	t 300.1	t 300.1
16	20	4	t 300	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1	t 300.1
Solved, $k = 16$			0	0	0	5	0	0	1
Solved, total			10	15	13	22	12	14	19